

**Formation :**

# Containers, enjeux, usages et solutions

**Formateur**

Ahmed Hosni

Technology manager

Contact: ahmedhosni.contact@gmail.com



# Containers, enjeux, usages et solutions

---

Un État de l'art des solutions d'orchestration de containers et de leur écosystème pour mettre en œuvre une plateforme de type CaaS (Container as a Service).

Il apporte des réponses sur le fonctionnement, la mise en place ou l'utilisation de conteneurs dans une organisation et apporte des conseils pour leur usage.

## **Participants**

Architectes, responsables des infrastructures IT, chefs de projet, administrateurs système et/ou réseau ou développeurs.

## OBJECTIFS PEDAGOGIQUES

---

1. Détailler les différents aspects de la technologie de containerisation, son écosystème
  2. Découvrir le fonctionnement de Kubernetes, ses composants internes et externes
  3. Comprendre les interactions avec le Cloud privé/public et le legacy.
  4. Appréhender les bénéfices et les limites des architectures micro-services en termes techniques et organisationnels
-

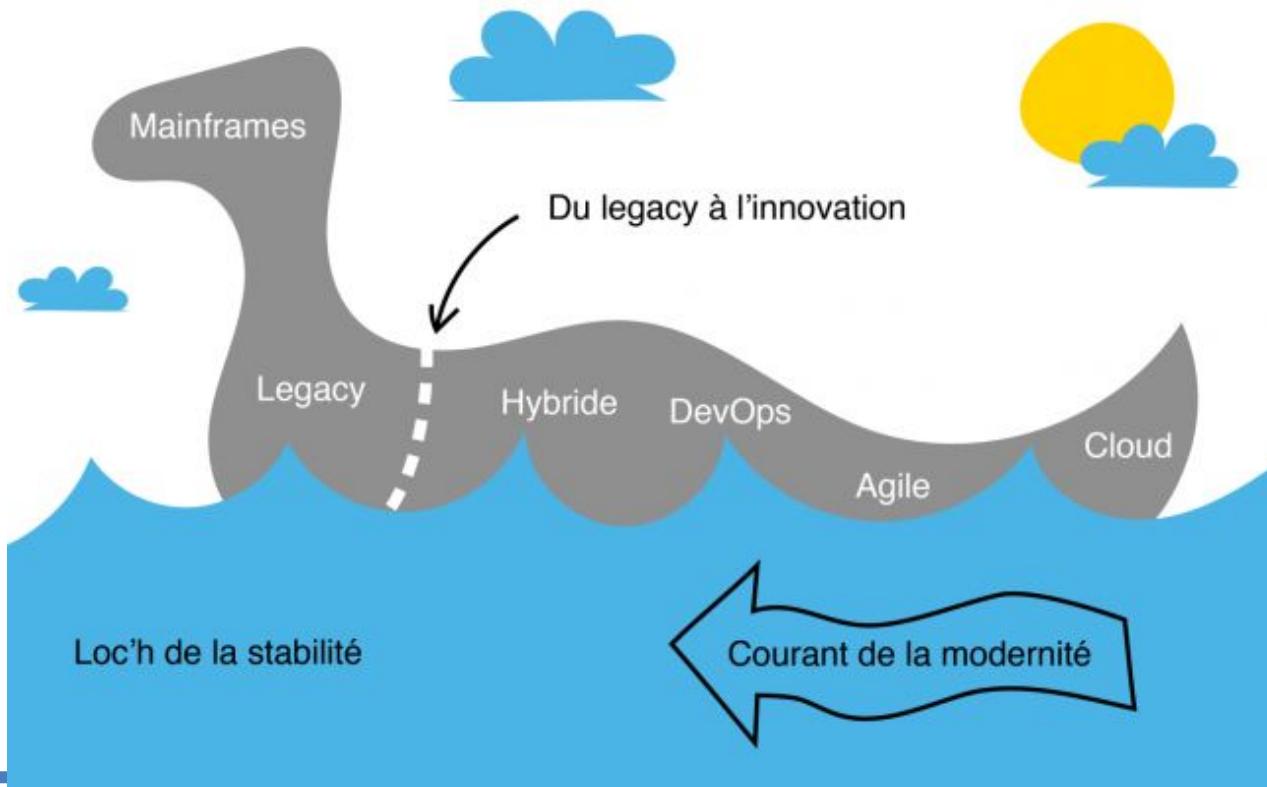
# Les fondamentaux

# Agenda

- 
1. IT Bimodale et DevOps.
  2. Application en 12 facteurs, intégration, déploiement continu (CI/CD), applications Cloud-native.
  3. SaaS, PaaS, IaaS, Stockage objet et bloc. Cloud privé, public, hybride : problématique du lock-in.
  4. Architecture élastique, Cattle versus Pet, Infrastructure as Code.
  5. Outils existants (Terraform, Ansible). Apport des containers versus Machines Virtuelles.
-

# IT Bimodale et DevOps

Bimodal IT : le meilleur des deux mondes informatiques



# IT Bimodale et DevOps



## Bimodal IT = Marathon Runners + Sprinters

Think  
Marathon Runner

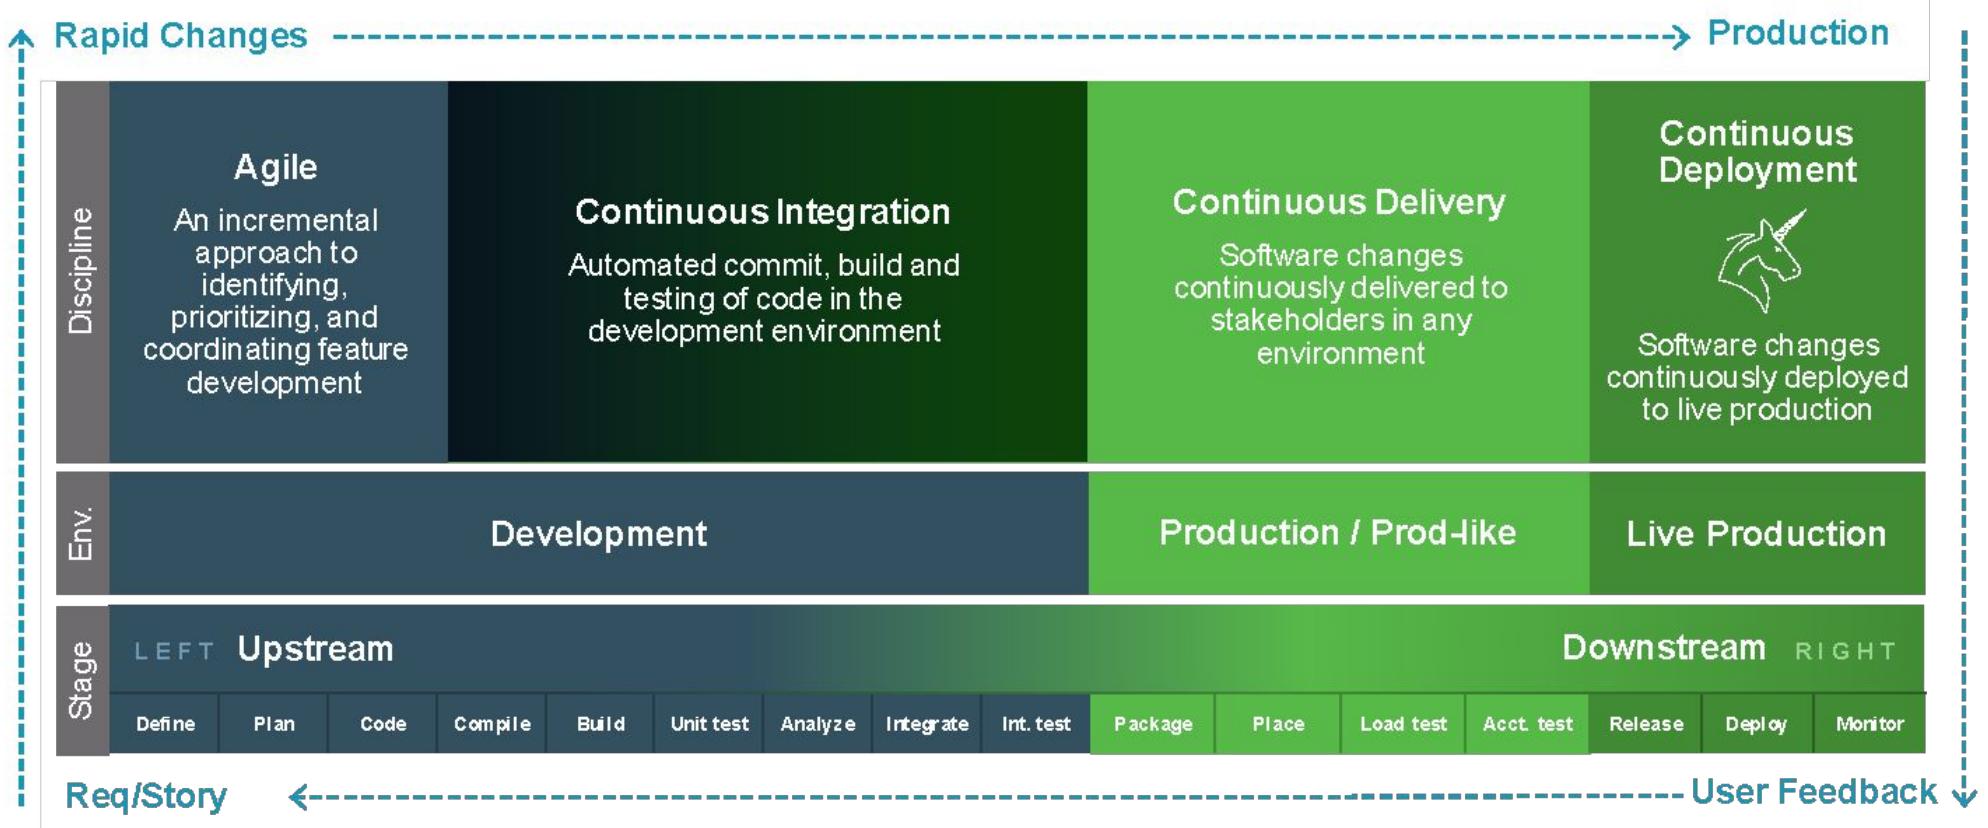


Mode 1	Mode 2	
Reliability	Goal	Agility
Price for performance	Value	Revenue, brand, customer experience
Waterfall, V-Model, high-ceremony IID	Approach	Agile, Kanban, low-ceremony IID
Plan-driven, approval-based	Governance	Empirical, continuous, process-based
Enterprise suppliers, long-term deals	Sourcing	Small, new vendors, short-term deals
Good at conventional process, projects	Talent	Good at new and uncertain projects
IT-centric, removed from customer	Culture	Business-centric, close to customer
Long (months)	Cycle times	Short (days, weeks)

Think  
Sprinter



# Context DevOps



# Application en 12 facteurs



**I. Codebase** Une base de codes suivie dans l'SCM, de nombreux déploiements.

**II. Dependencies** Déclarer et isoler explicitement les dépendances

**III. Config** Stockage de la configuration dans l'environnement

**IV. Backing services** Traiter les backing services comme des ressources rattachées

**V. Build, release, run** Étapes de construction et d'exécution sont strictement séparées

**VI. Processes** Exécuter l'application comme un ou plusieurs stateless processus

# Application en 12 facteurs



**VII. Port binding** Services d'exportation via le port binding

**VIII. Concurrency** Mise à l'échelle via le modèle de processus

**IX. Disposability** Maximisez la robustesse avec un démarrage rapide et un arrêt gracieux.

**X. Dev/prod parity** Maintenir le développement, la mise en scène et la production aussi semblables que possible

**XI. Logs** Traiter les journaux comme des flux d'événements

**XII. Admin processes** Exécuter des tâches administratives ou de gestion en tant que processus ponctuels

## Qu'est ce que l'intégration Continue (Ci) ?

---

Continuous Integration is a software development *practice* where members of a team integrate their work *frequently*. Usually each person integrates at least daily - leading to *multiple* integrations per day.

– Martin Fowler - Continuous Integration

---

# Continuous



- Each integration is verified by an **automated** build (including test)
- Integrate code **often**, at least daily, to make integration a **non-event**
- **Continuously** build and integrate
  - Trigger builds automatically based on commit and merge actions and success of upstream builds

# Common pitfalls

## **Not putting everything into the code repository:**

Everything that's needed to build and configure the application and the system should be in your repository.

## **Not automating the build process:**

Manual steps create opportunities for mistakes and leave steps undocumented.

## **Not fixing broken builds right away:**

A key goal of CI is having a stable build from which everyone can develop.

If the build can't be fixed in a few minutes, the offending change should be identified and reverted.

## **Not triggering quick tests on every change:**

Full end-to-end tests are necessary, but quick tests (unit tests) are also important to enable fast feedback.

## **Not merging into trunk often enough:**

Many organizations have automated tests and builds, but don't enforce a daily merge into trunk.

This leads to long-lived branches that are harder to integrate, and to long feedback loops for the developers.

# Ways to measure CI

Factor to test	What to measure
Code commits trigger a build of the software	The percentage of code commits that result in a software build without manual intervention.
Code commits trigger a series of automated tests	The percentage of code commits that result in a suite of automated tests being run without manual intervention.
Automated builds and tests are executed successfully every day	The percentage of automated builds and the percentage of automated tests that are executed successfully every day.
Current builds are available to testers for exploratory testing	The availability of builds to testers, or its inverse, namely the unavailability of builds to testers.
Developers get feedback from the acceptance and performance tests every day	The availability of feedback to developers from acceptance and performance tests—that is, the percentage of tests that provide feedback that is available to developers within a day.
Broken builds are fixed immediately	The time it takes between the build breaking and having it fixed, either with a check-in that fixes the problem, or by reverting the breaking change.

# Continuous Delivery (Cd)

*How long would it take to your organization to deploy a change that involves just one single line of code?*

Mary Poppendieck



Tom Poppendieck

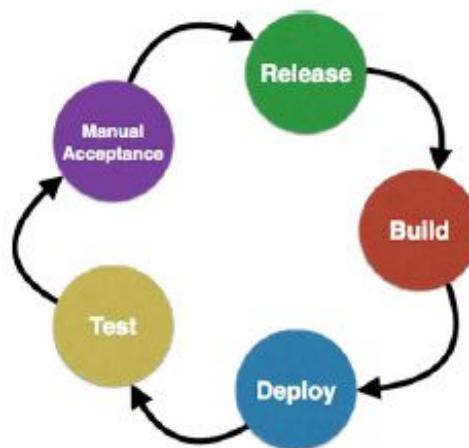


- Reduce deployment risks
- Allow more frequent user feedback
- Make progress observable for everyone

# Why Use Continuous Delivery?

Continuous Delivery is the next step after Continuous Integration:

- Each change to the system **can** be released for **production**
- Delivery can be done at **any** time, on **any** environment
- CD Pipeline produces a validated deployable artifact



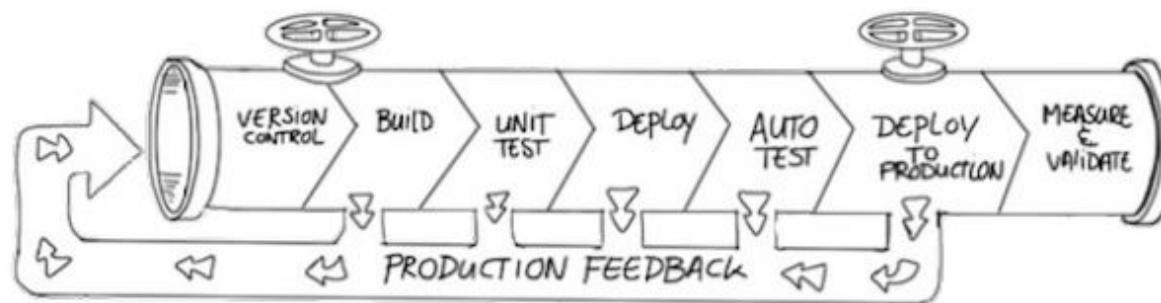
Your team prioritizes keeping the software **deployable** over working on new features

– Martin Fowler



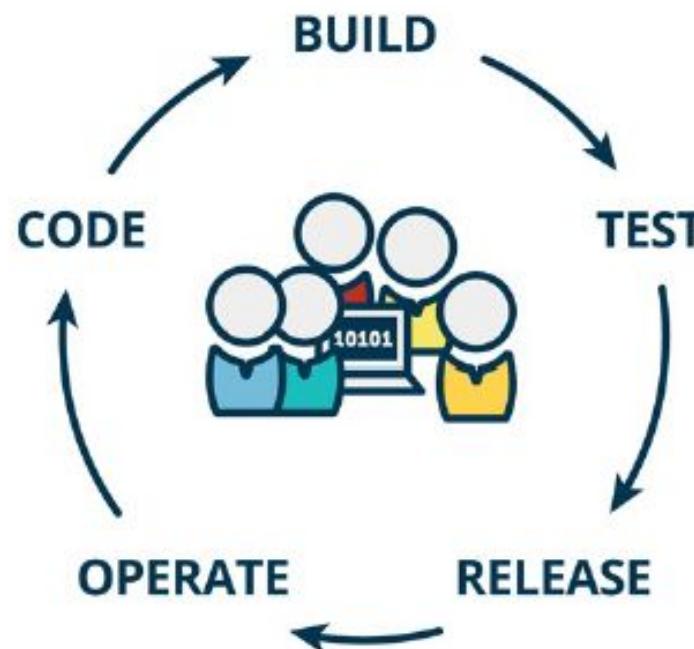
# How to Do Continuous Delivery?

- Have a collaborative working relationship with **everyone** involved
  - DevOps is a popular model i.e. 2 Pizza Team
- Use **Delivery Pipelines**, which are automated implementations of your application's build lifecycle process. This is discussed more later in this course.



# What is Continuous Deployment ?

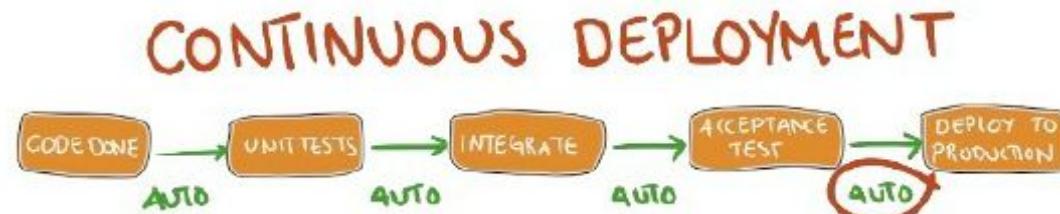
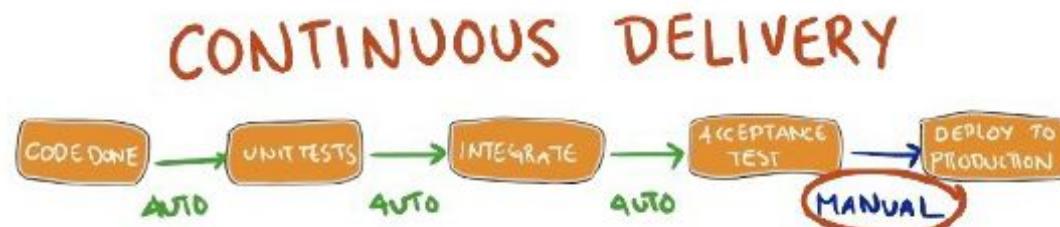
- Continuous Deployment is the last step after Continuous Delivery.
  - Used by Amazon and Netflix
- Any change to the system is automatically pushed to production.
  - Not appropriate in all scenarios



# Continuous Delivery Versus Deployment

These two terms are easily confused:

- Continuous Delivery needs a **manual** step before it is deployed to production.
- Continuous Deployment deploys the modified and tested software to production automatically.
  - Full automated testing and reporting is essential to trusting Continuous Deployment



# WHAT IS A CI/CD PIPELINE?

- Model of the value stream of your software production business
- Defines and models build, testing, and/or deployment processes as continuous workflows
- Fundamental to achieving Continuous Delivery
  - Continuous Delivery is necessary to achieve Continuous Deployment



## WHY USE A CI/CD PIPELINE?

---

- Using a "CD Pipeline" gives same outcome as a value or supply chain in other industries
    - This model is easier to analyze and manage
  - It is driven by the **fail-fast** concept
    - "Fast is cheap": Pipeline detects failures as early as possible
-

- Composed of *stages*, mapped to atomic steps
  - Scoped to cacheable result
  - You should never have to rerun a successful step, just retrieve the cached result
- Transition between 2 stages is a *gate*
  - Manual or Automated
  - A transition that confirms a requirement is met is a **Validation Gate**
  - *gates* trigger next stage
- Stages can be run sequentially and/or in parallel
  - Thus *gates* can branch and even be conditional
- Driven by atomic change: generally SCM single commit

- Deployment is the set of actions that makes a software system ready for usage
- Deploying software is done in environments
- Environments are isolated runtimes with specific properties:
  - Production environments are where software lives for end users
  - Testing environments are locations where tests are run against software
  - Disaster Recovery Environment is an emergency location from which software can be run after a disaster occurs on a production environment

# Automated Deployment

---

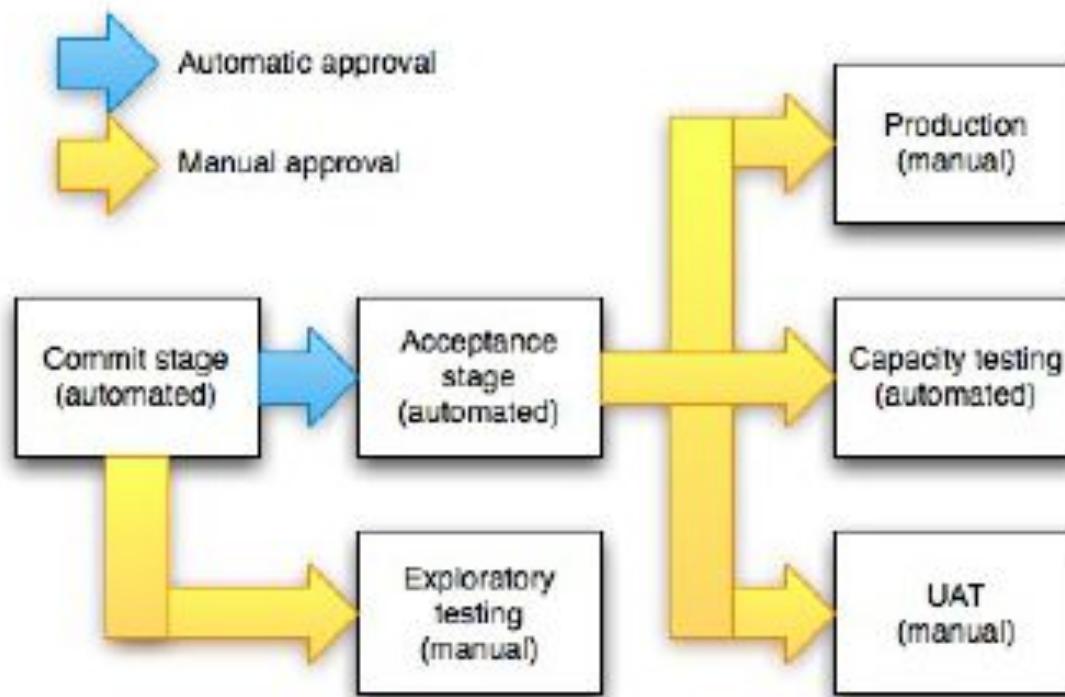
Using *automation* for a deployment stage:

- Make deployment repeatable and less error-prone
  - Bring ROI when deploying on new environment
  - Make frequent releases possible by bringing *confidence*
  - Still can require a manual (read: *human*) trigger to have validation gate
-

# Distributing

- The *build* stage:
  - Fully automated stage.
  - Automated Gates to trigger: event from the SCM (commit, merge, pull-request)
- Easy to distribute to be efficient:
  - Parallelize builds of different target architectures or OSes (ARM, Windows, Android...)
  - Parallelize builds on independent modules of your code before assembling (Map & Reduce)
  - Unit and Integration tests are tied to the code: run them in parallel !
- The *deployment* stage:
  - Should be automated, but can be manual
  - Gates can be manual or automated
- Easy to distribute the automated ones
- You can also distribute manual gate-based if they are on different environments

# An Optimized Pipeline Example



# Always Build Upon Good Foundations

- Introducing Binary reuse: "*Only Build Your Binaries Once*"
  - Compiling again a given binary violates the efficiency rule, (shorter build time)
  - Ensure same version of codebase is used across the pipeline
  - If the binary does not comply with "deployable" rule, then it must fail a test
  - Focus on failing-fast

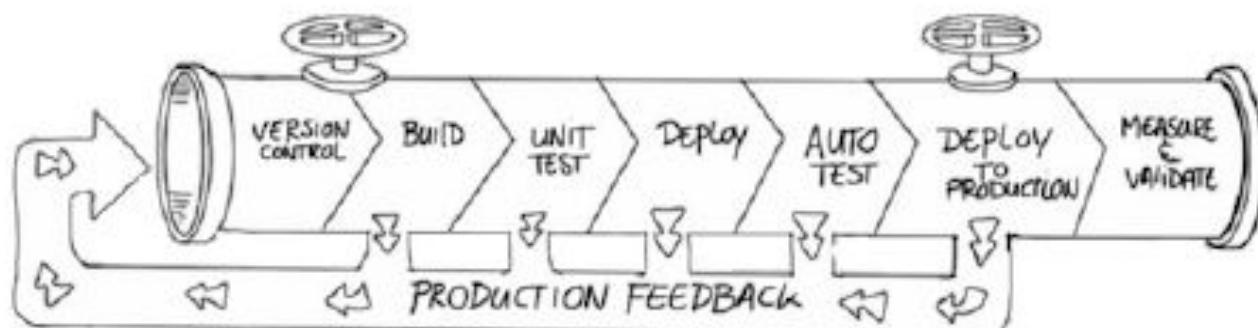
# Why Have a Feedback Loop?

---

- *Context*: Continuous Delivery implemented with CD Pipeline
  - *Target*: Focusing on fast feedback to react quickly
  - *Need* systems that give *quick* feedback to pipeline's *actors*
    - Key to success on the *efficiency* rule
    - Give *confidence* to the actor: they know when it breaks
-

# La Feedback Loop?

- The feedback loop is the set of tools that:
  - Provides **status** of each step of the pipeline
  - Gives the **right** message to the **right** actor of the pipeline
  - Is also a great source of **measurements**



## SaaS, PaaS, IaaS

---

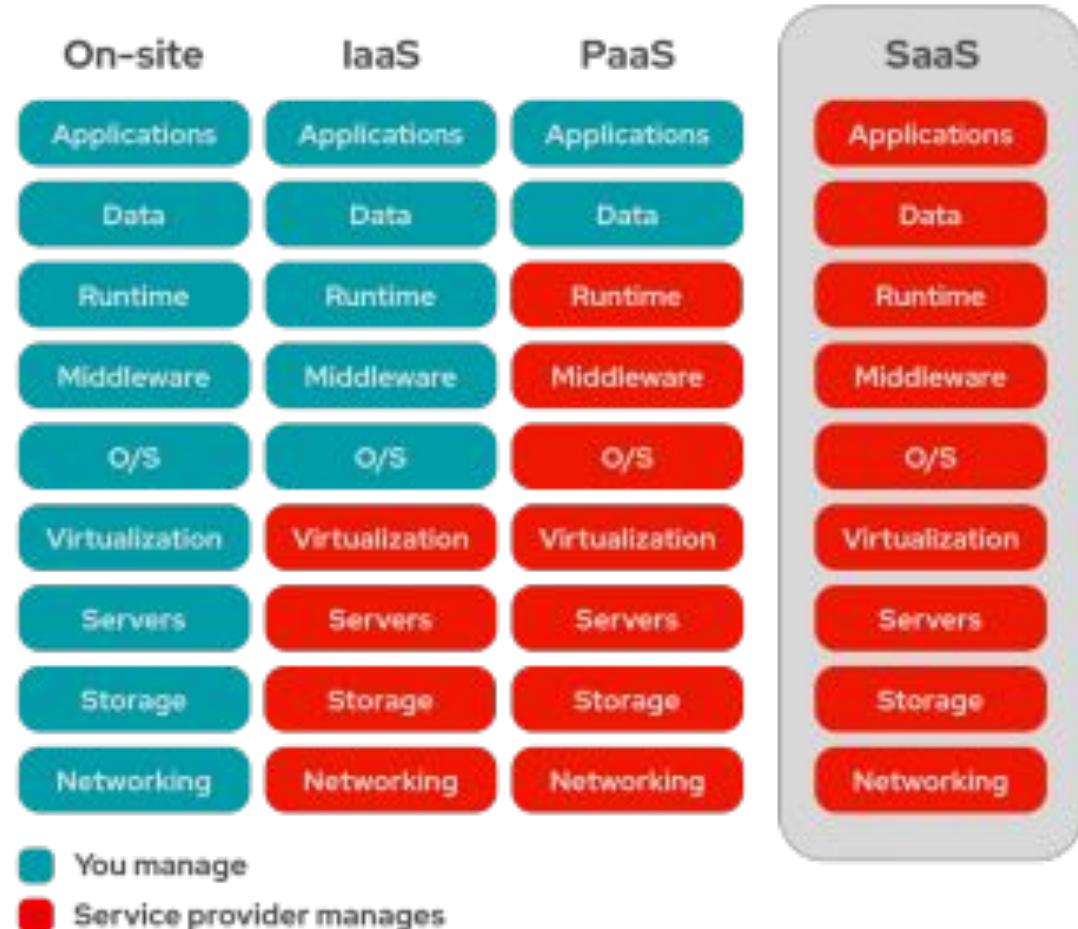
# SaaS



Le SaaS est un service qui fournit une application logicielle - que le fournisseur de services cloud gère - à ses utilisateurs.

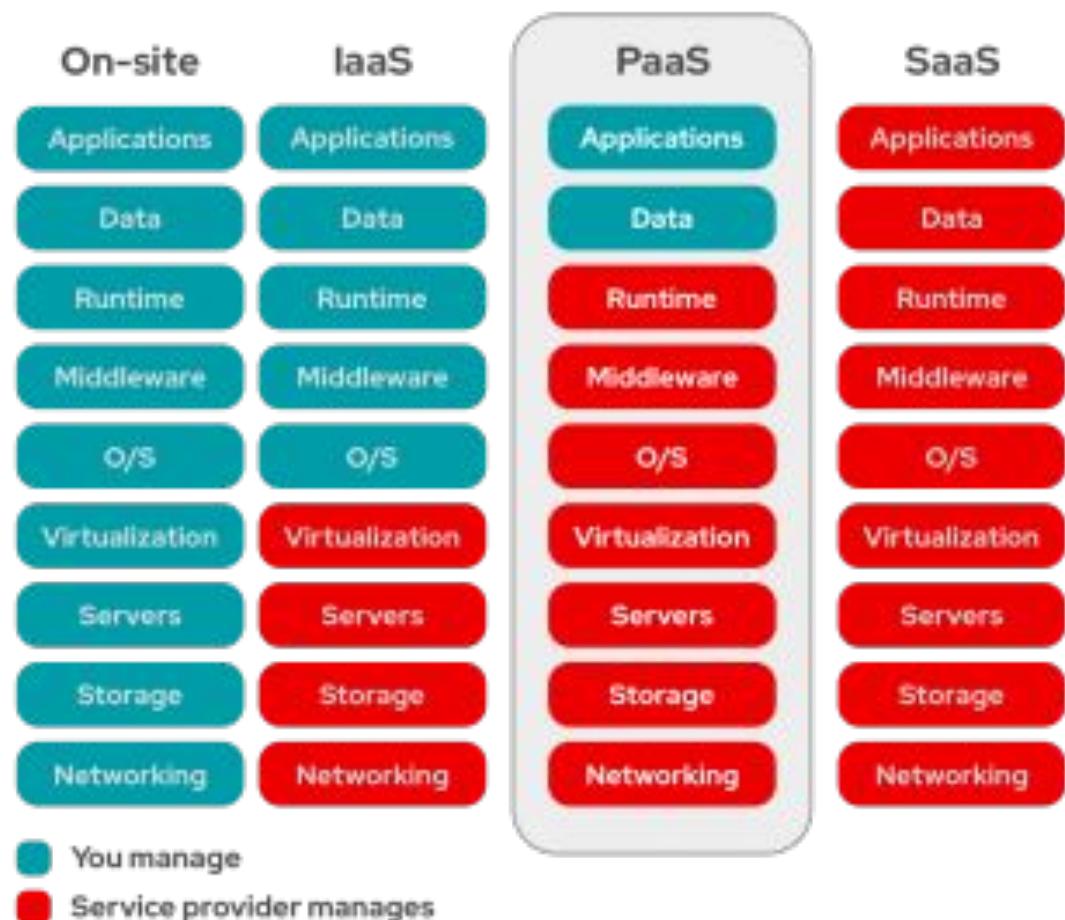
En générale, les applications SaaS sont des applications Web ou des applications mobiles auxquelles les utilisateurs peuvent accéder via un navigateur Web.

Les mises à jour logicielles, les corrections et toute maintenance logicielle générale sont prises en charge pour l'utilisateur et se connectent aux applications cloud via un tableau de bord ou une API.



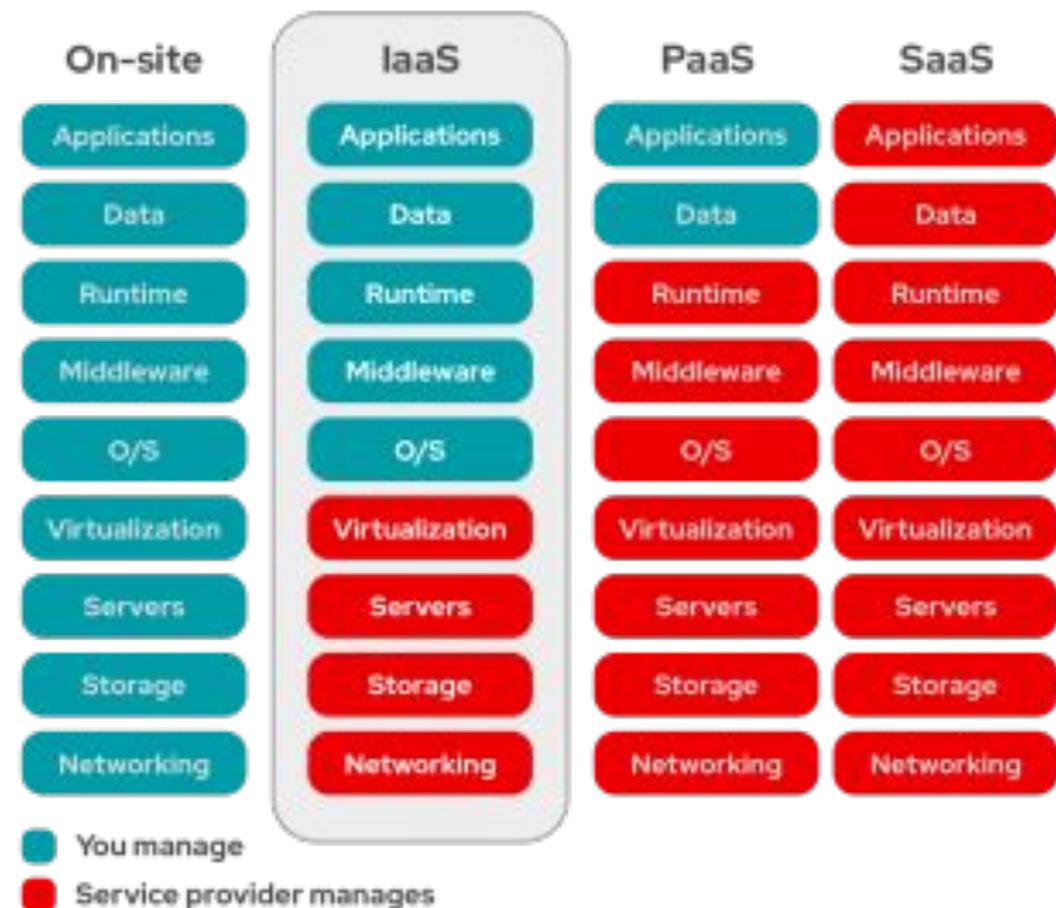
# PaaS

PaaS signifie que le matériel et une plate-forme logicielle d'application sont fournis et gérés par un fournisseur de services cloud externe, mais que l'utilisateur gère les applications exécutées au-dessus de la plate-forme et les données sur lesquelles l'application s'appuie.



# IaaS

- IaaS signifie qu'un fournisseur de services cloud gère l'infrastructure pour vous (les serveurs, le réseau, la virtualisation et le stockage de données) via une connexion Internet.
- L'utilisateur a accès via une API ou un tableau de bord, et loue essentiellement l'infrastructure.



# Cloud privé, public, hybride :

## problématique du lock-in.

---

# Cloud privé



- Les clouds privés sont vaguement définis comme des environnements cloud uniquement dédiés à un seul utilisateur final ou à un seul groupe, où l'environnement s'exécute généralement derrière le pare-feu de cet utilisateur ou groupe.
- Tous les clouds deviennent des clouds privés lorsque l'infrastructure informatique sous-jacente est dédiée à un seul client avec un accès complètement isolé.
- Mais les clouds privés ne doivent plus provenir d'une infrastructure informatique sur site.
- Les organisations construisent maintenant des clouds privés sur des centres de données loués et appartenant aux fournisseurs situés hors site, ce qui rend obsolètes les règles de localisation et de propriété.

# Cloud public

- Les clouds publics sont des environnements cloud généralement créés à partir d'une infrastructure informatique n'appartenant pas à l'utilisateur final.
  - Certains des plus grands fournisseurs de cloud public incluent Alibaba Cloud, Amazon Web Services (AWS), Google Cloud, IBM Cloud et Microsoft Azure.
  - Les clouds privés ne doivent plus provenir d'une infrastructure informatique sur site. Les organisations construisent maintenant des clouds privés sur des centres de données loués et appartenant aux fournisseurs situés hors site, ce qui rend obsolètes les règles de localisation et de propriété.
-

# Cloud hybride

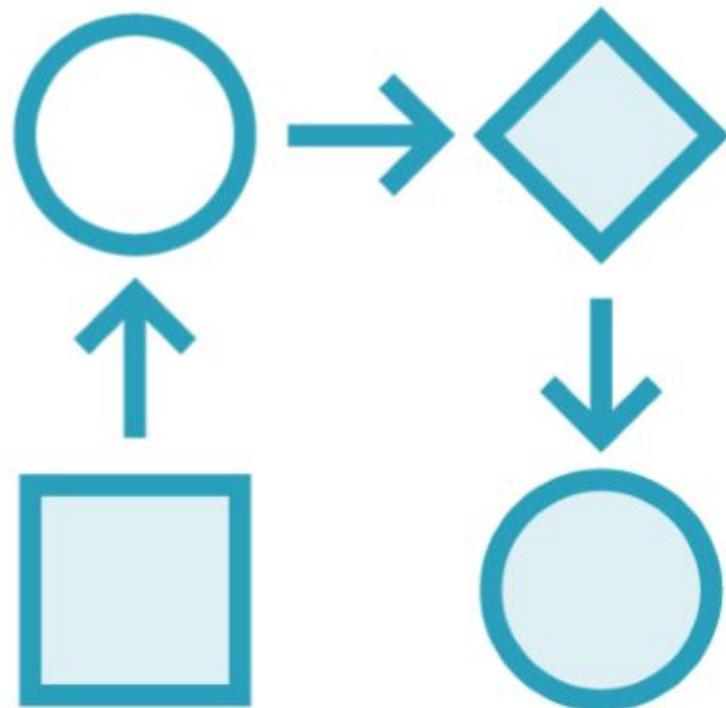


- Un cloud hybride est un seul environnement informatique transparent créé à partir de plusieurs environnements connectés via des réseaux locaux (LAN), des réseaux étendus (WAN), des réseaux privés virtuels (VPN) et / ou des API.
- Les caractéristiques des clouds hybrides sont complexes et les exigences peuvent différer selon la personne à qui vous le demandez. Par exemple, un cloud hybride peut inclure:
  - ✓ Au moins 1 cloud privé et au moins 1 cloud public
  - ✓ 2 nuages privés ou plus
  - ✓ 2 clouds publics ou plus
  - ✓ Un environnement bare-metal ou virtuel connecté à au moins 1 cloud public ou privé

# Infrastructure as Code

---

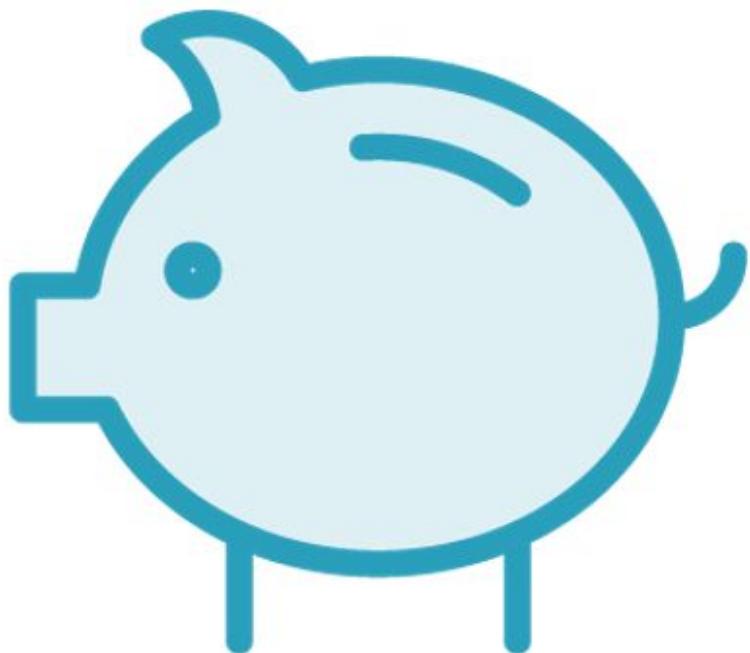
# IaC: le Pipeline



The workflow of moving a project through stages of work

Common names: Release, Build, Deployment

## Le gain pour le business et les clients



Change and release management  
Improved security  
Reduced recovery time

# Le gain pour le business et les clients

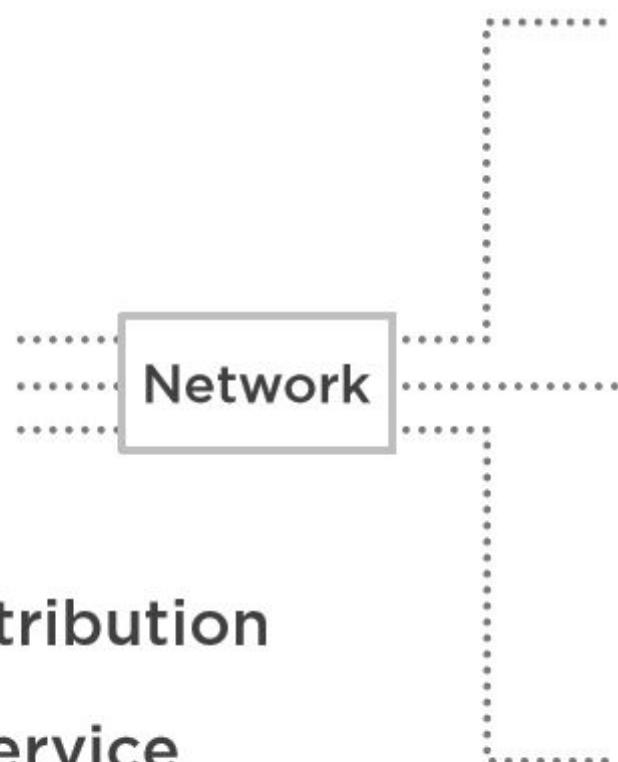
Authentication

Messaging

Switch/Router



Network



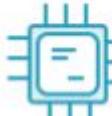
Content Distribution

Customer Service

System 1



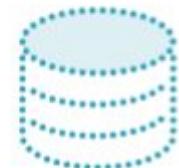
CPU



RAM



Storage



System 2



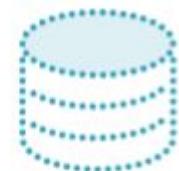
CPU



RAM



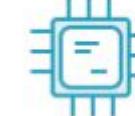
Storage



System 3



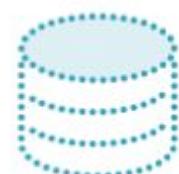
CPU



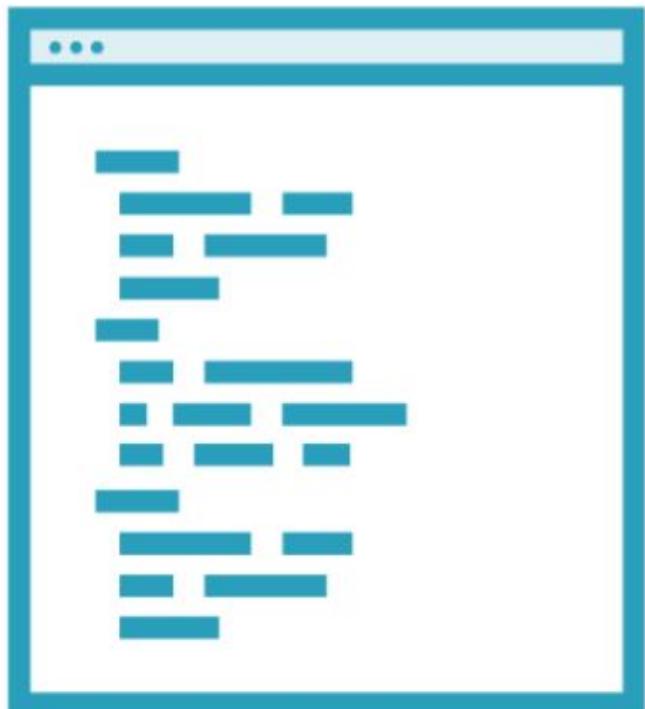
RAM



Storage



# La configuration



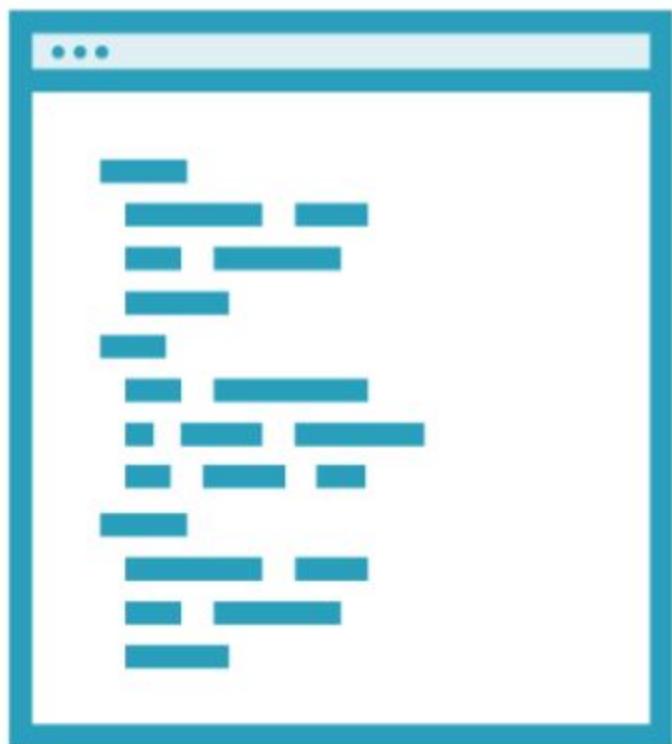
- | **Read by an automation process**
- | **Contains configuration properties**
- | **Imperative – How to do**
- | **Declarative – What you want**

# Impératif – Démarrer une voiture



1. Open driver side car door
2. Get in car
3. Locate ignition switch
4. Insert key into ignition switch
5. Turn ignition
6. Slightly press on the gas pedal

# Déclaratif – Démarrer une voiture



1. Start the car

# Cibles de livraison Push/Pull

Push

Best for testing, great production solution

Deployment



Pull

Great production solution

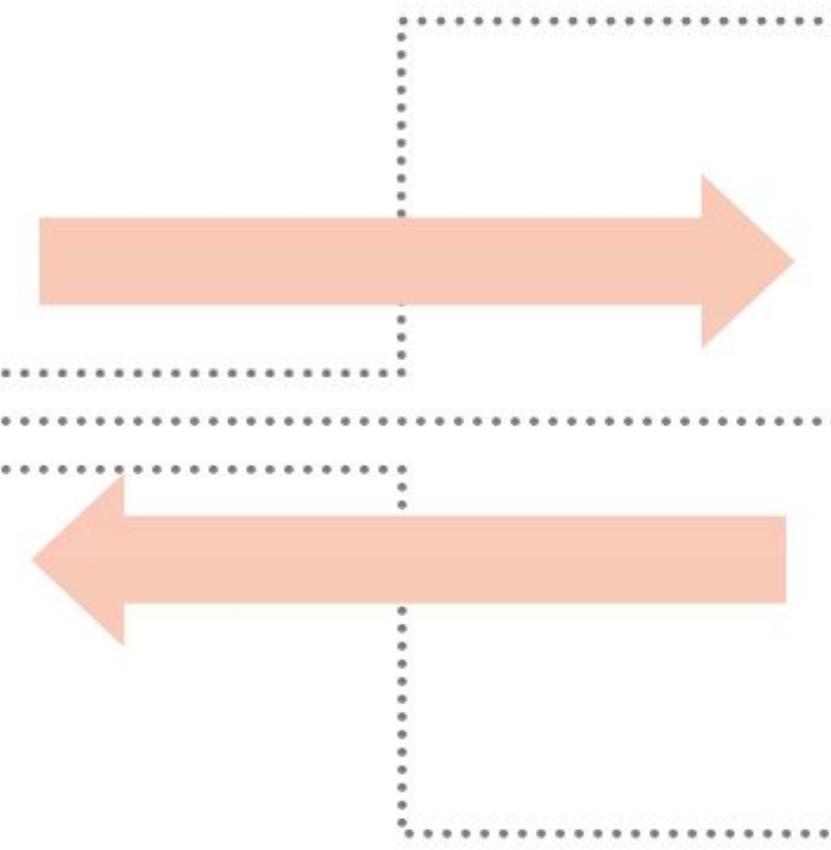
Target 1



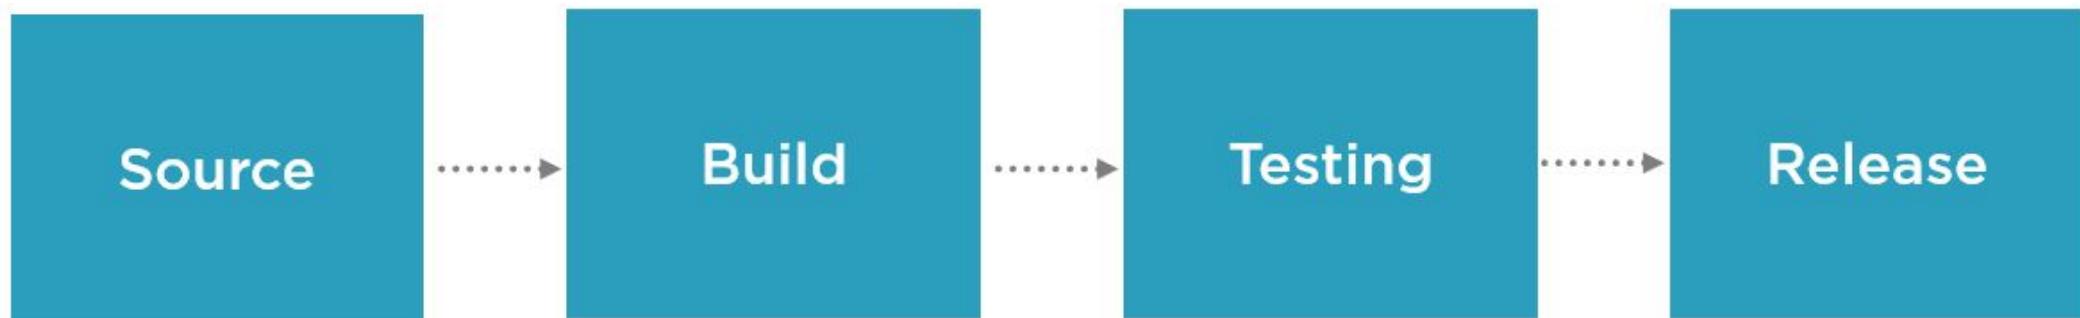
Target 2



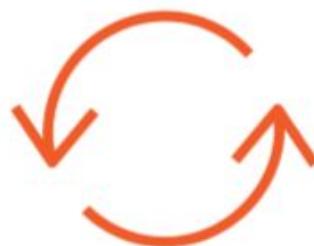
Target 3



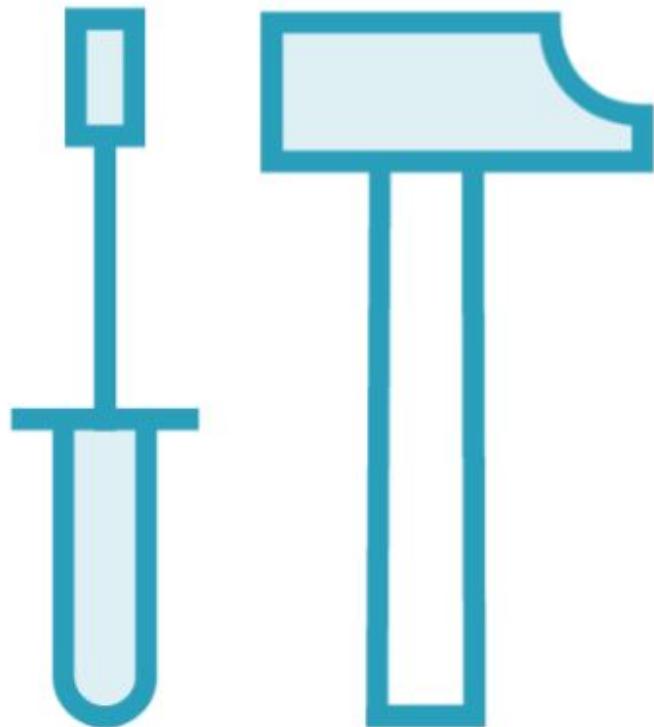
# Cibles de livraison Push/Pull



Automation and  
orchestration -  
Monitoring and  
validating



## Outils CCA - Continuous Configuration Automation



Referred to as the DevOps Toolchain  
Tools that automate and orchestrate  
May be specific to a platform  
Need to meet your security and compliance needs  
Needs to scale

# Une nouvelle vision de “l'échec”



Due to inconsistency procedures, lack of automation, lack of change control  
- IT Ops managers are slow to make change because of failure

Infrastructure From Code provides a reliable, repeatable process that can learn.

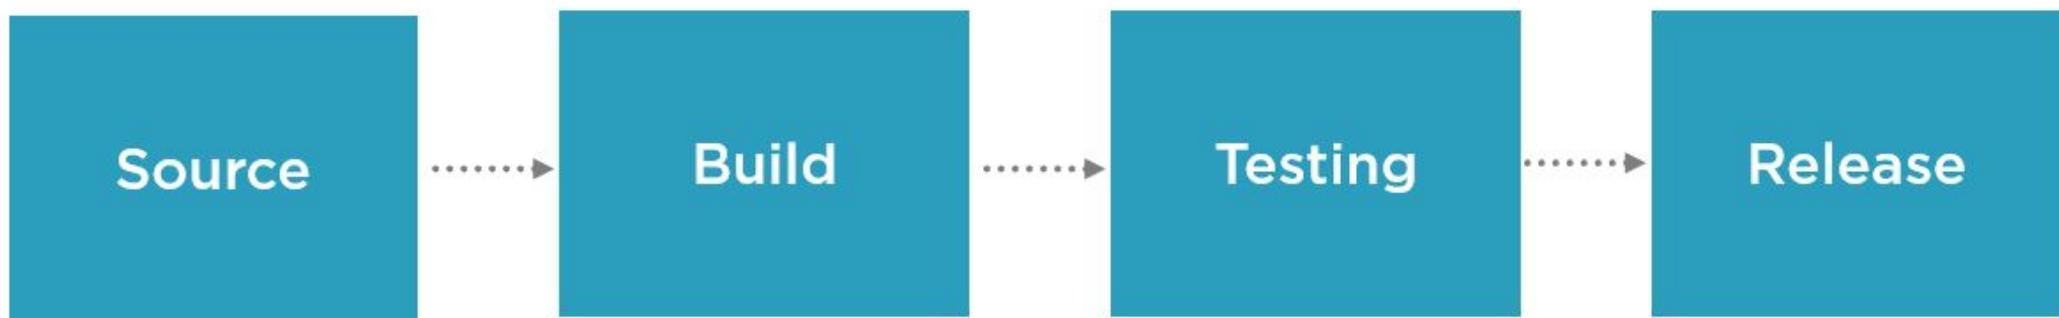
Rush to failure!

# What Does Release Mean?



- After the process successfully completes
- Configuration change is in Source
- Build process completes
- Testing has passed
- A build artifact is ready to be deployed

# Le pipeline de Release



GitHub  
Apache Subversion  
Microsoft Visual Studio  
Team Foundation server

Team City  
AppVeyor  
Jenkins  
Microsoft Visual Studio

ServerSpec, Rspec  
Chef InSpec  
Pester  
Test-Kitchen  
Vagrant

Environment and  
situation  
dependent

# Continuous Delivery (CD)



Also known as Continuous Deployment with slight changes

Best with immutable servers – every server matches exact

Server can be re-deployed entirely if failure occurs

Example: Web servers in a farm

# Incremental Deployment



**For servers that must be maintained over a long period of time**

**For servers that cannot be re-deployed**

**Configuration changes are delivered as updates**

**Example: Database, Domain Controllers**

# Continuous Integration



**Bringing the process all together with a workflow (orchestration)**

**Add quality gates (testing) and approval**

**Can suspend new changes until failures are resolved**

# Démarrer avec l'Infrastructure as Code



**Start small**

**Start with configuration management**

**Add source control**

**Add testing**

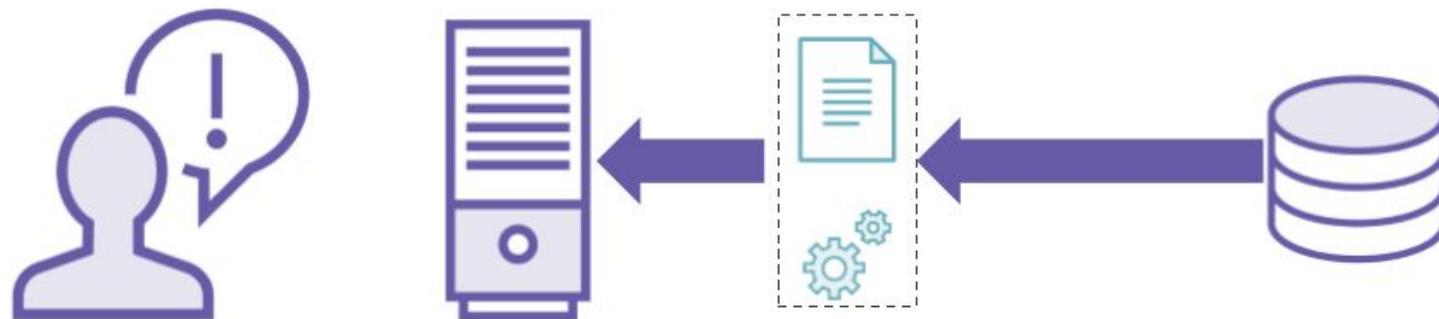
**Build the Pipeline**

# Infrastructure Immutable

Immutable

Adjective:

Unchanging over time, or unable to be changed



# Pets vs Cattle



Mittens!



XB306-1

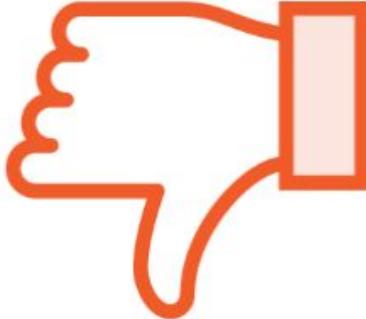
XB306-2

XB306-3

...

# Pets vs Cattle

Cats	Cattle
Kept alive at all costs	Expendable
Need manual intervention	Work “out of the box”
Difficult to Scale	Easy to Scale
High-stress	Low-stress



# Autres avantages

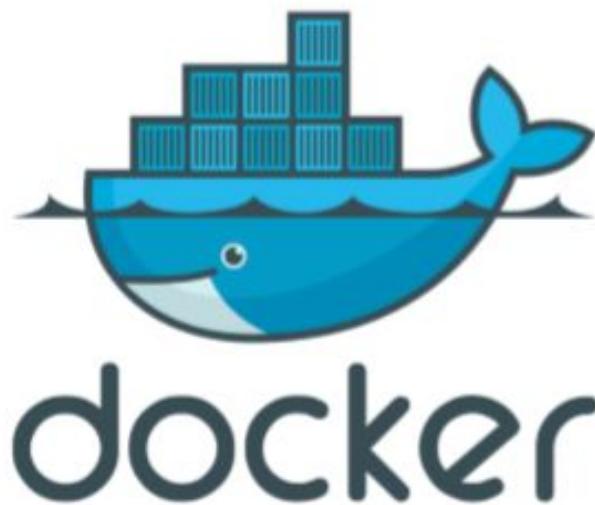
Testable Infrastructure

Reproduce Production

Unit of Deployment

Confident Changes

# Docker?



~~No Windows support~~

Application re-architecture

Orchestration

Sometimes worth it, sometimes not

# Packer!



Utilize native tools

Integrate with configuration management

Easy to convert to Docker

# Un Template Packer

```
{  
  "builders": [  
  ],  
  "provisioners": [  
  ],  
  "post-processors": [  
  ]  
}
```



# Un Template Packer

```
{  
  "type": "amazon-ebs",  
  "name": "MyCowServer-AWS",  
  "region": "us-west-2"  
}
```

## Builders

Generate your image

Provider-specific

Multiple builders = cross-provider builds!

# Un Template Packer

```
{  
  "type": "shell",  
  "script": "makeMyCowAwesome.sh",  
  "only": ["MyCowServer-AWS"]  
}
```

Provisioners

Customize your image

Scripts or configuration management

Can be builder-specific

# Un Template Packer

```
{  
  "type": "compress",  
  "output": "MyCowServer-VirtualBox.tar.gz",  
  "only": ["MyCowServer-VirtualBox"]  
}
```

## Post-Processors

Put on the finishing touches

Integration with other services

Key to Docker conversion

# Provisioners

---

## Provisioners

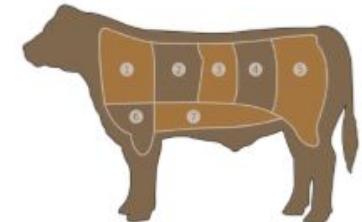
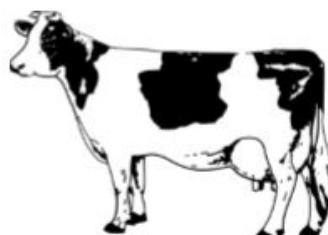
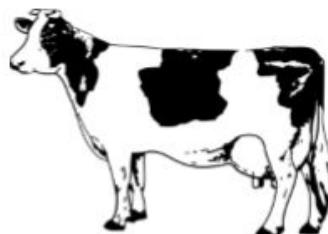
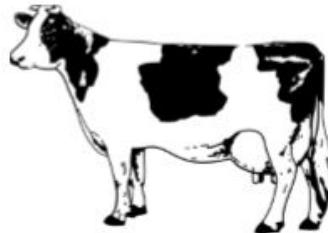
- Reuse your Existing Work
- Fun with Scripts!
- Move My Files!

## User Variables

## Bash and Ansible

---

# Provisioners



# Provisioners

```
{  
  "type": "shell",  
  "script": "makeMyCowAwesome.sh",  
  "only": ["MyCowServer-AWS"]  
}
```

Provisioner Types

Configuration Management

File

Script

# Provisioners de type Configuration Management



# Provisioners de type Script

Local Shell

Remote Shell

Powershell

Windows Shell

Windows Restart

# Variables

## User Variables

```
{  
  "variables": {  
    "aws_region": ""  
  },  
  "builders": [ {  
    "type": "amazon-ebs",  
    "region": "{{user `aws_region`}}"  
  }]  
}
```

# Variables

## User Variables - Environment Variables

```
{  
  "variables": {  
    "aws_region": "{{env `AWS_REGION`}}"  
  },  
  "builders": [{  
    "type": "amazon-ebs",  
    "region": "{{user `aws_region`}}"  
  }]  
}
```

# Built-In Functions

## Built-In Functions

`build_name`

`build_type`

`isotime [FORMAT]`

`lower`

`pwd`

`clean_ami_name`

`timestamp`

`uuid`

`...`

## clean-ami-name



Amazon Only

Remove unwanted characters

Apply with pipe

- `{} "10:02:36" | clean_ami_name {}`
- Output: 10-02-36

Commonly used with isotime

# Timestamp



Replace with Unix Timestamp  
Keep Image Names from Colliding

# Configuration des variables utilisateur

```
> packer build -var 'aws_region=us-west-1' template.json

> packer build -var-file=variables.json template.json

{

  "aws_region": "us-west-1"

}
```

# La transition de Packer vers Docker

---

## Packer vs Docker, Fight!

- Differences
- Why Switch? When?

## Packer and Docker – Playing Nice

- Transition to Docker
  - Hybrid Environments
-

# Packer – Qu'est ce qu'on obtient?



**Images for cloud and on-premise**

- Near-identical

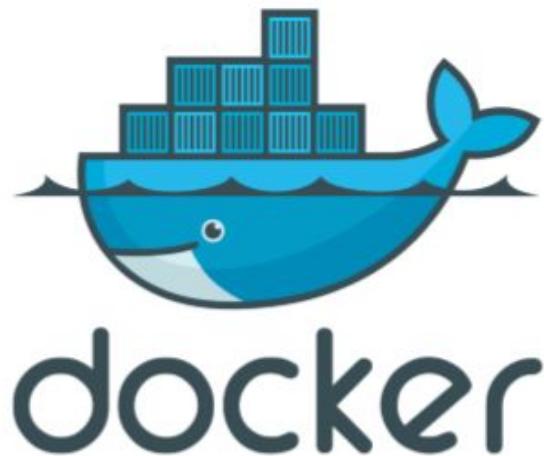
**Source controlled**

**Easy updates, easy rebuilds**

**Minimal code duplication**

- Integration with existing systems

# Docker - Qu'est ce qu'on obtient de plus?



## Binpacking

- Deployment flexibility
- Maximized server utilization

## Speed

## Built-In Versioning

## Containers – Identical Across Environments

# Les modules sont idempotents

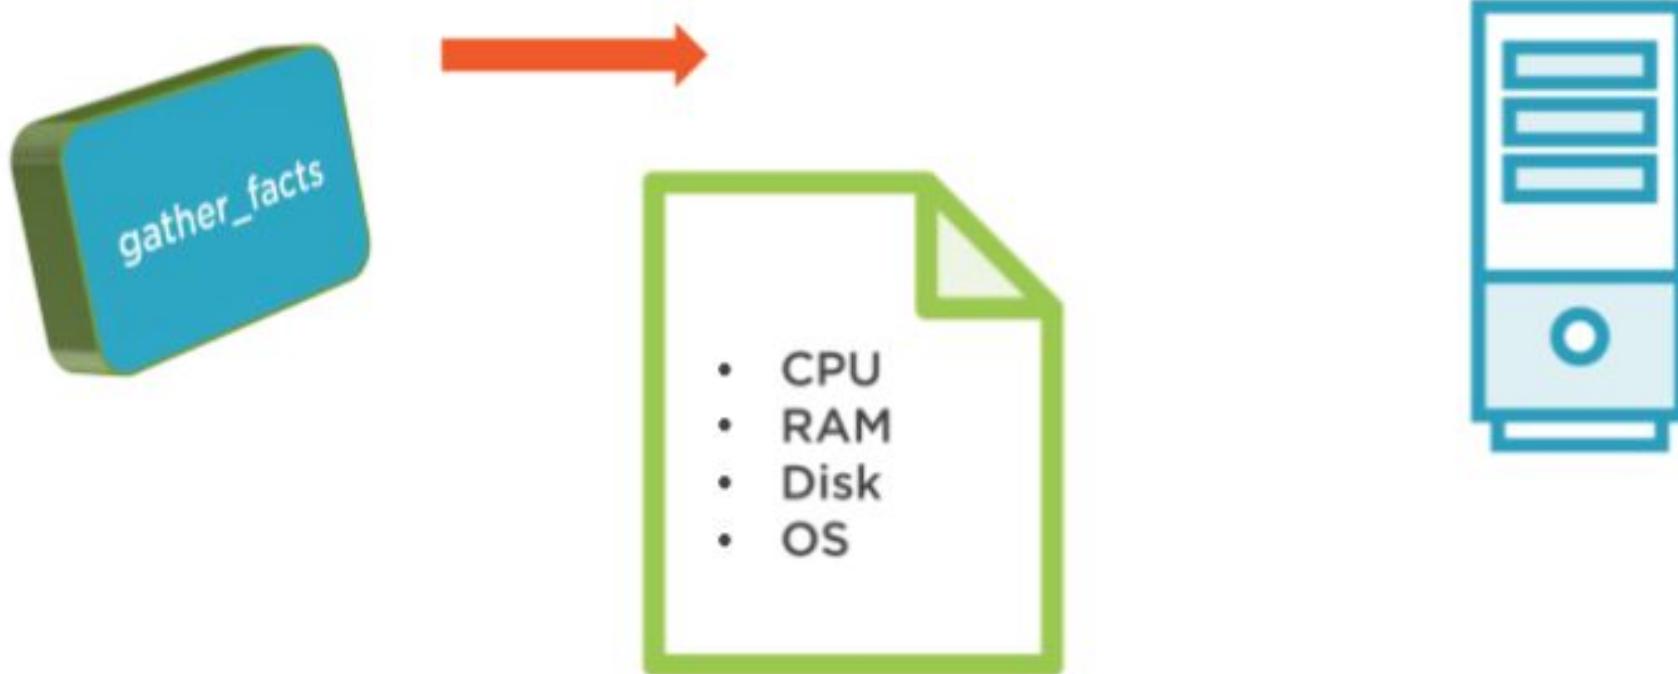


Running same Ansible commands...



“Desired State”

# Le module gather facts



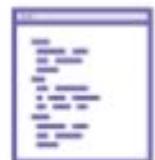
# Découvrons les Playbooks



Playbooks are groups of tasks that usually call modules



Playbooks are written in YAML



May contain variables, lists of hosts, groups, and tasks

## Les playbooks sont sauvegardés dans un SCM



## playbook.yaml

## Une play

```
---
- hosts: linux_servers
  vars:
    foo: "bar"
  tasks:
    - name: Copy a file
      copy:
        src: "/tmp/config.json"
        dest: "/etc/config.json"
```

This is one play

## Hosts

### playbook.yaml

```
---
- hosts: linux_servers ←
  vars:
    foo: "bar" ←
  tasks:
    - name: Copy a file
      copy:
        src: "/tmp/config.json"
        dest: "/etc/config.json"
```

# Variables

Scope of “version” variable

playbook.yaml

```
---
```

- hosts: linux\_servers
- vars:
  - foo: "bar" ← Variable
- tasks:
  - name: Copy a file
  - vars:
    - version: "beta" ← Variable
  - copy:
    - src: "/tmp/config.json"
    - dest: "/etc/config.json"

Scope of “foo” variable

## playbook.yaml

```
---
```

```
- hosts: linux_servers
```

```
  vars:
```

```
    foo: "bar"
```

```
  tasks:
```

```
#1 - name: Copy a file
```

```
    copy:
```

```
      src: "/tmp/config.json"
```

```
      dest: "/etc/config.json" ] }
```

```
#2 - name: Start Apache
```

```
    service:
```

```
      name: httpd
```

```
      state: started ] }
```

The code block shows a YAML playbook named "playbook.yaml". It defines a single host group "linux\_servers" with a variable "foo" set to "bar". It contains two tasks: task #1 uses the "copy" module to move a file from "/tmp/config.json" to "/etc/config.json", and task #2 uses the "service" module to start the "httpd" service in a started state. Annotations with arrows point to specific parts of the code: an orange arrow points to the "name" key in the first task, a blue arrow points to the "copy" module name, and a purple arrow points to the closing brace of the arguments for the "copy" task. Similar annotations are present for the second task.

## Tasks

## playbook.yaml

```
---
- hosts: linux_servers
  vars:
    foo: "bar"
  tasks:
    - name: Copy a file
      copy:           ← Module
        src: "/tmp/config.json"
        dest: "/etc/config.json" }
```

The code snippet shows a YAML playbook named "playbook.yaml". It defines a single host group "linux\_servers" with a variable "foo" set to "bar". A task is defined to copy a file from "/tmp/config.json" to "/etc/config.json". The "copy" command is highlighted with a red rounded rectangle and a blue arrow points to it from the text "Module". The "src" and "dest" arguments are highlighted with a purple brace and the text "Arguments".

# Tasks

## Ecrire des playbooks Ansible

- Playbooks are essentially sets of instructions (plays) that you send to run on a single target or groups of targets (hosts).
- For this example, the target host will be a RHEL/CentOS 7 base install.
- There will be a web server installed (NGINX) and then an index.html file will be created in the default webroot.
- After the install and file tasks are completed the service will be started.

# Ecrire des playbooks Ansible

Playbooks start with the YAML three dashes (---) followed by:

- **Name:** good for keeping the Playbooks readable
- **Hosts:** identifies the target for Ansible to run against
- **Become statements:** true statement is included here to ensure nginx installs without a problem (it's not always required)

```
1 | ---  
2 | - name: Install nginx  
3 |   hosts: host.name.ip  
4 |   become: true
```

# Ecrire des playbooks Ansible

## The state:

- Present statement lets Ansible check the state on the target first before performing any further action.
- In both cases if the repo or package is already present, Ansible knows it doesn't have to do any more for this task and continues.

```
1 tasks:
2   - name: Add epel-release repo
3     yum:
4       name: epel-release
5       state: present
6
7   - name: Install nginx
8     yum:
9       name: nginx
10      state: present
```

# Ecrire des playbooks Ansible

- Insert the index page in the default Nginx location

```
1 | - name: Insert Index Page
2 |   template:
3 |     src: index.html
4 |     dest: /usr/share/nginx/html/index.html
```

- The last thing the Playbook will do is make sure that the nginx service has been started (and if not, start it).

```
1 | - name: Start NGINX
2 |   service:
3 |     name: nginx
4 |     state: started
```

# Ecrire des playbooks Ansible

```
1  ---
2  - name: Install nginx
3    hosts: host.name.ip
4    become: true
5
6    tasks:
7      - name: Add epel-release repo
8        yum:
9          name: epel-release
10         state: present
11
12     - name: Install nginx
13       yum:
14         name: nginx
15         state: present
16
17     - name: Insert Index Page
18       template:
19         src: index.html
20         dest: /usr/share/nginx/html/index.html
21
22     - name: Start NGINX
23       service:
24         name: nginx
25         state: started
26
27
```

# Ecrire des playbooks Ansible



Roles help organize all the files and folders of your code



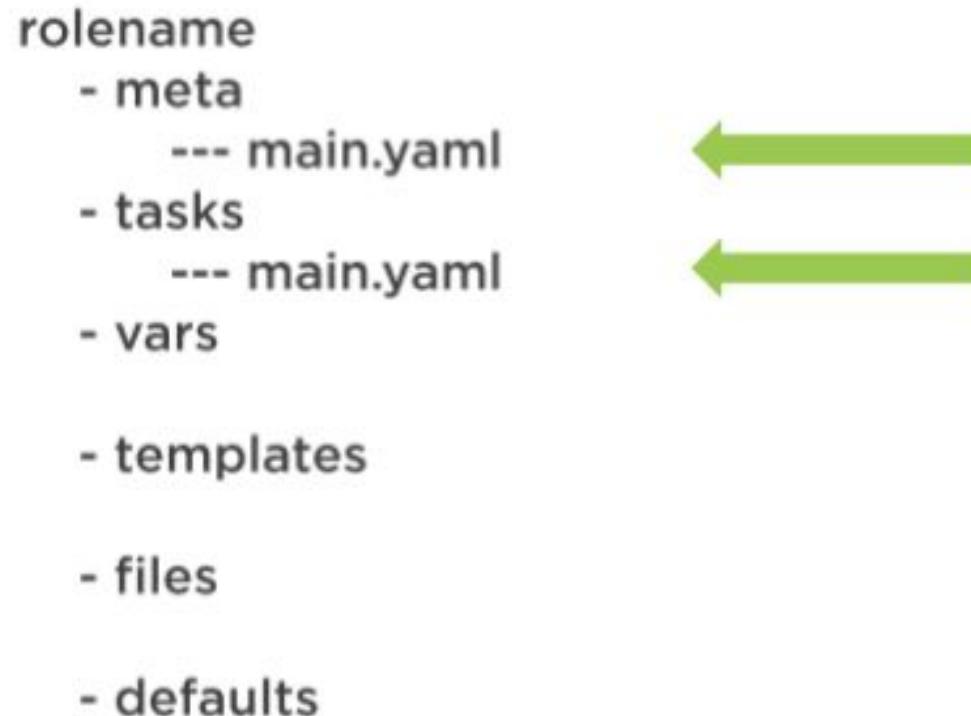
Pre-defined conventions of folders and files



No special project files, just folder conventions

# Structure d'un Role Ansible

```
rolename
  - meta
    --- main.yaml
  - tasks
    --- main.yaml
  - vars
  - templates
  - files
  - defaults
```



# Le dossier Meta



Descriptions of the role, dependencies, Ansible Galaxy information



Information about what roles this one depends on



Do we need to run a role before this one

# Le dossier Task



Primary entry for the actions to start happening



Same format as typical tasks in a playbook



Can have multiple task files included

# D'autres dossiers Ansible Role

rolename

- meta

- tasks

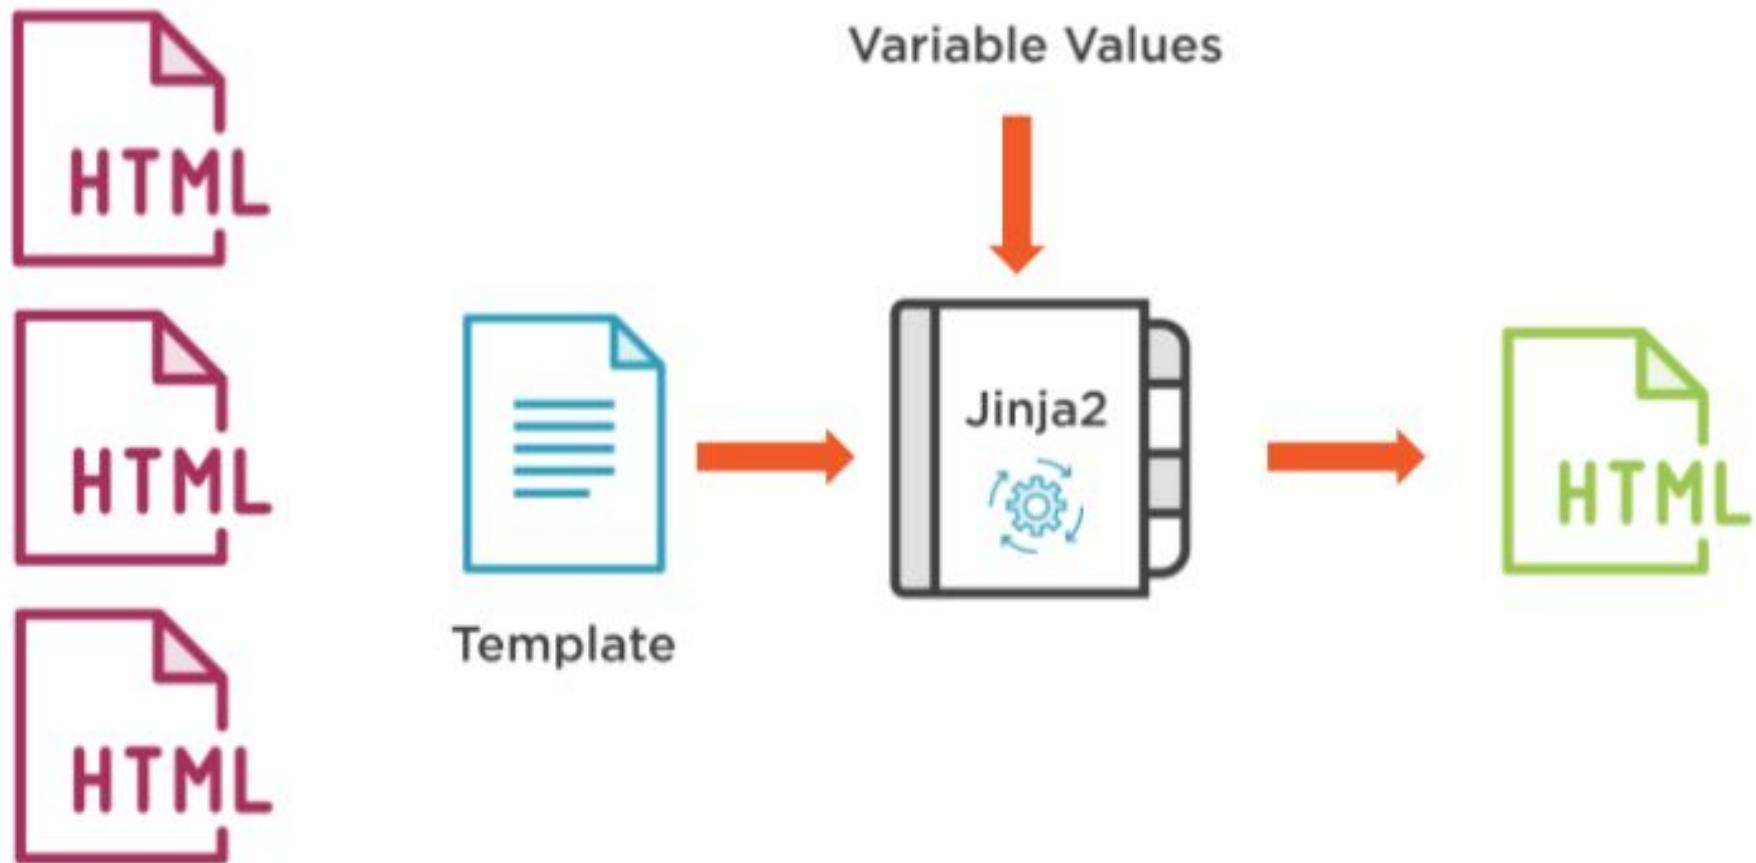
- vars

- templates  **Templated configuration recipes**

- files  **Hold files that need to be copied to hosts**

- defaults  **Good for variable defaults with overrides**

# Templates Ansible avec Jinja2



# Terraform



- Terraform est un outil permettant de créer, modifier et versionner une infrastructure de manière sûre et efficace.
- Terraform peut gérer différents fournisseurs d'infrastructure, allant des Cloud Providers (AWS, Azure, GcP, Alibaba Cloud, etc ...)
- Ainsi que les solutions internes personnalisées (VmWare, Kubernetes, etc ...).

# Terraform

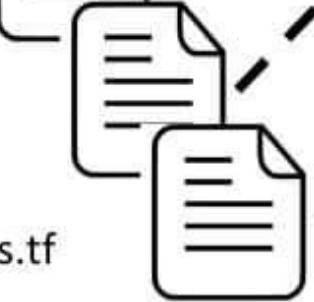
terraform.state



terraform.tfvars



terraform-provider.tf



terraform-instances.tf



HashiCorp  
**Terraform**



# Démonstration avec

# Terraform et Aws

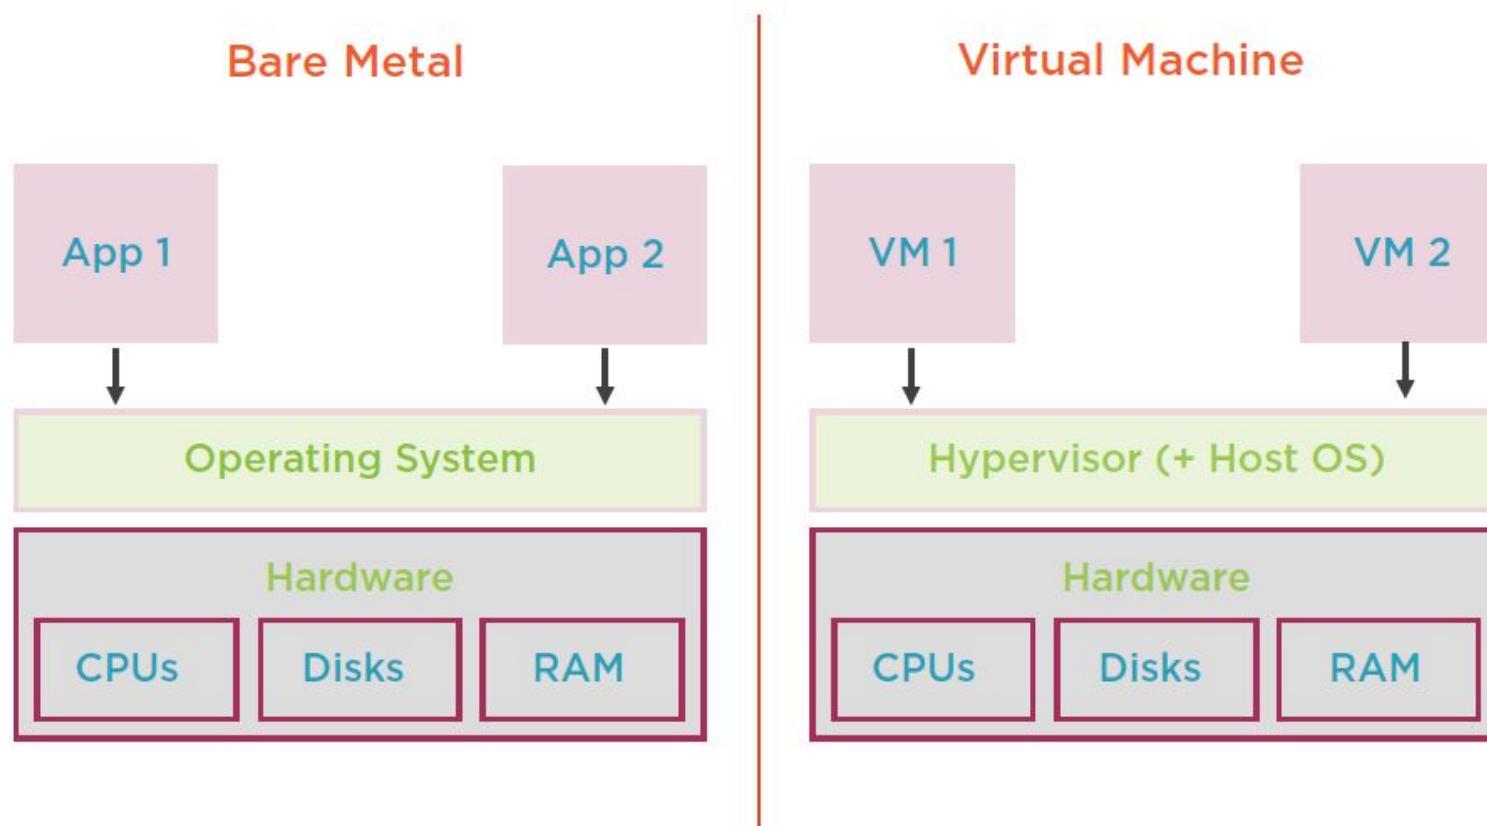
---

# Containers avec Docker

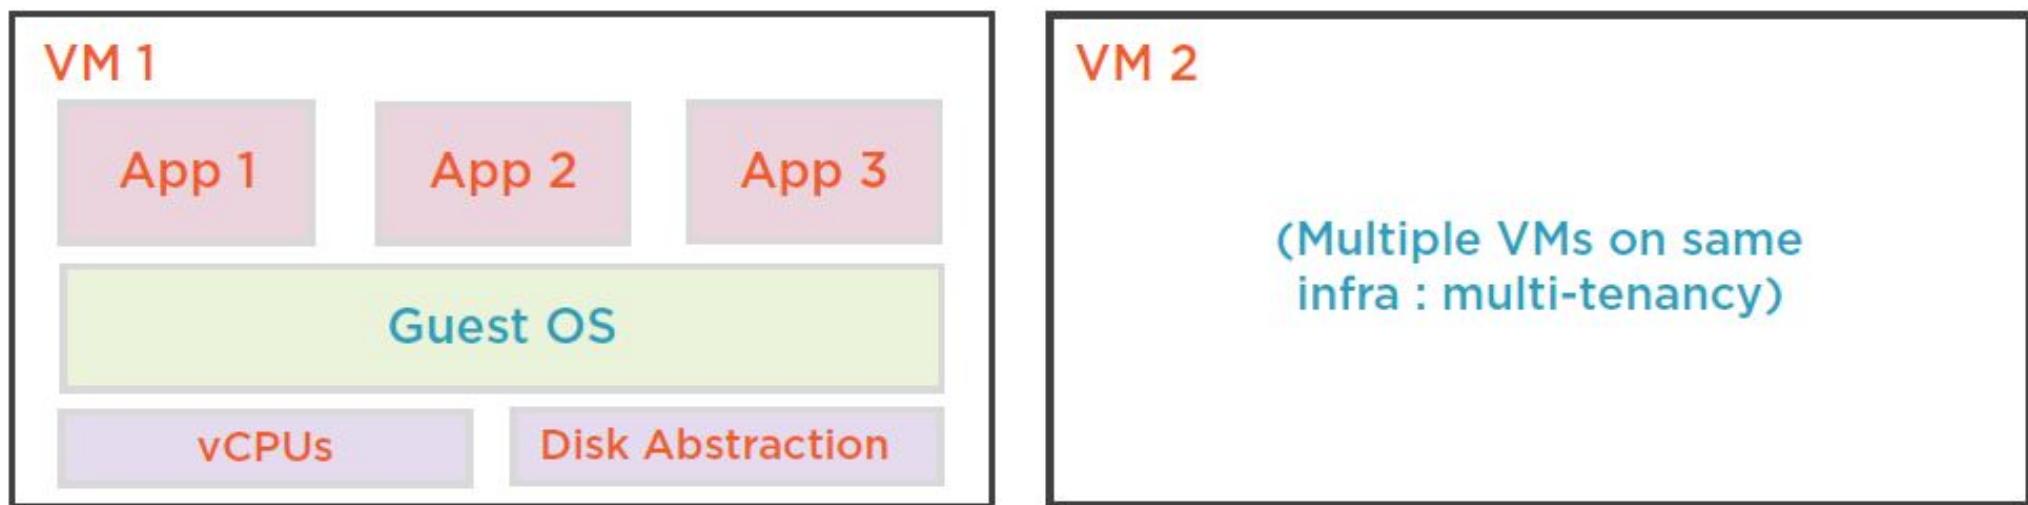
# Agenda

- 1. Concepts de base : immutabilité, image, layers, registry, problématique réseau et stockage.**
- 2. Automatisation avec Dockerfile/docker-compose, intégration avec Github, Jenkins, DockerHub.**
- 3. Bénéfices attendus : reproductibilité, manageabilité.**
- 4. Apports en termes d'élasticité, Agilité, évolutivité.**
- 5. Impacts sur les équipes de développement et d'infrastructure.**

## Bare Metal and Virtual Machines



## Apps on Virtual Machines



Hypervisor (+ Host OS)

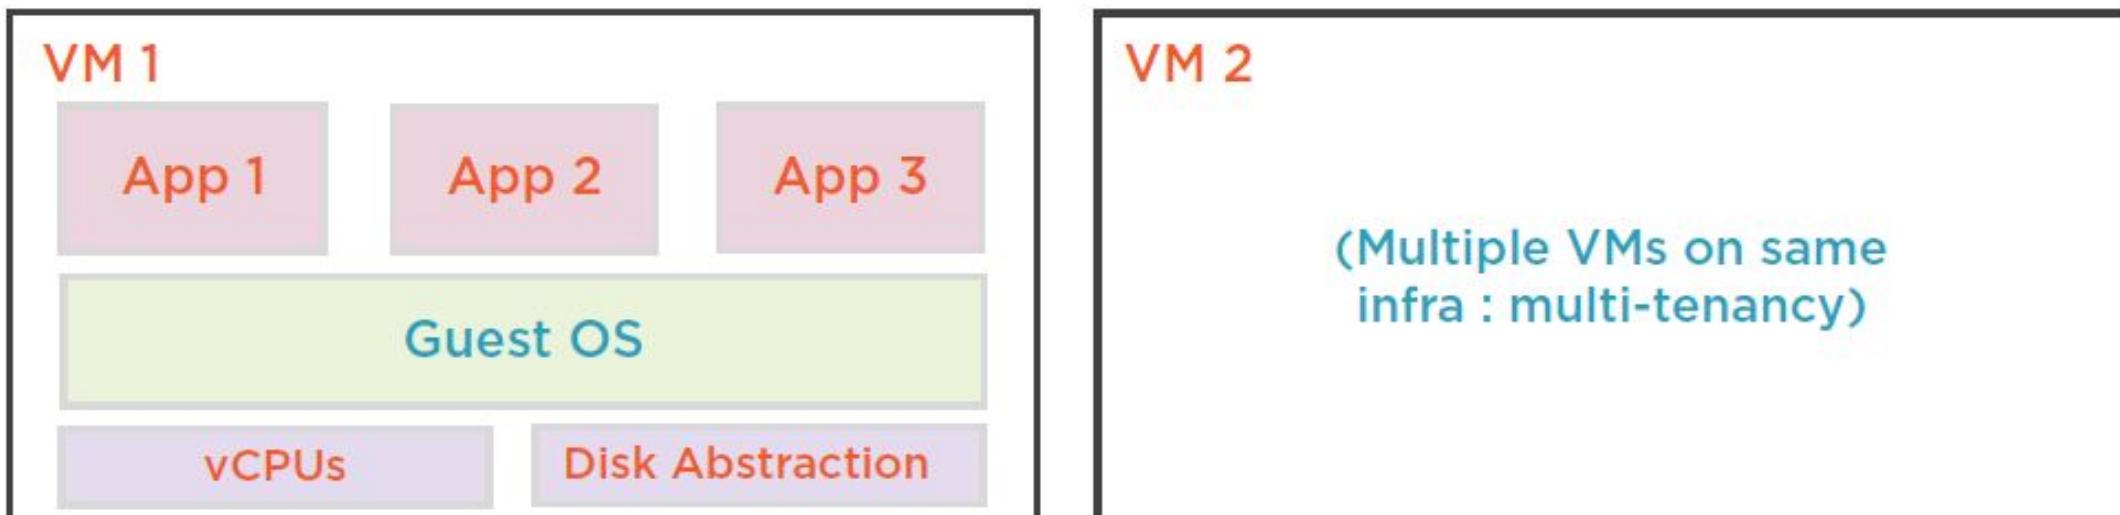
Hardware

CPUs

Disks

RAM

## Apps on Virtual Machines



Hypervisor (+ Host OS)

Hardware

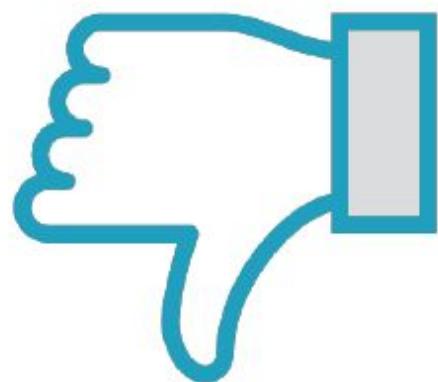
CPUs

Disks

RAM

VMs are far more lightweight than bare metal environments and are often used to run applications

## Drawbacks of VMs



### Contain guest OS

- Introduces platform dependency
- Bloats image size to GB (apps far smaller)

### Heavyweight

- Slow to boot up

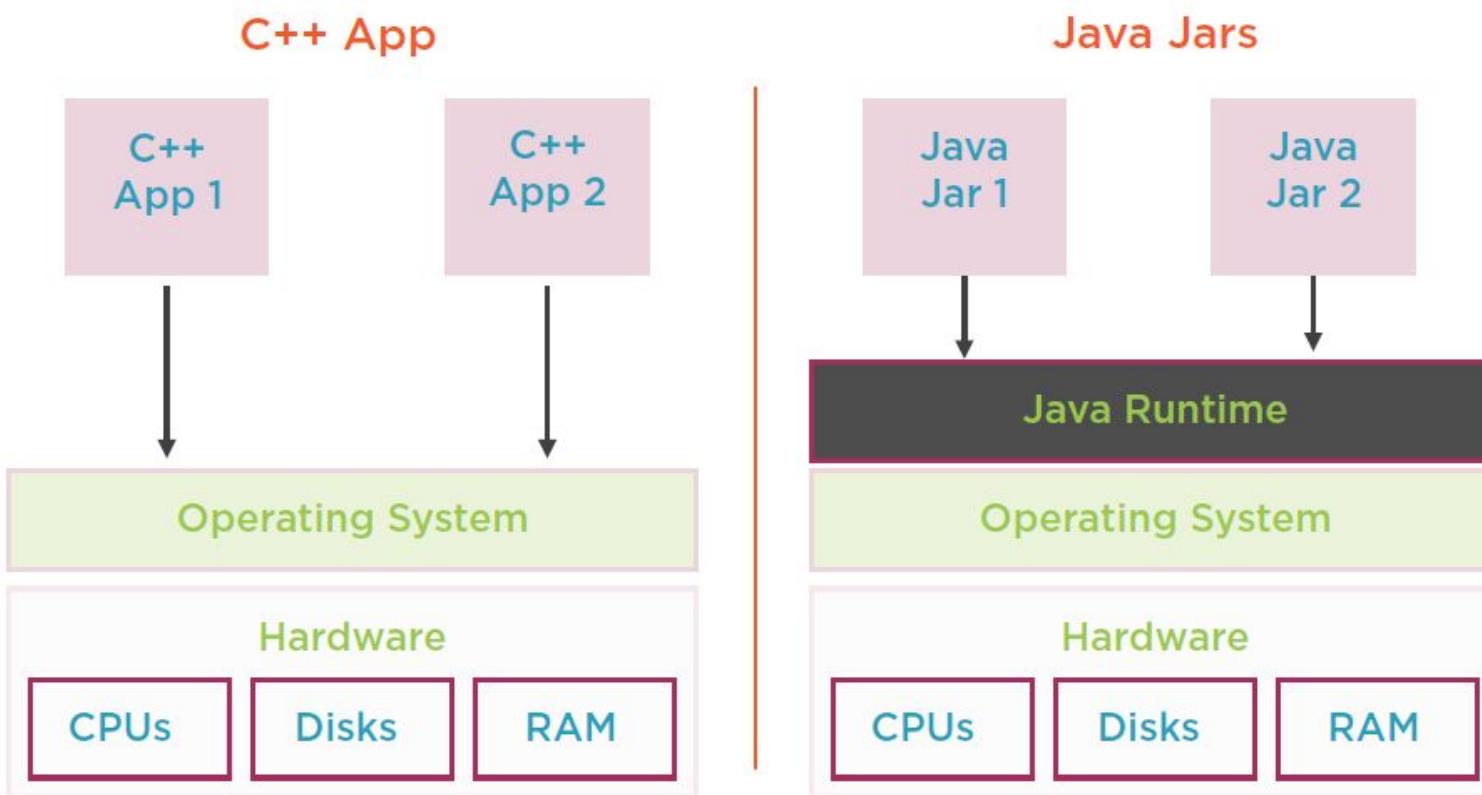
### Not trivial to migrate

- VM migration tools needed

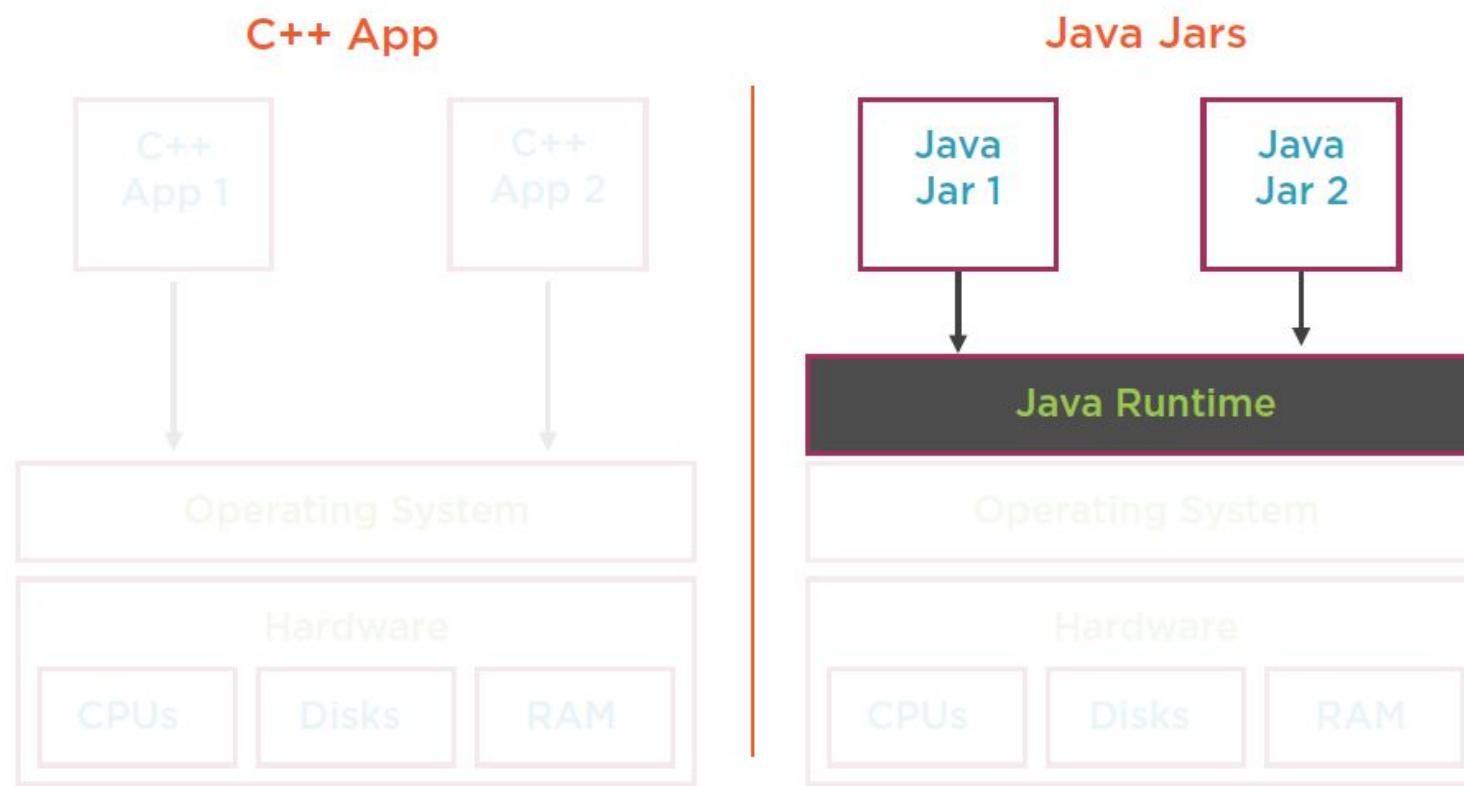
# Container

A container image is a lightweight, stand-alone, executable package of a piece of software that includes everything needed to run it: code, runtime, system tools, system libraries, settings

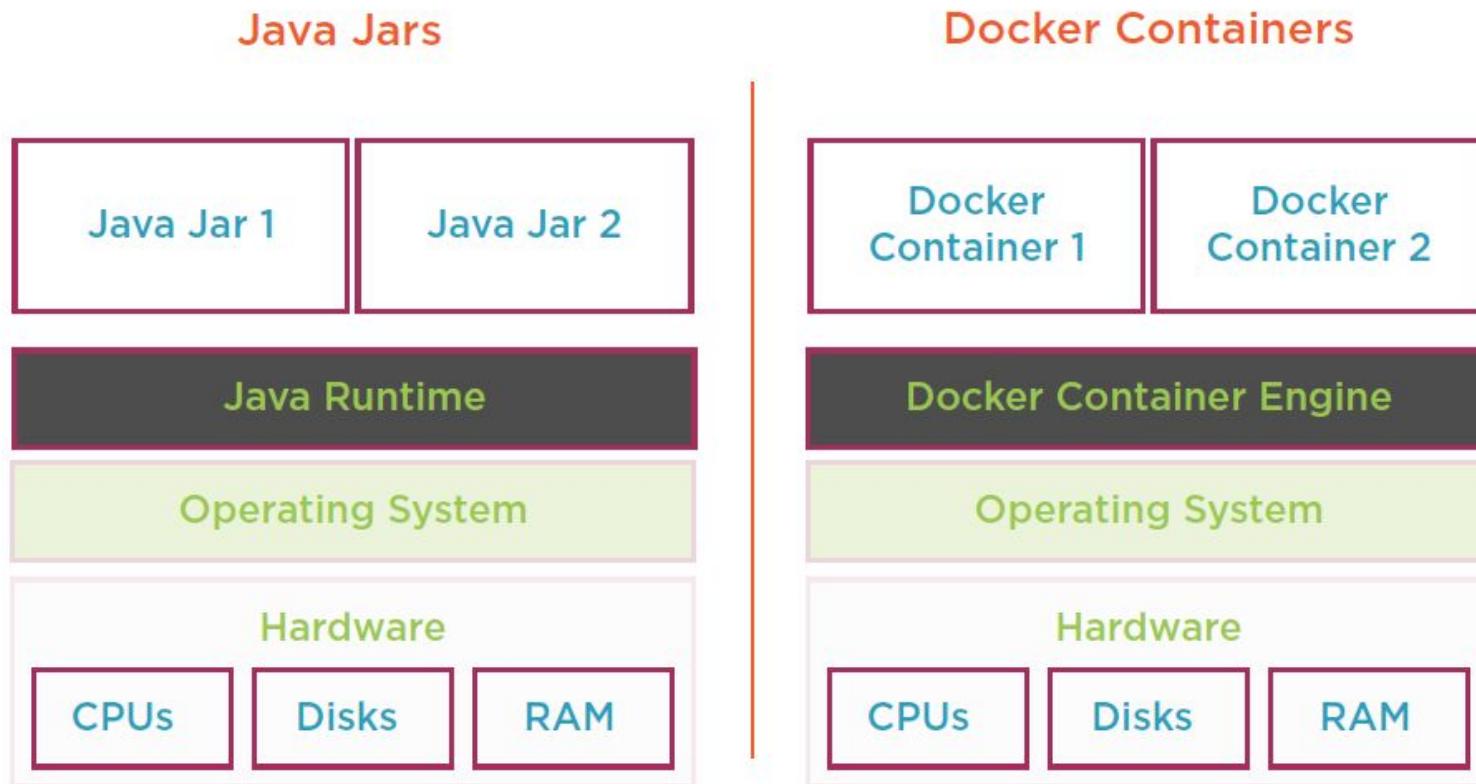
## Containers Are ‘Like’ Jars



## Containers Are ‘Like’ Jars



## Containers Are ‘Like’ Jars



## Java Jars and Containers

### Java Jar Files

- Files containing apps
- Platform independent
- Run on layer of abstraction
- Java Runtime

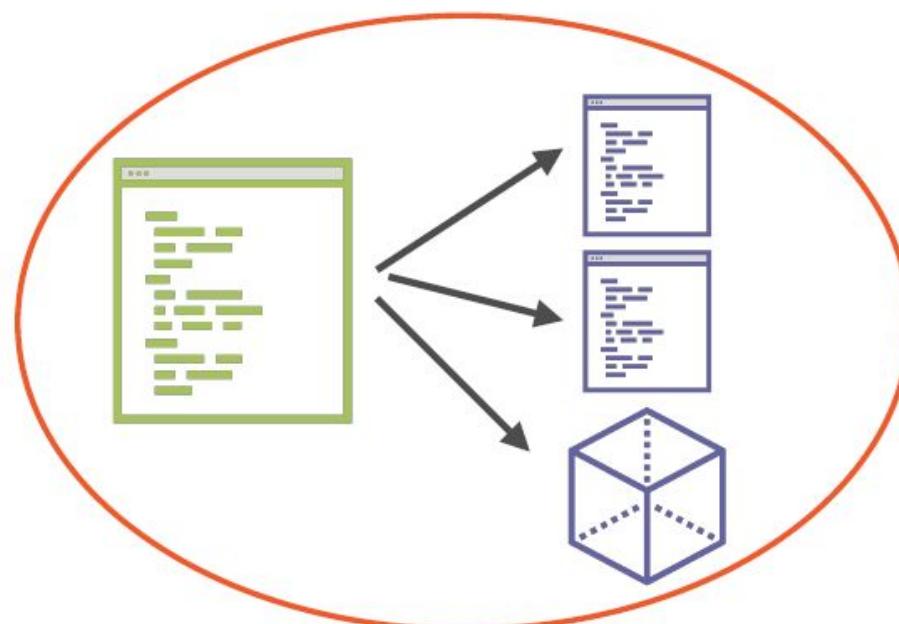
### Containers

- Files containing apps
- Platform independent
- Run on layer of abstraction
- Docker Runtime

# Container

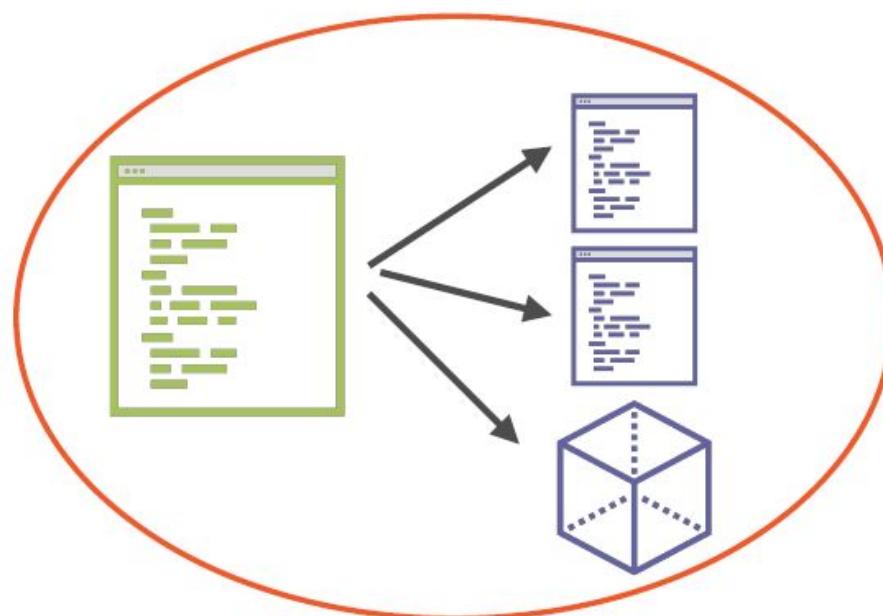
A container image is a lightweight, stand-alone, executable package of a piece of software that includes everything needed to run it: code, runtime, system tools, system libraries, settings

## Docker Containers



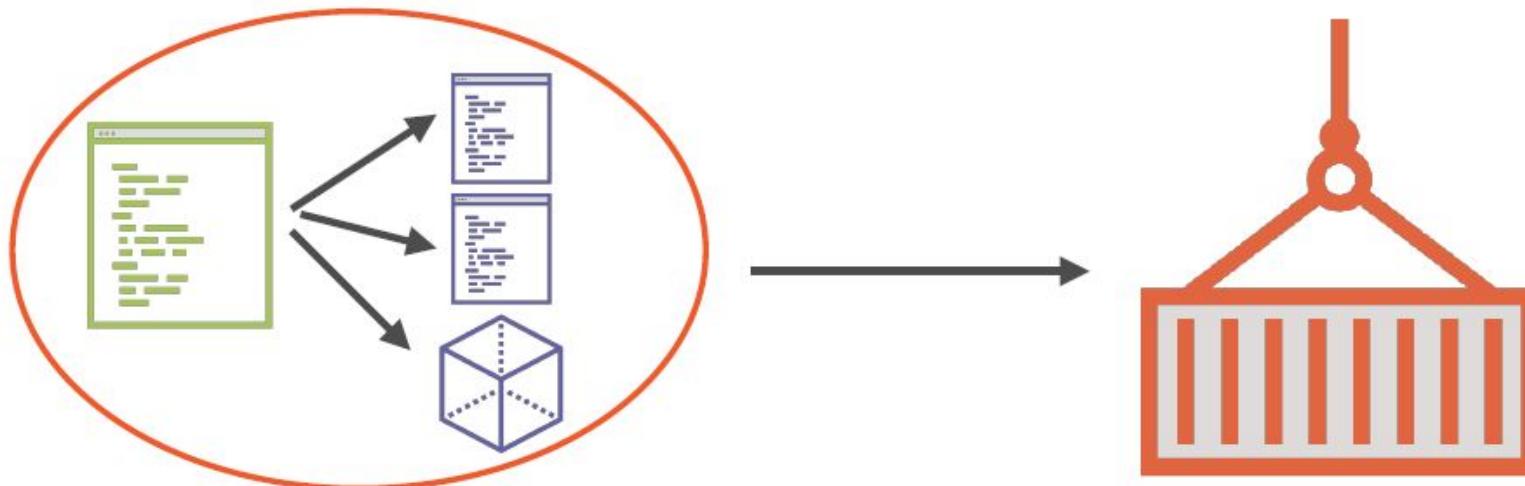
A container packages code and its  
dependencies into an image

## Docker Containers



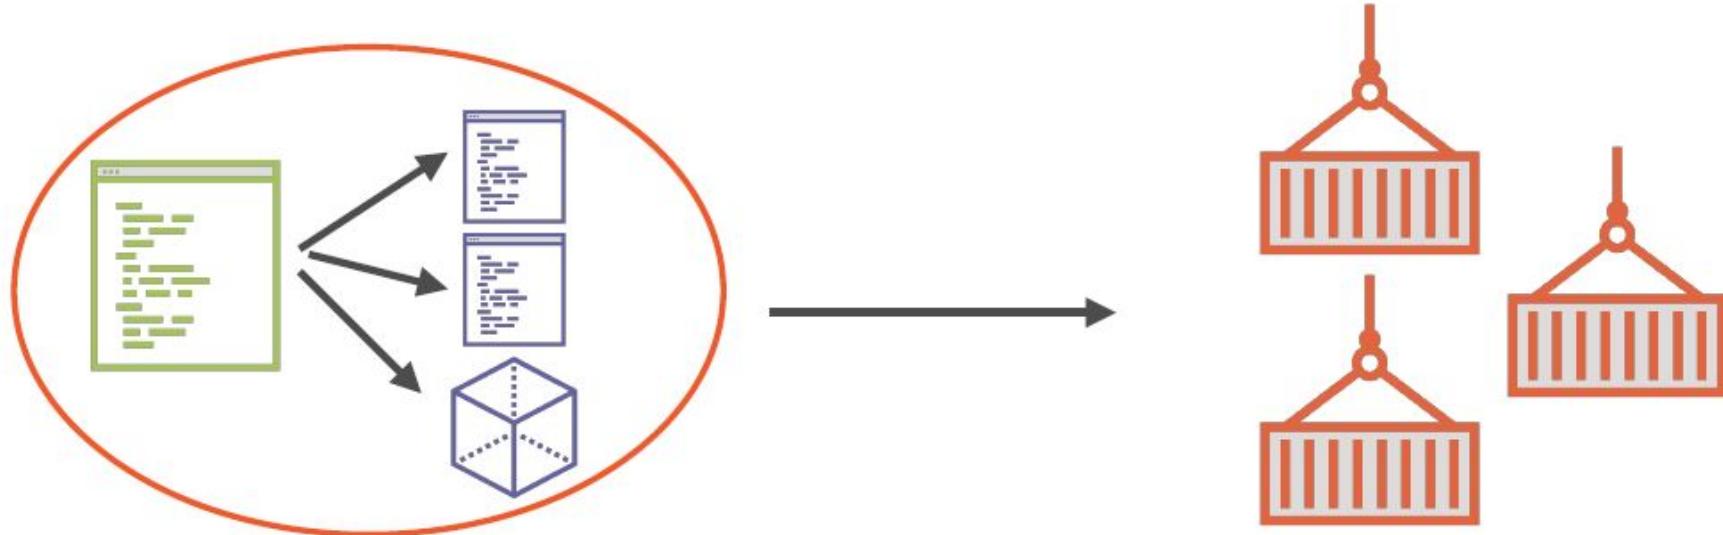
A fully self-contained environment

## Docker Containers



Code is abstracted away from the machine

## Docker Containers



Create as many containers as you want  
from the same image

# Démonstration:

# intégration avec Github

---

Lab:

Qu'est-ce qu'un conteneur?

---

# Qu'est-ce qu'un conteneur?



## Processus

Les conteneurs ne sont que des processus Linux normaux avec une configuration supplémentaire appliquée. Lancez le conteneur Redis suivant afin que nous puissions voir ce qui se passe sous les couvertures.

```
> docker run -d --name=db redis:alpine
```

```
$ docker run -d --name = db redis: alpin
Impossible de trouver l'image 'redis: alpine' localement
alpin: Extraire de la bibliothèque / redis
801bfaa63ef2: Tirage complet
9a8d0188e481: Tirage complet
8a3f5c4e0176: Tirage complet
3f7cb00af226: Tirage complet
e421f2f8acb5: Tirage terminé
f41cc3c7c3e4: Pull terminé
Digest: sha256: 2cd821f730b90a197816252972c2472e3d1fad3c42f052580bc958d3ad641f96
Statut: Image plus récente téléchargée pour redis: alpin
3a1def9d32b79cf98868f85465c38acc52cb5be1b8a5e731e0f442464333bc97
```

Le conteneur Docker lance un processus appelé redis-server. Depuis l'hôte, nous pouvons voir tous les processus en cours d'exécution, y compris ceux démarrés par Docker.

```
> ps aux | grep redis-server
```

```
$ ps aux | grep redis-serveur
999 1769 0,2 1,2 28968 12892? Ssl 21:56 0:00 serveur redis
racine 1996 0,0 0,0 14220 1008 pts / 0 S + 21:58 0:00 grep --couleur = auto serveur redis
```

# Qu'est-ce qu'un conteneur?

- Docker peut nous aider à identifier les informations sur le processus, y compris le PID (ID de processus) et PPID (ID de processus parent)  
`Docker top db`
- Qu'est ce que le PPID? Utilisez `ps aux | grep <ppid>` pour trouver le processus parent. Probablement Containerd.
- La commande `pstree` listera tous les sous-processus. Voir l'arborescence des processus Docker à l'aide de `pstree -c -p -A $(pgrep dockerd)`

```
root@user-HP-ProBook-440-G5:/home/user# pstree -c -p -A $(pgrep dockerd)
dockerd(17504)-+-{dockerd}(17505)
              |-{dockerd}(17506)
              |-{dockerd}(17507)
              |-{dockerd}(17508)
              |-{dockerd}(17509)
              |-{dockerd}(17510)
              |-{dockerd}(17511)
              |-{dockerd}(17512)
              |-{dockerd}(17513)
              |-{dockerd}(17514)
              |-{dockerd}(17515)
              |-{dockerd}(17516)
              |-{dockerd}(17517)
              |-{dockerd}(17518)
              |-{dockerd}(17519)
              |-{dockerd}(17639)
```

# Répertoire de processus

- Linux n'est qu'une série de fichiers et de contenus magiques, ce qui rend amusant l'exploration et la navigation pour voir ce qui se passe sous les couvertures, et dans certains cas, modifiez le contenu pour voir les résultats.
- La configuration de chaque processus est définie dans le /proc répertoire. Si vous connaissez l'ID de processus, vous pouvez identifier le répertoire de configuration.
- La commande ci-dessous répertorie tout le contenu de /proc et stocke le PID Redis pour une utilisation future  
`DBPID=$(pgrep redis-server)` `echo Redis is $DBPID` `ls /proc`

```
root@user-HP-ProBook-440-G5:/home/user# DBPID=$(pgrep redis-server)
root@user-HP-ProBook-440-G5:/home/user# echo Redis is $DBPID
Redis is 23692
```

- Chaque processus a sa propre configuration et ses propres paramètres de sécurité définis dans différents fichiers.
- Par exemple, vous pouvez afficher et mettre à jour les variables d'environnement définies pour ce processus. `cat /proc/$DBPID/environ`

```
> docker exec -it db env
```

Lab :

Lancement de conteneurs à l'aide  
de Podman et Libpod

---

## Installer podman sur Ubuntu 18

- Installer les packages podman:

```
> sudo apt-get install software-properties-common -y  
> sudo add-apt-repository -y ppa:projectatomic/ppa  
> sudo apt-get install podman -y  
> podman info
```

- Travailler avec Podman:

Avant de commencer, vous devrez définir les registres docker.io dans le fichier **/etc/containers/registries.conf**.

```
> echo -e "[registries.search]\nregistries = ['docker.io']" | sudo tee /etc/containers/registries.conf
```

Ensuite, téléchargez et exécutez l'image hello-world avec la commande suivante :

```
> podman run hello-world
```

# Installer podman sur Ubuntu 18

```
root@user-HP-ProBook-440-G5:/etc/containers# podman run hello-world
Trying to pull docker.io/library/hello-world...
Getting image source signatures
Copying blob 0e03bdcc26d7 done
Copying config bf756fb1ae done
Writing manifest to image destination
Storing signatures

Hello from Docker!
This message shows that your installation appears to be working correctly.

To generate this message, Docker took the following steps:
 1. The Docker client contacted the Docker daemon.
 2. The Docker daemon pulled the "hello-world" image from the Docker Hub.
    (amd64)
 3. The Docker daemon created a new container from that image which runs the
    executable that produces the output you are currently reading.
 4. The Docker daemon streamed that output to the Docker client, which sent it
    to your terminal.

To try something more ambitious, you can run an Ubuntu container with:
$ docker run -it ubuntu bash

Share images, automate workflows, and more with a free Docker ID:
https://hub.docker.com/

For more examples and ideas, visit:
https://docs.docker.com/get-started/
```

# Lancement de conteneurs à l'aide de Podman et Libpod



## Exercice:

libpod fournit une bibliothèque pour les applications qui cherchent à utiliser le concept Container Pod popularisé par Kubernetes. libpod contient également un outil appelé podman pour gérer les pods , les conteneurs et les images de conteneurs.

```
> podman --help
```

L'interface de ligne de commande pour Podman est compatible avec Docker, ce qui signifie qu'elle devrait être familière. La commande ci-dessous lancera un nouveau conteneur.

```
> podman login docker.io
```

```
root@user-HP-ProBook-440-G5:/home/user# podman login docker.io
Username: hosniah
Password:
Login Succeeded!
```

```
> podman run -d --name http-noports hosniah/docker-http-server:1.0
```

```
root@user-HP-ProBook-440-G5:/home/user# podman run -d --name http-noports hosniah/docker-http-server:1.0
Trying to pull docker.io/hosniah/docker-http-server:1.0...
Getting image source signatures
Copying blob f139eb4721ae skipped: already exists
Copying config 0d6ee549ae done
Writing manifest to image destination
Storing signatures
56c92b0f01de9b68c0bcc5419538a1a0939018712abf9bdc5498c5fd4967c0e6
```

# Lancement de conteneurs à l'aide de Podman et Libpod



> *podman ps*

root@user-HP-ProBook-440-G5:/home/user# podman ps						
CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS	NAMES
56c92b0f01de	docker.io/hosniah/docker-http-server:1.0	/app	5 minutes ago	Up 5 minutes ago		http-noports

Étant donné que ces conteneurs ne s'exécutent pas via Docker, ils ne seront pas visibles via l'interface de ligne de commande Docker.

> *docker ps*

Les conteneurs Libpod ont les mêmes principes en termes de configuration et de sécurité. Si nous voulions que les ports soient accessibles, nous devons les lier explicitement.

> *podman run -d --name http -p 80:80 hosniah/docker-http-server:1.0*

> *curl localhost*

## CLI compatible Docker

---

La configuration d'un conteneur peut être sortie via inspect. La sortie est compatible avec l'API Docker.

```
> podman inspect http
```

Comme nous exécutons un service compatible, il est possible de définir simplement un alias pour remplacer la CLI Docker et avoir la même expérience.

```
> alias docker=podman
```

Désormais, lorsque nous exécuterons Docker ps, il sera basé sur les conteneurs exécutés par Podman.

```
> docker ps
```

---

# CLI compatible Docker



```
root@user-HP-ProBook-440-G5:/home/user# podman inspect http
[
  {
    "Id": "28866519f23bc5185d78c0a3bf36580ab6fcfe63a7595778da2f4e24b25b1790",
    "Created": "2021-01-24T20:52:05.513358189+01:00",
    "Path": "/app",
    "Args": [
      "/app"
    ],
    "State": {
      "OciVersion": "1.0.1-dev",
      "Status": "configured",
      "Running": false,
      "Paused": false,
      "Restarting": false,
      "OOMKilled": false,
      "Dead": false,
      "Pid": 0,
      "ExitCode": 0,
      "Error": "",
      "StartedAt": "0001-01-01T00:00:00Z",
      "FinishedAt": "0001-01-01T00:00:00Z",
      "Healthcheck": {
        "Status": "",
        "FailingStreak": 0,
        "Log": null
      }
    },
    "Image": "0d6ee549ae1314ff0c0b8fea18042f3891250ba6b920fc714563be624e3c62a3",
    "ImageName": "docker.io/hosniah/docker-http-server:1.0",
    "Rootfs": ""
  }
]
```

# Sécurité des conteneurs Podman



---

Les conteneurs peuvent également être exécutés au premier plan, par exemple en lançant un shell basé sur Alpine.

```
> podman run -it alpine sh
```

Une fois que vous avez un shell, la sécurité du conteneur peut être testée en utilisant "Am I Contained" par Jessfraz.

```
> wget -O amicontained  
https://github.com/jessfraz/amicontained/releases/download/v0.3.0/amicontained-linux-amd64; chmod +x amicontained; ./amicontaine
```

---

# Sécurité des conteneurs Podman



```
root@user-HP-ProBook-440-G5:/home/user# podman run -it alpine sh
Trying to pull docker.io/library/alpine...
Getting image source signatures
Copying blob 596ba82af5aa done
Copying config 7731472c3f done
Writing manifest to image destination
Storing signatures
/ # wget -O amicontained https://github.com/jessfraz/amicontained/releases/download/v0.3.0/amicontained-linux-amd64; chmod +x amicontained; ./amicontained
Connecting to github.com (140.82.121.4:443)
Connecting to github.com (140.82.121.3:443)
Connecting to github-production-release-asset-2e65be.s3.amazonaws.com (52.217.85.92:443)
saving to 'amicontained'
amicontained          100% |*****| 1764k  0:00:00 ETA
'amicontained' saved
Container Runtime: not-found
Has Namespaces:
  pid: true
  user: false
AppArmor Profile: unconfined
Capabilities:
  BOUNDING -> chown dac_override fowner fsetid kill setgid setuid setpcap net_bind_service net_raw sys_chroot mknod audit_write setfcap
  AMBIENT -> chown dac_override fowner fsetid kill setgid setuid setpcap net_bind_service net_raw sys_chroot mknod audit_write setfcap
Chroot (not pivot_root): false
Seccomp: filtering
```

# Kubernetes orchestrateur de containers

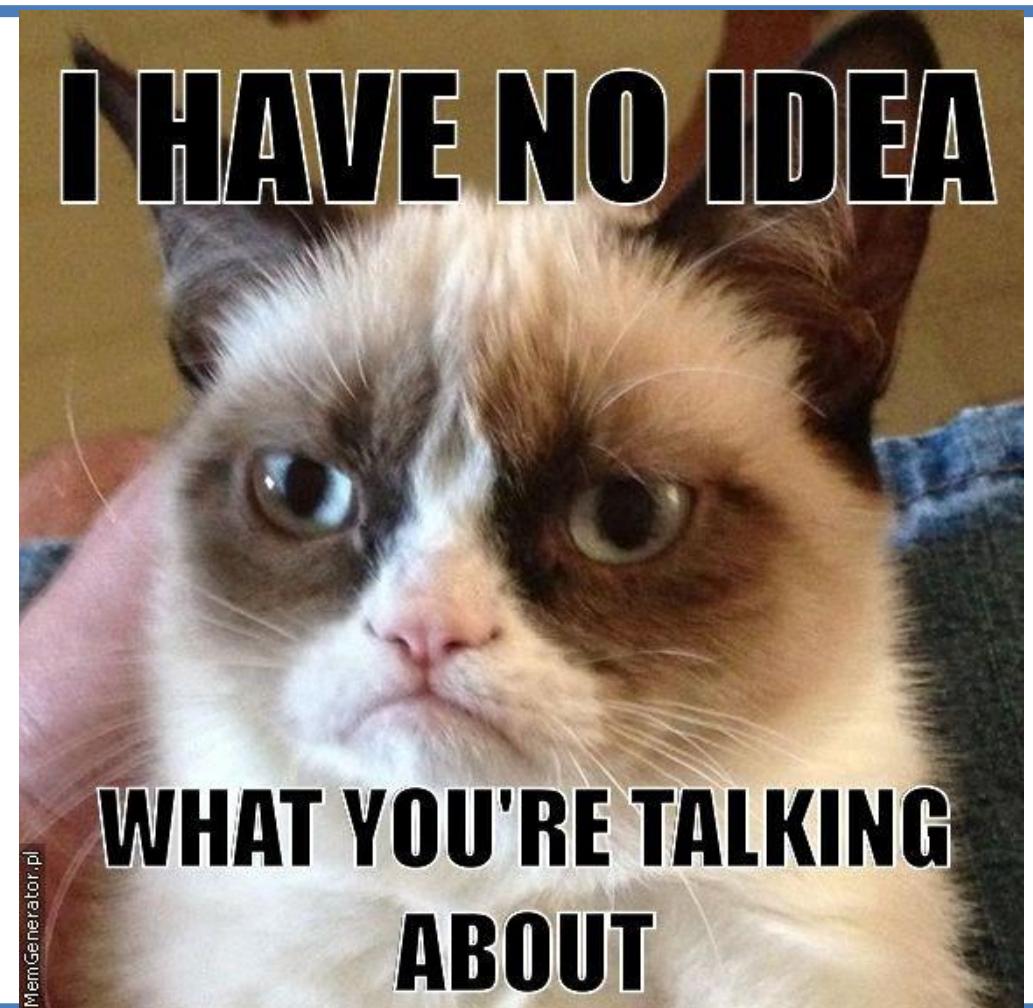
# Agenda

1. Nœuds Master/Workers, concepts de Pods, service, différents types d'Ingress Controller.
2. Stockage : stateful, stateless, shared (NFS, GlusterFS, CEPH, rook).
3. Gestion de configuration. Usage des Jobs et DaemonSets.
4. Composant interne (etcd, kubelet, kube-dns, kube-proxy, apiserver), complémentaire (Helm/Tiller, envoy, side-car proxy).
5. Service Discovery/Mesh (Istio), calico, cilium.

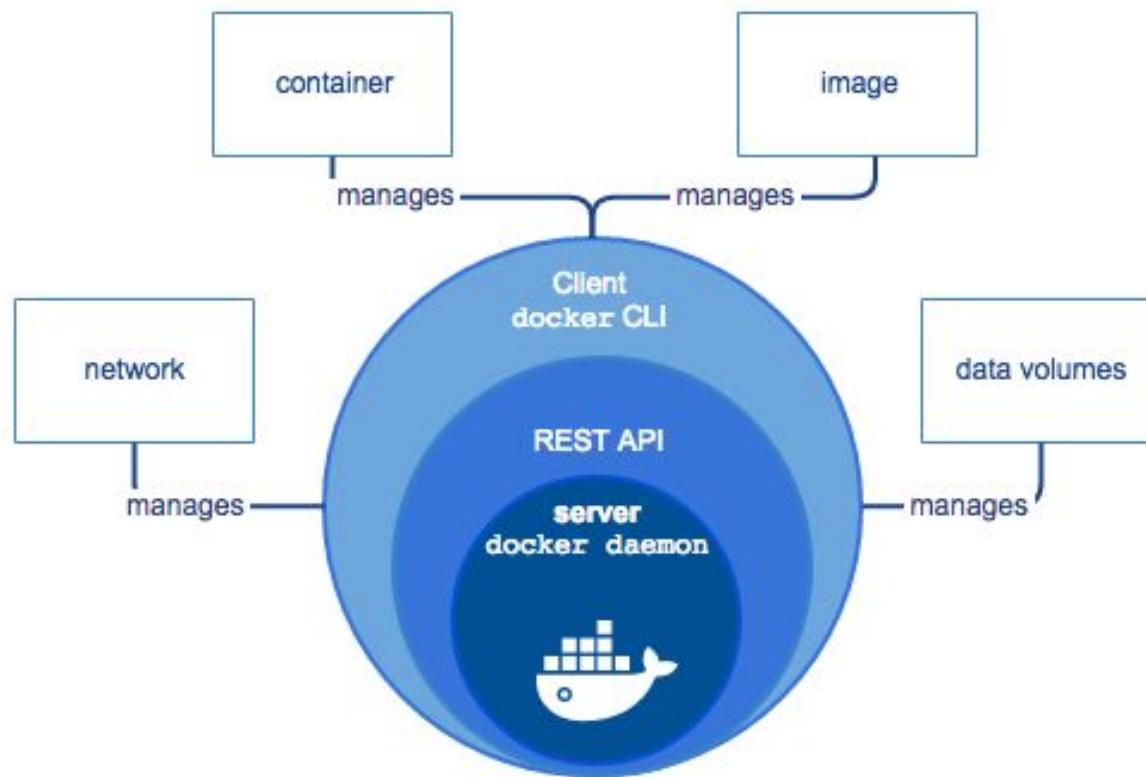


# Containers?

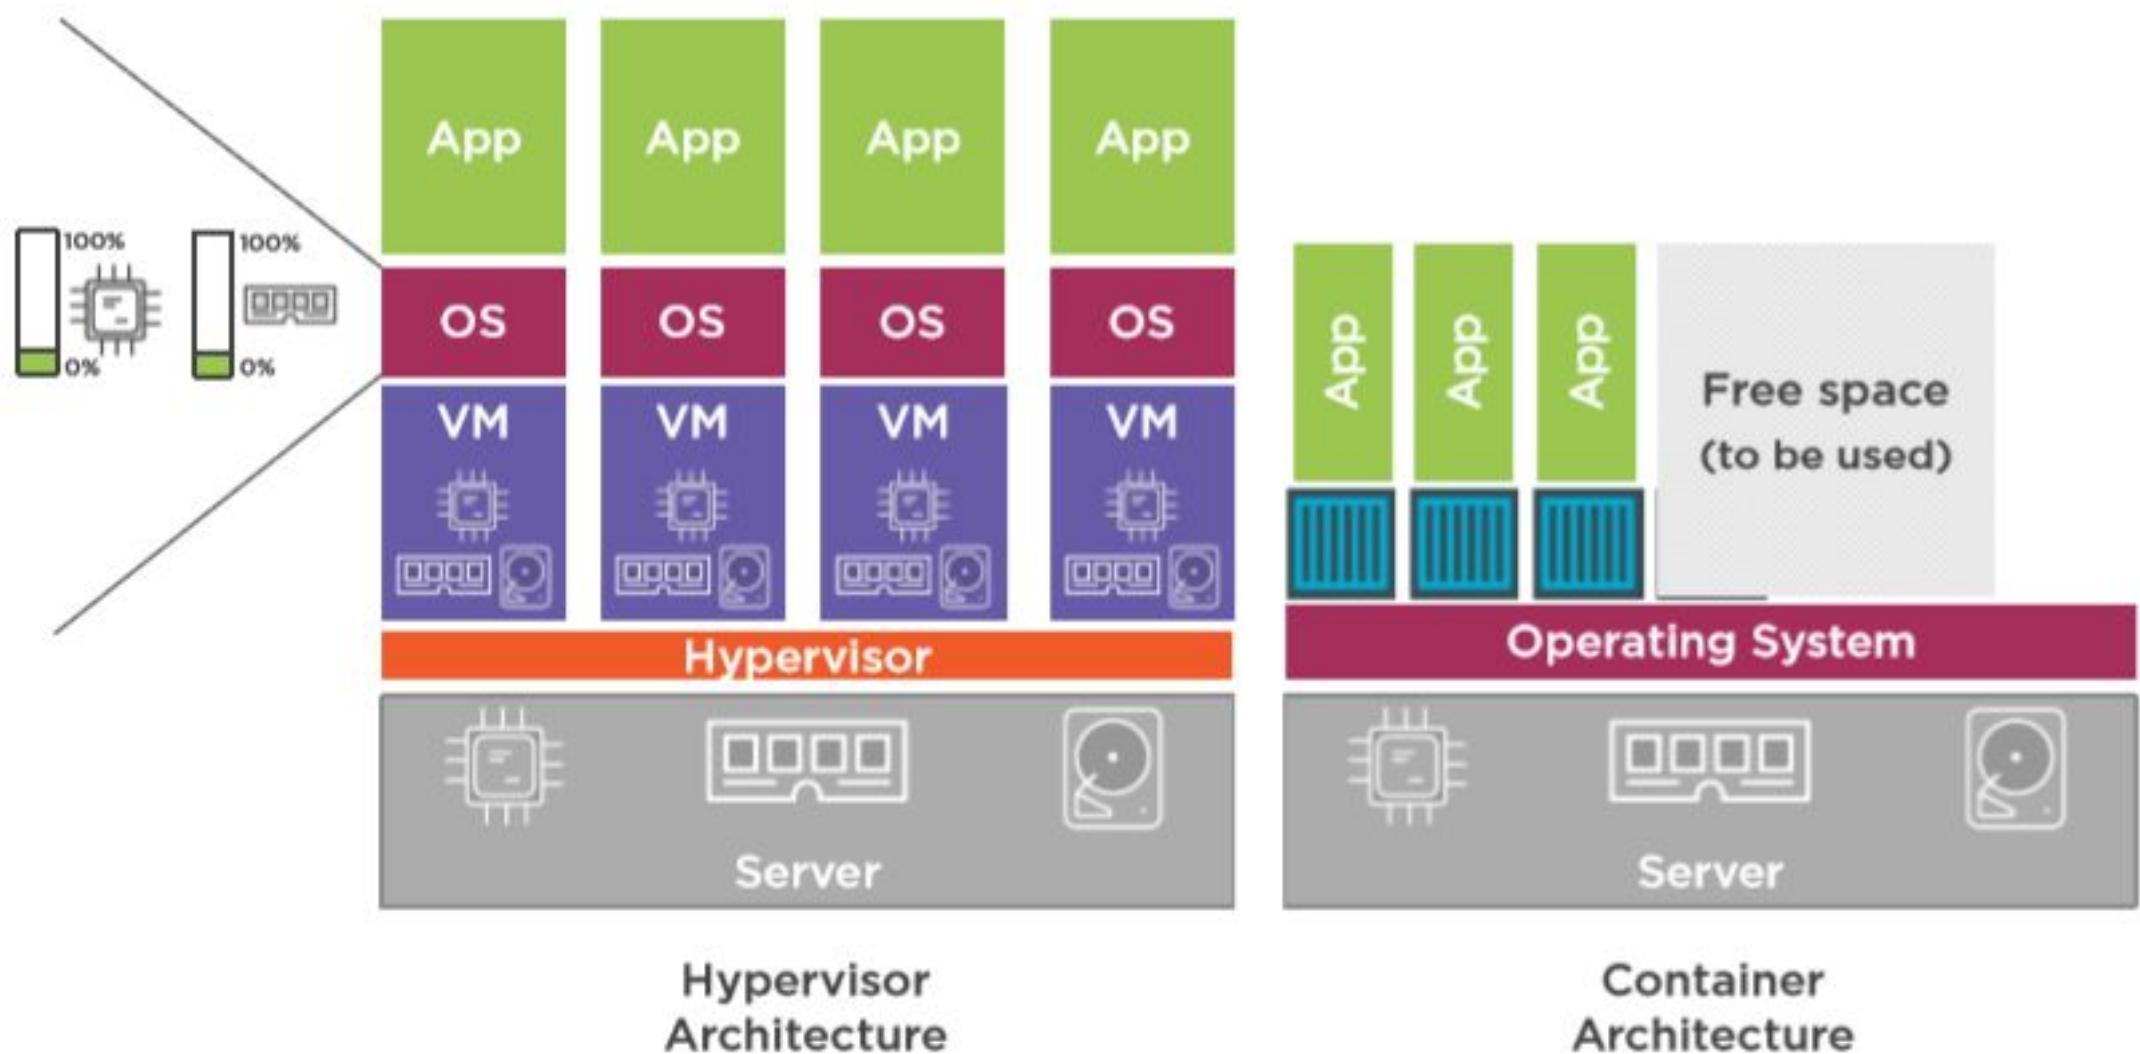
- Namespace?
- Cgroups?
- Layered filesystems  
(Rootfs)?



# Docker



# Architecture de base de Kubernetes



# Architecture de base de Kubernetes

Past  
(physical servers)



Wasted!

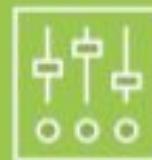


Wasted!



Wasted!

Present  
(Hypervisor virtualization)



More efficient  
than physical  
servers

Could be better!

Present/Future  
(Containers)



More efficient  
than Hypervisor  
virtualization

Virtualization 2.0

Less mature than  
Hypervisor virtualization

Less mature than  
Hypervisor ecosystem

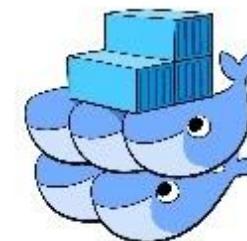
# Orchestration de containers?

## Container Orchestration Tools



**MESOS**

Marathon (Mesosphere)



Docker Swarm



Nomad (HashiCorp)



Kubernetes

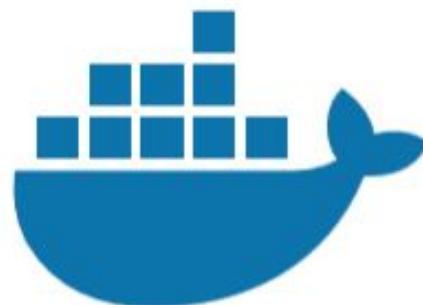
# Kubernetes?

{ vCenter }

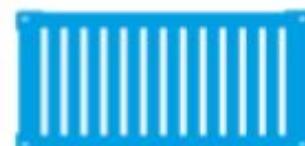


Scheduling, scaling, healing,  
updating...

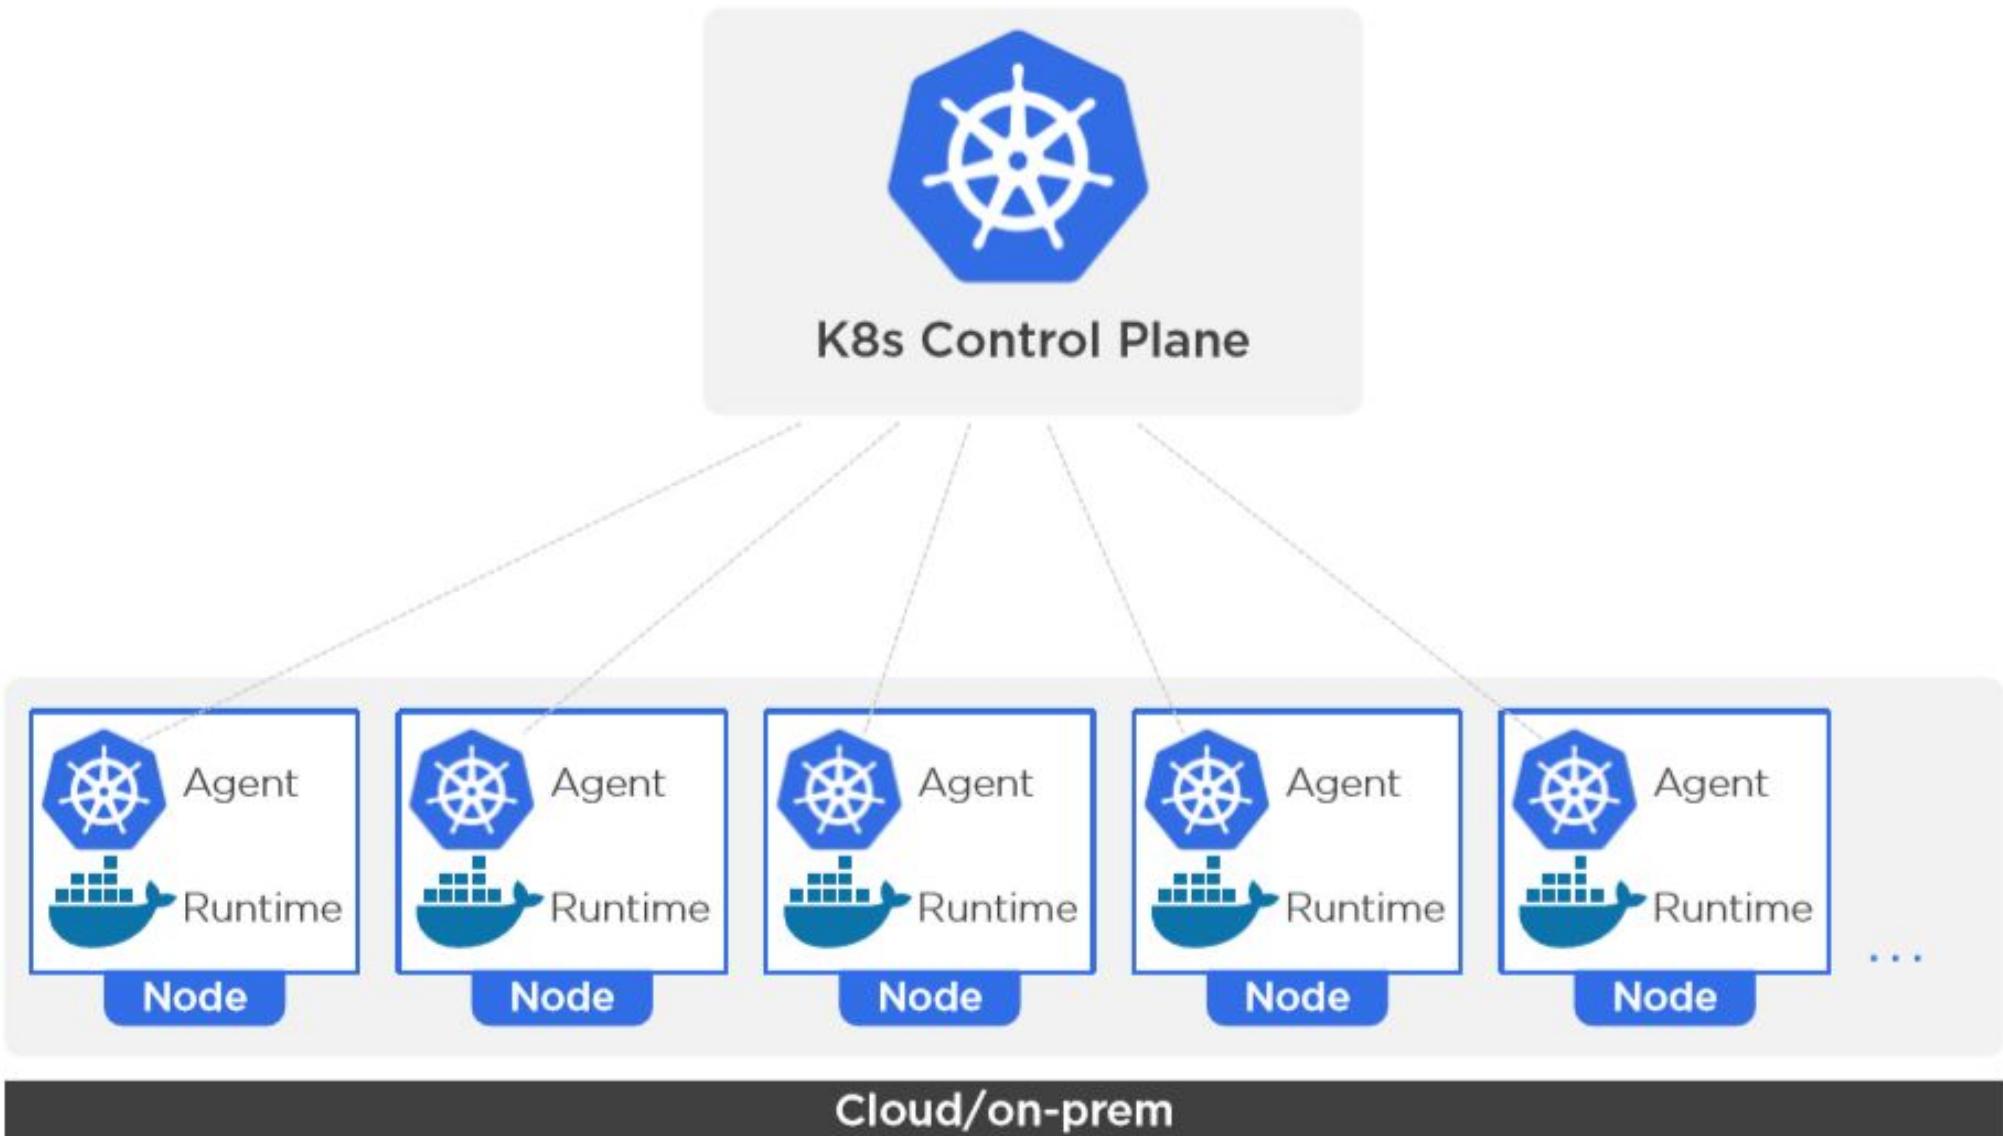
{ ESXi }



start | stop | delete ...

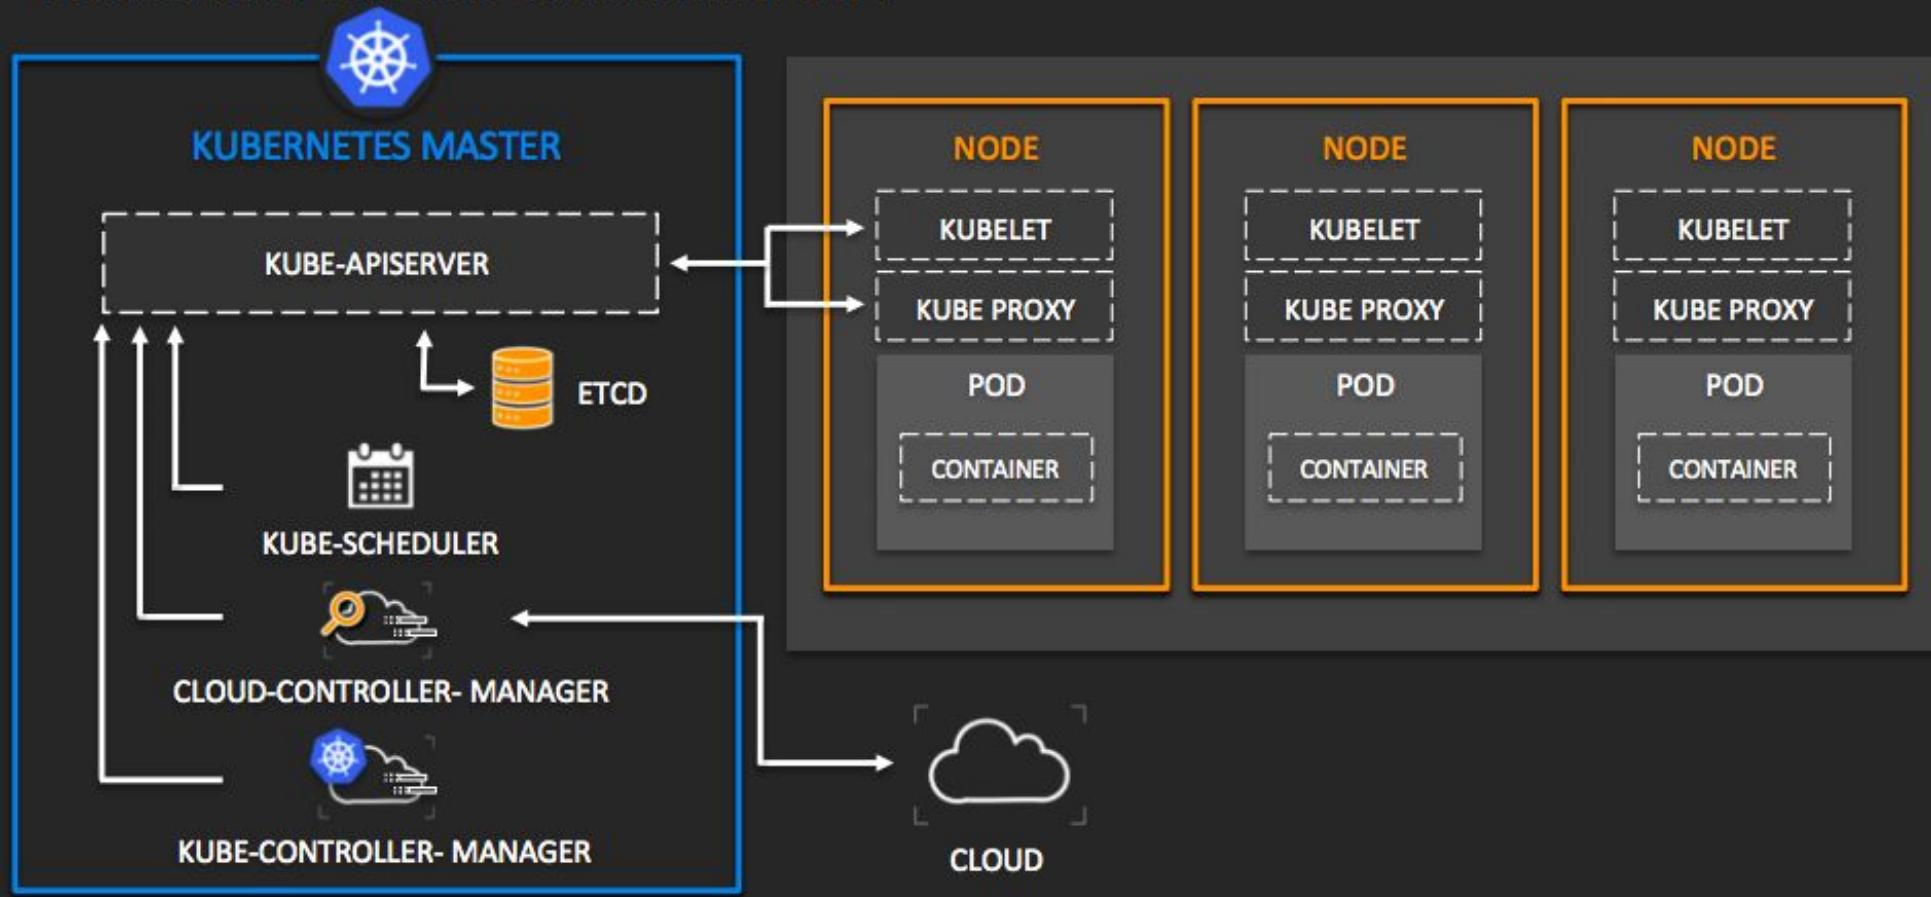


# Kubernetes?



# Architecture de base de Kubernetes

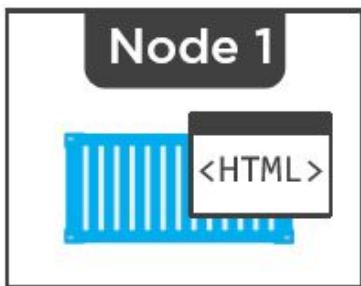
## BASIC KUBERNETES ARCHITECTURE



# Stateless vs Stateful

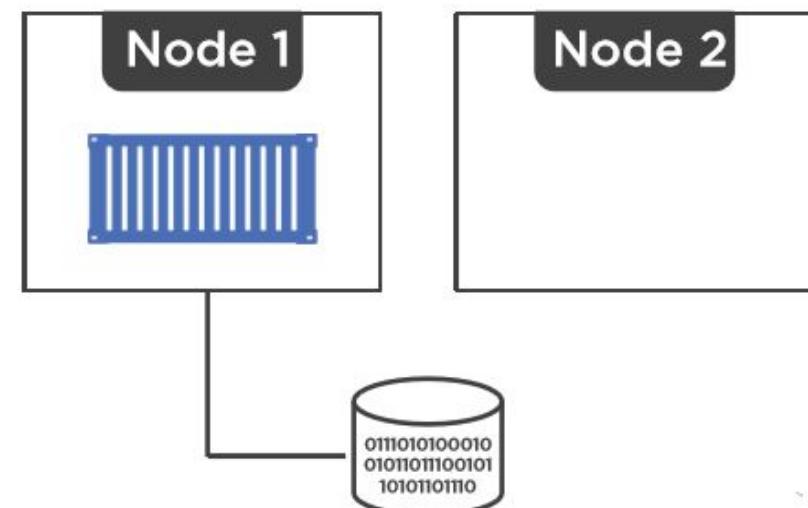
## ✓ Stateless

- Doesn't remember stuff



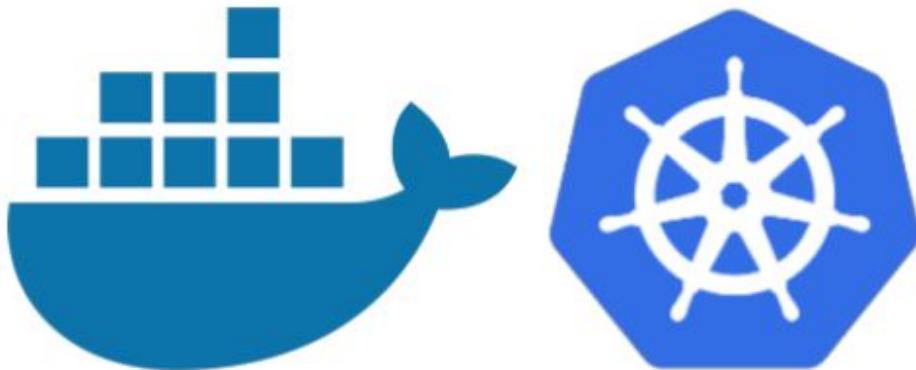
## Stateful ✓

- Has to remember stuff



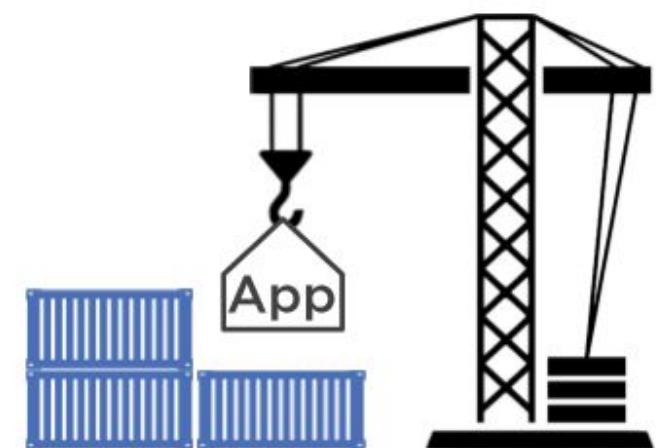
# Migrating d'applications legacy

Stateful



Added features for  
stateful apps

Legacy



Some apps can be  
migrated directly to  
containers

**Helm!**



## Attractions of Containers



### No guest OS

- Platform independent
- Considerably smaller than VM images

### Lightweight

- Small and fast
- Quick to start
- Speeds up autoscaling

### Hybrid, multi-cloud

- Hybrid: Work on-premise and on cloud
- Multi-cloud: Not tied to any specific cloud platform



## Standalone Container Limitations

### No autohealing

- Crashed containers won't restart automatically
- Need higher level orchestration

### No scaling or autoscaling

- Overloaded containers don't spawn more automatically
- Need higher level orchestration

### No load balancing

- Containers can't share load automatically
- Need higher level orchestration

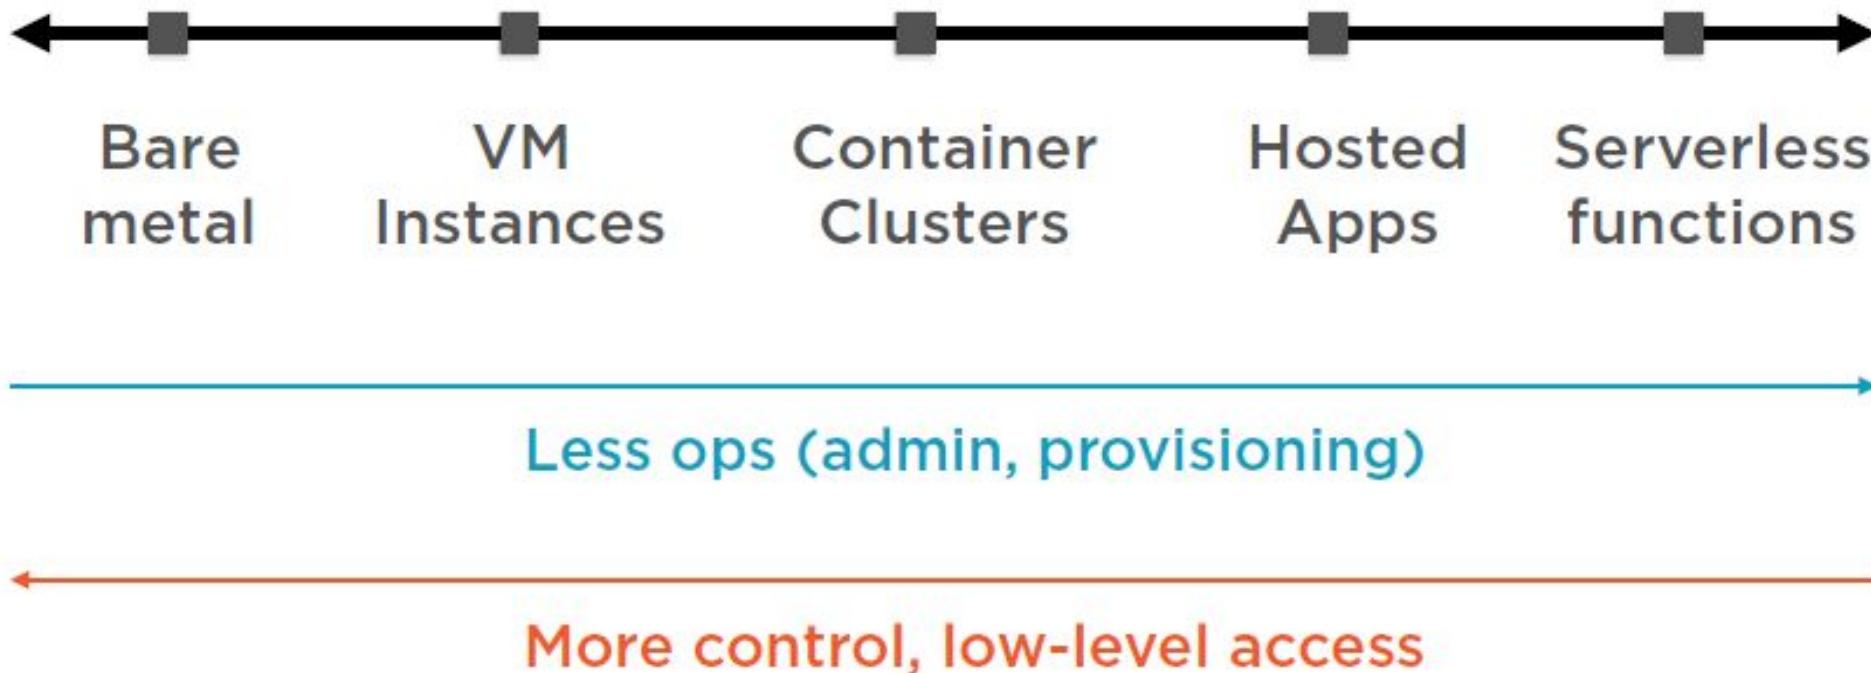
### No isolation

- Crashing containers can take each other down
- Need sandbox to separate them

# Kubernetes

Orchestration technology for containers - convert isolated containers running on different hardware into a cluster

# Compute Choices



Kubernetes is fast emerging as middle-ground between IaaS and PaaS in a hybrid, multi-cloud world



## Hybrid, Multi-cloud

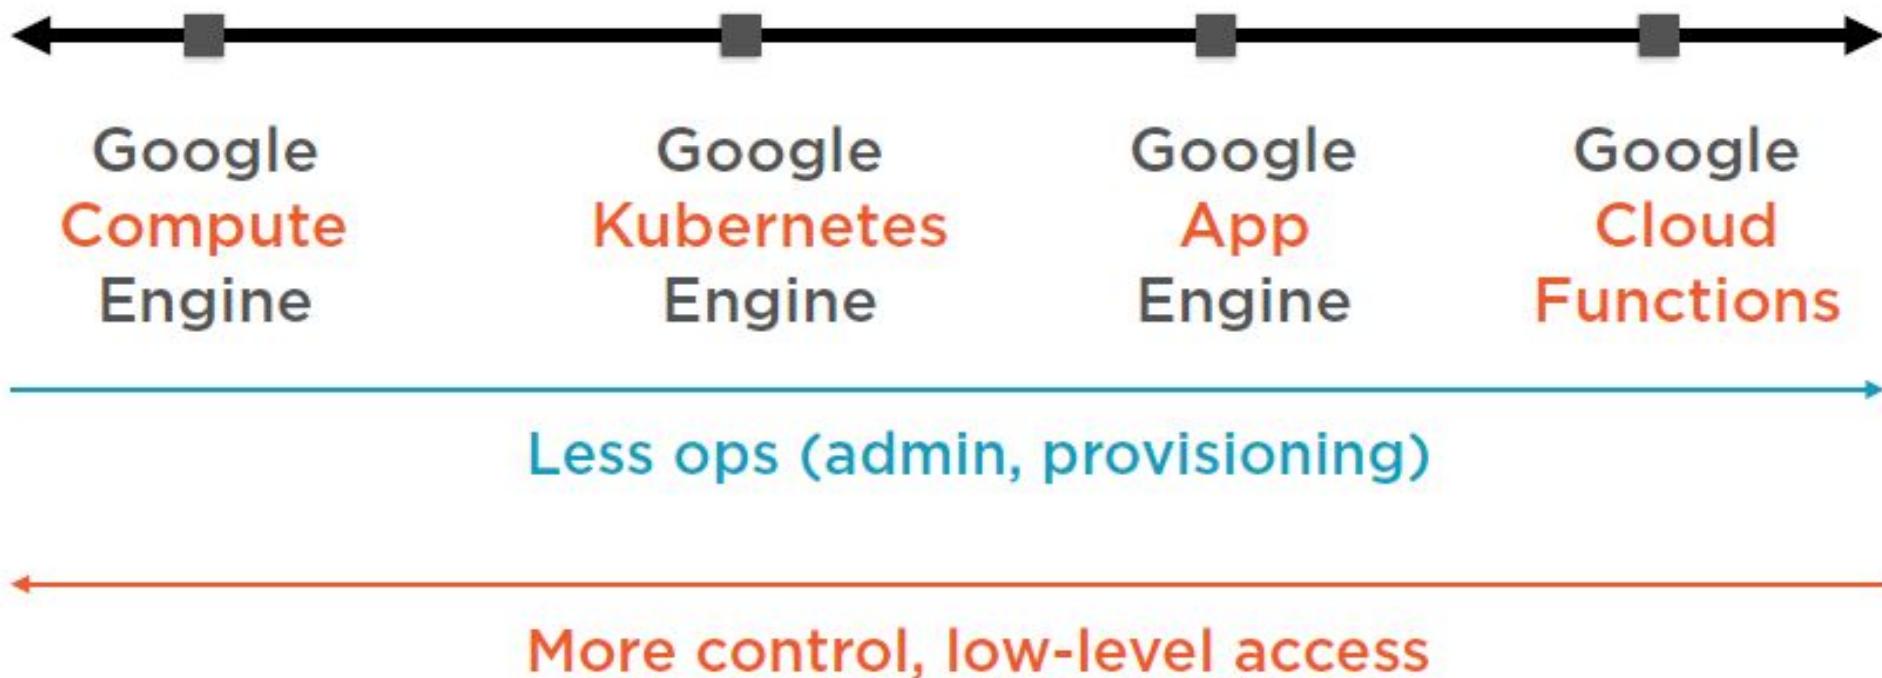
### **Hybrid: Runs on-premise and on cloud**

- Provides smooth migration path

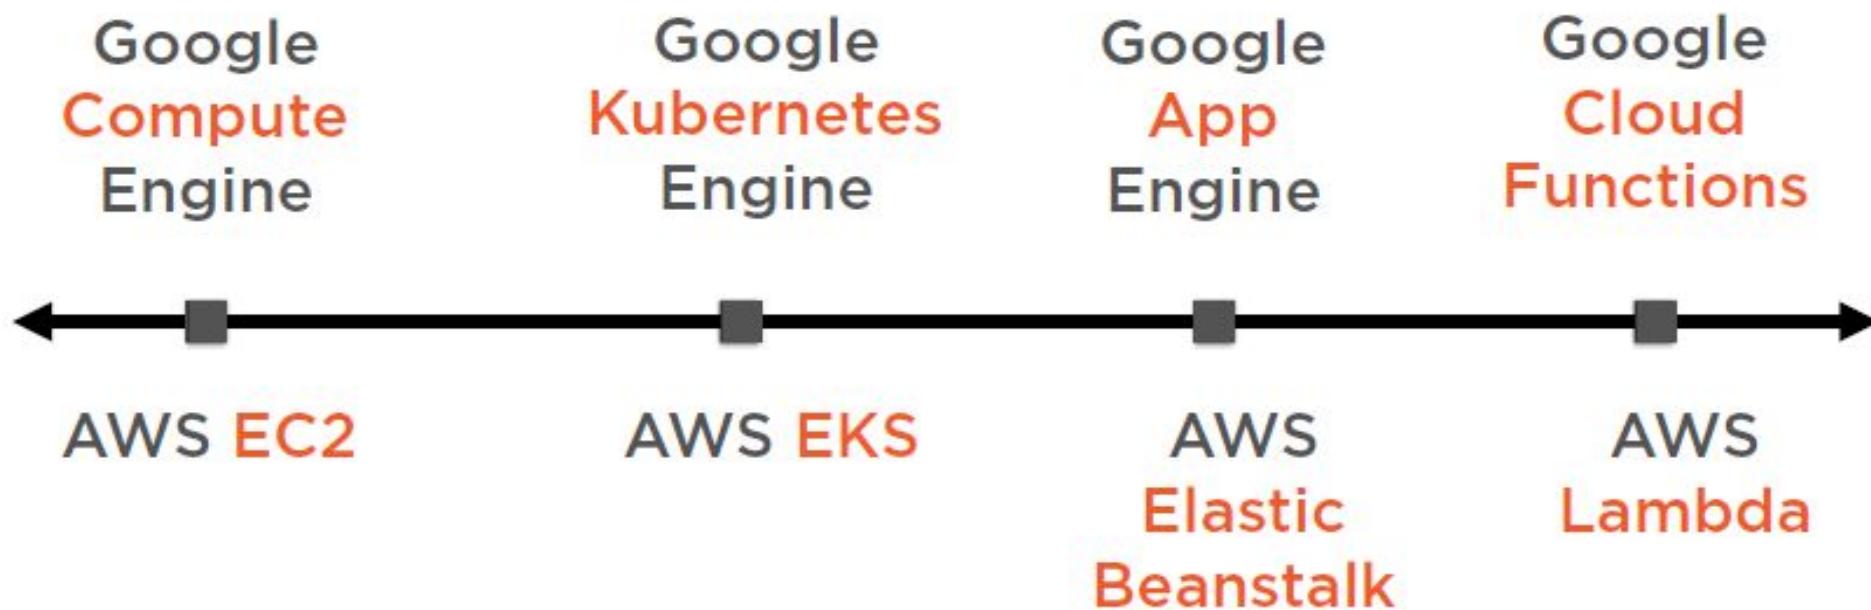
### **Multi-cloud: Supported by all big cloud platforms**

- Important for strategic independence
- Amazon-Whole Foods merger

# GCP Compute Choices



# GCP Compute Choices



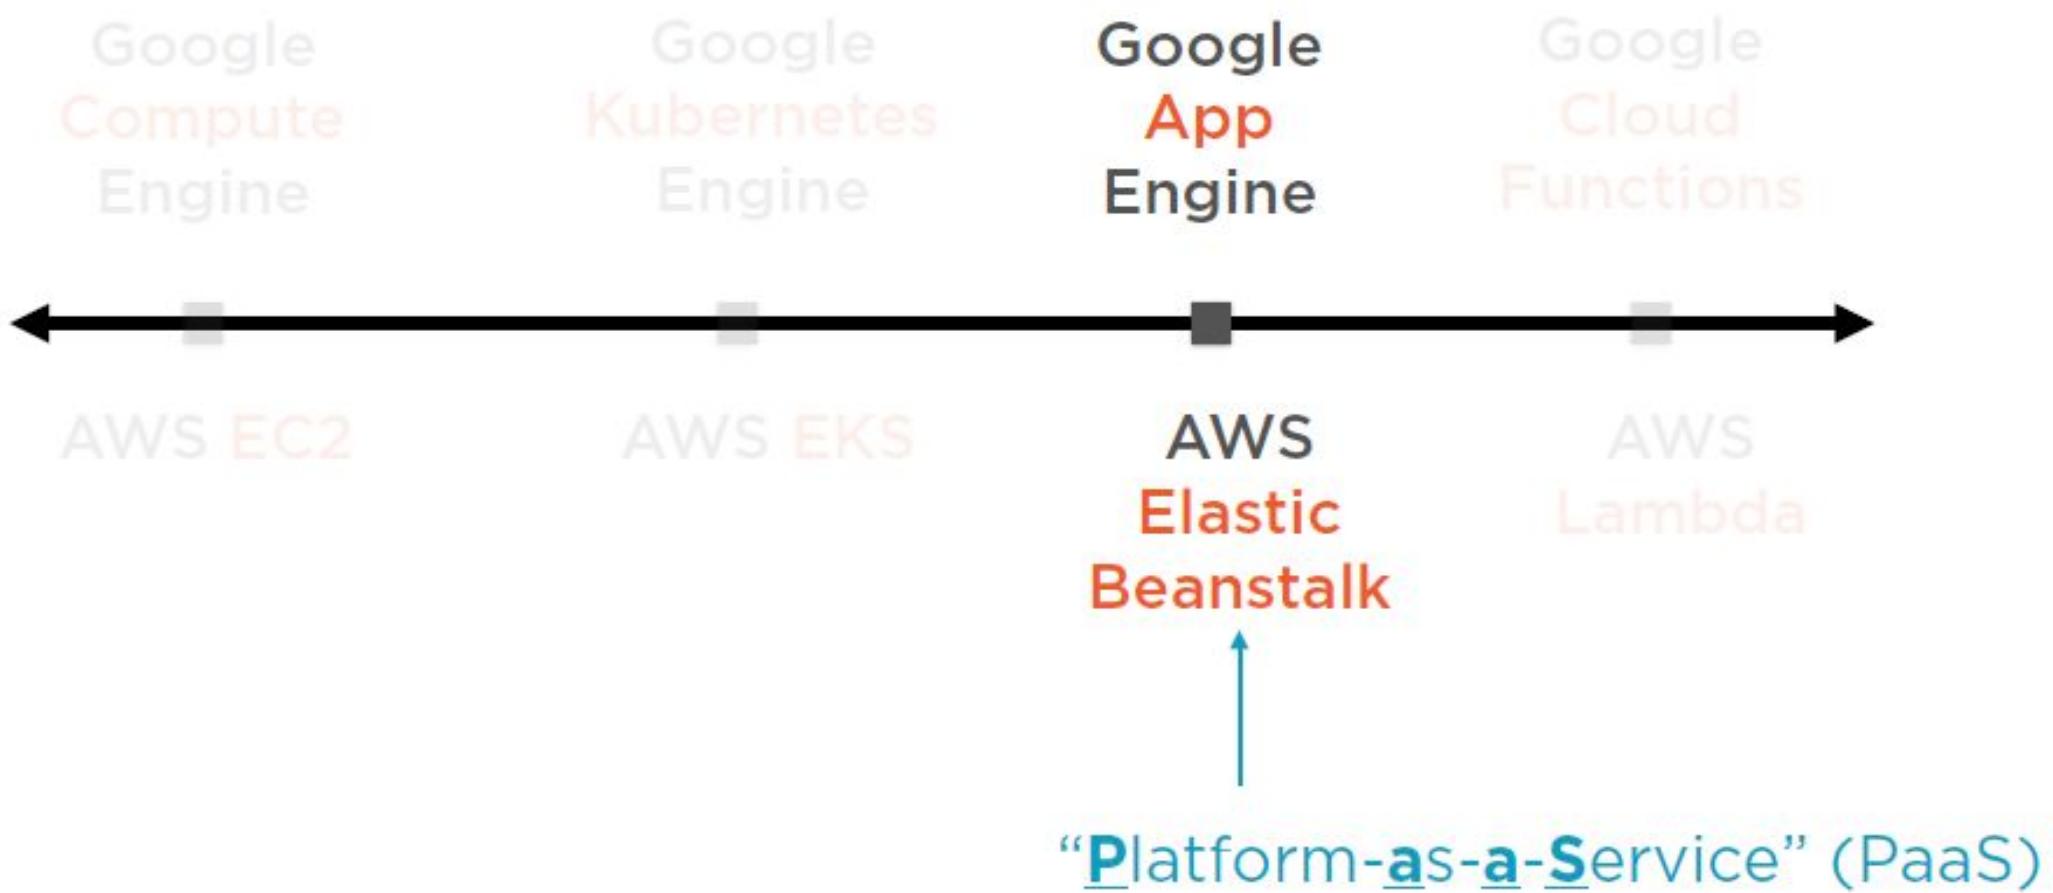
Every major cloud platform supports the same range of compute choices

# GCP Compute Choices



“Infrastructure-as-a-Service” (IaaS)

# GCP Compute Choices



# IaaS vs. PaaS

## Infrastructure-as-a-Service

Heavy operational burden

Migration is hard

## Platform-as-a-Service

Provider lock-in

Migration is very hard

# Compute Choices

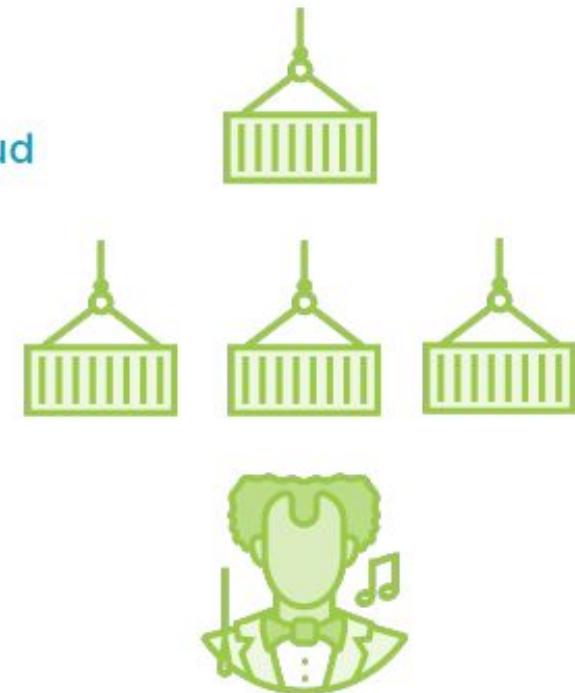


**Containers**

hybrid, multi-cloud

**Container Clusters**

**Kubernetes**

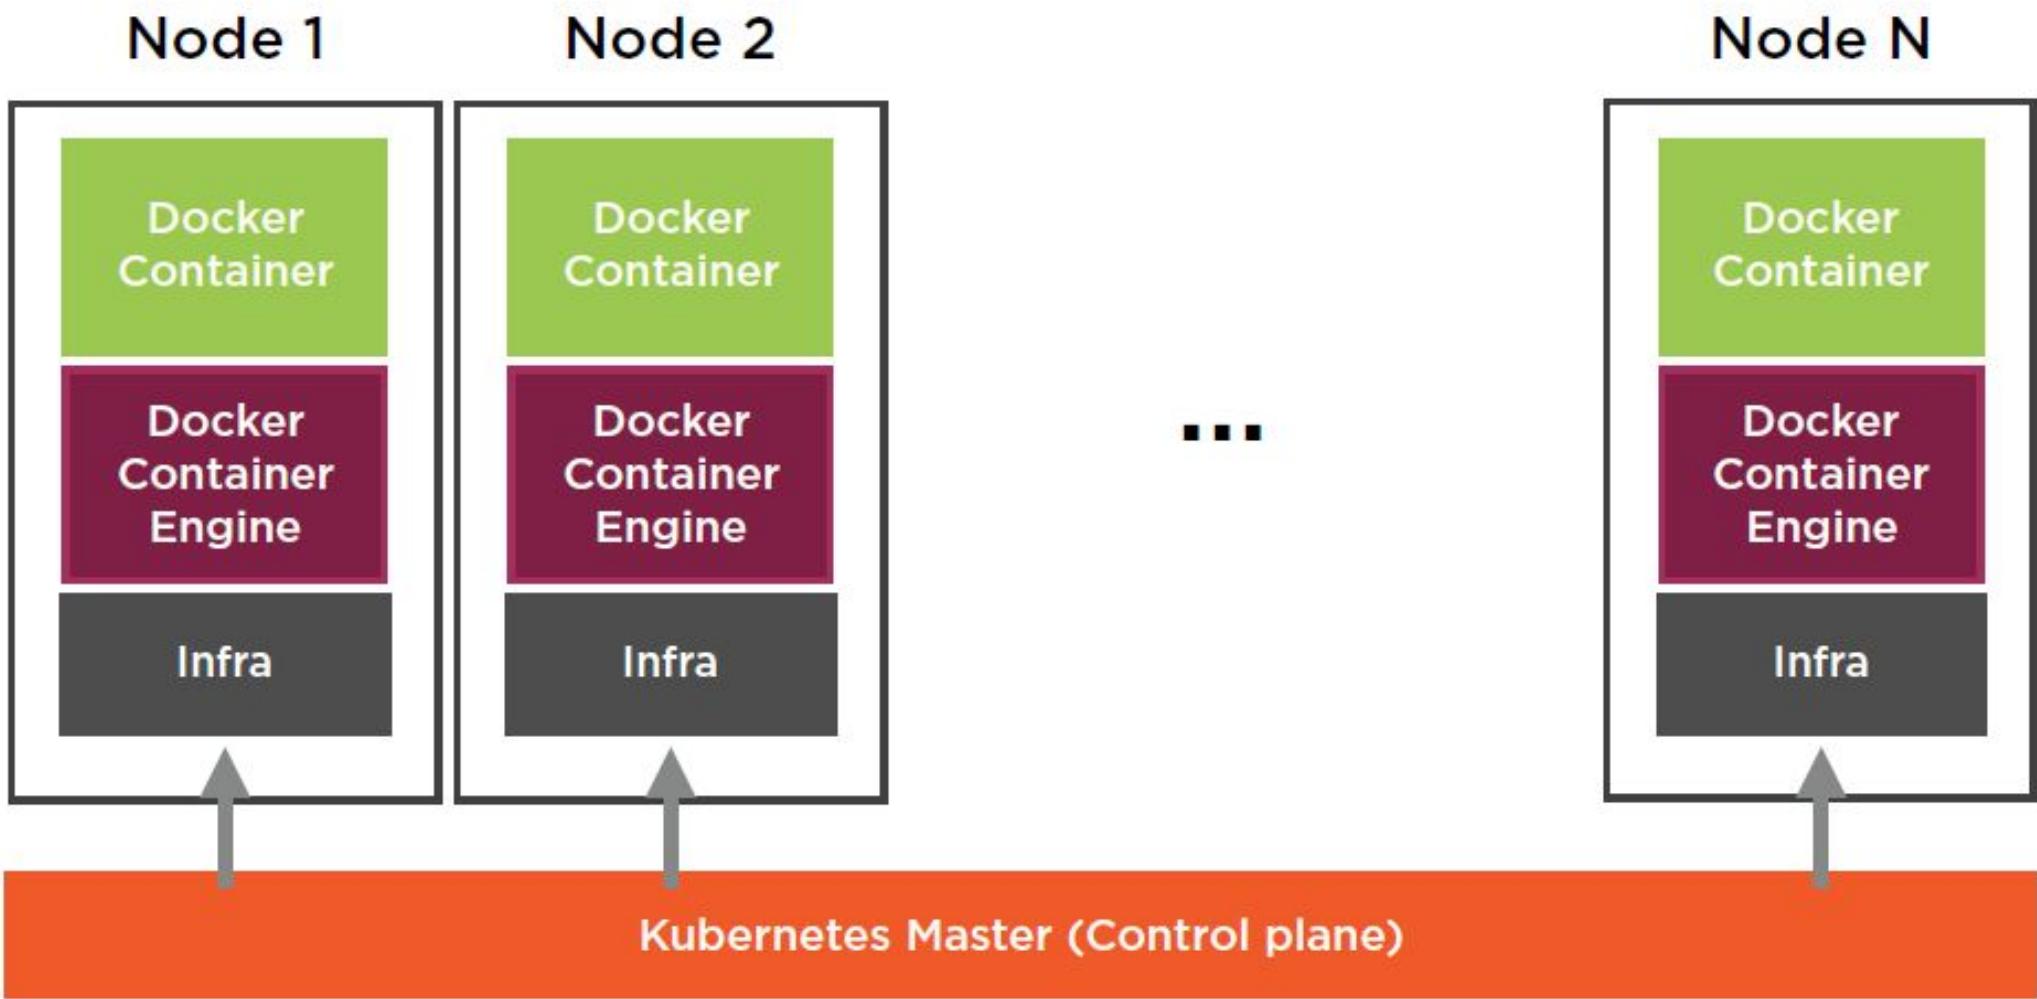


# Kubernetes as Orchestrator

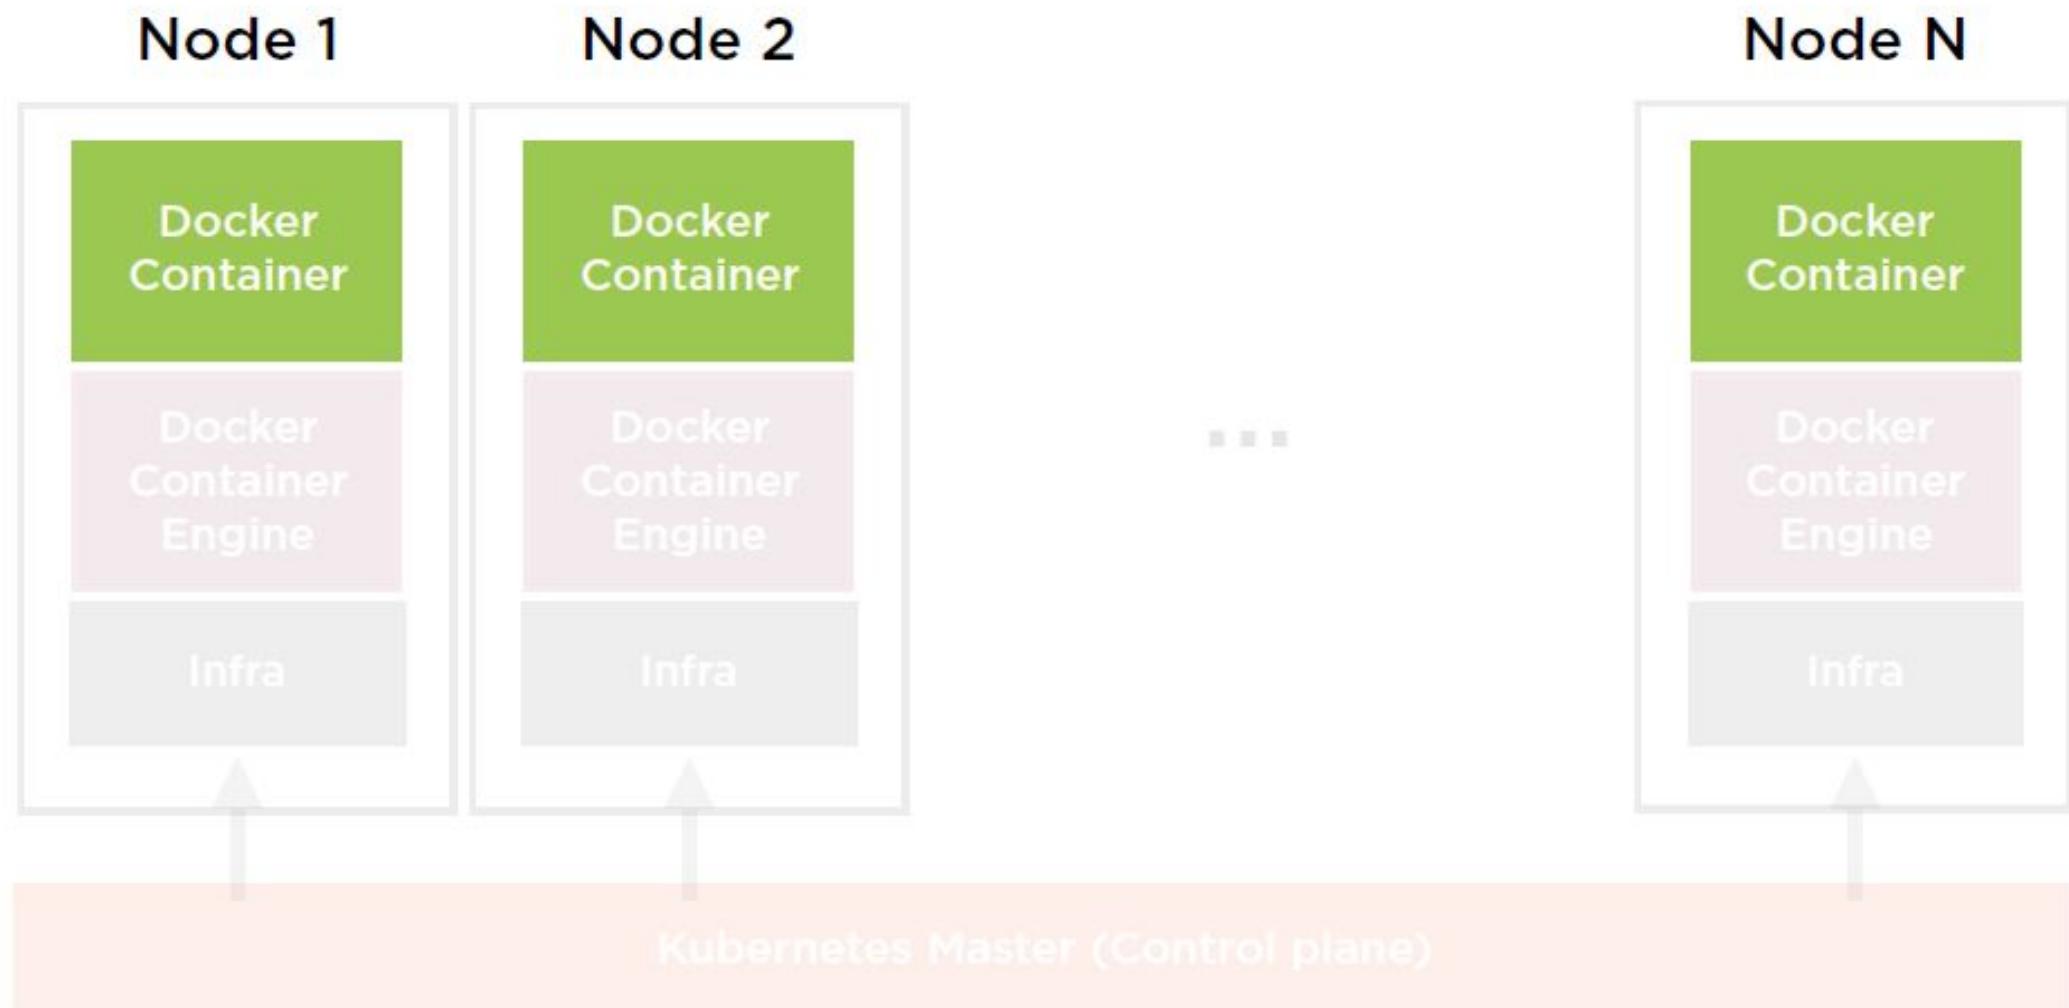


- Fault-tolerance**
- Autohealing**
- Isolation**
- Scaling**
- Autoscaling**
- Load balancing**

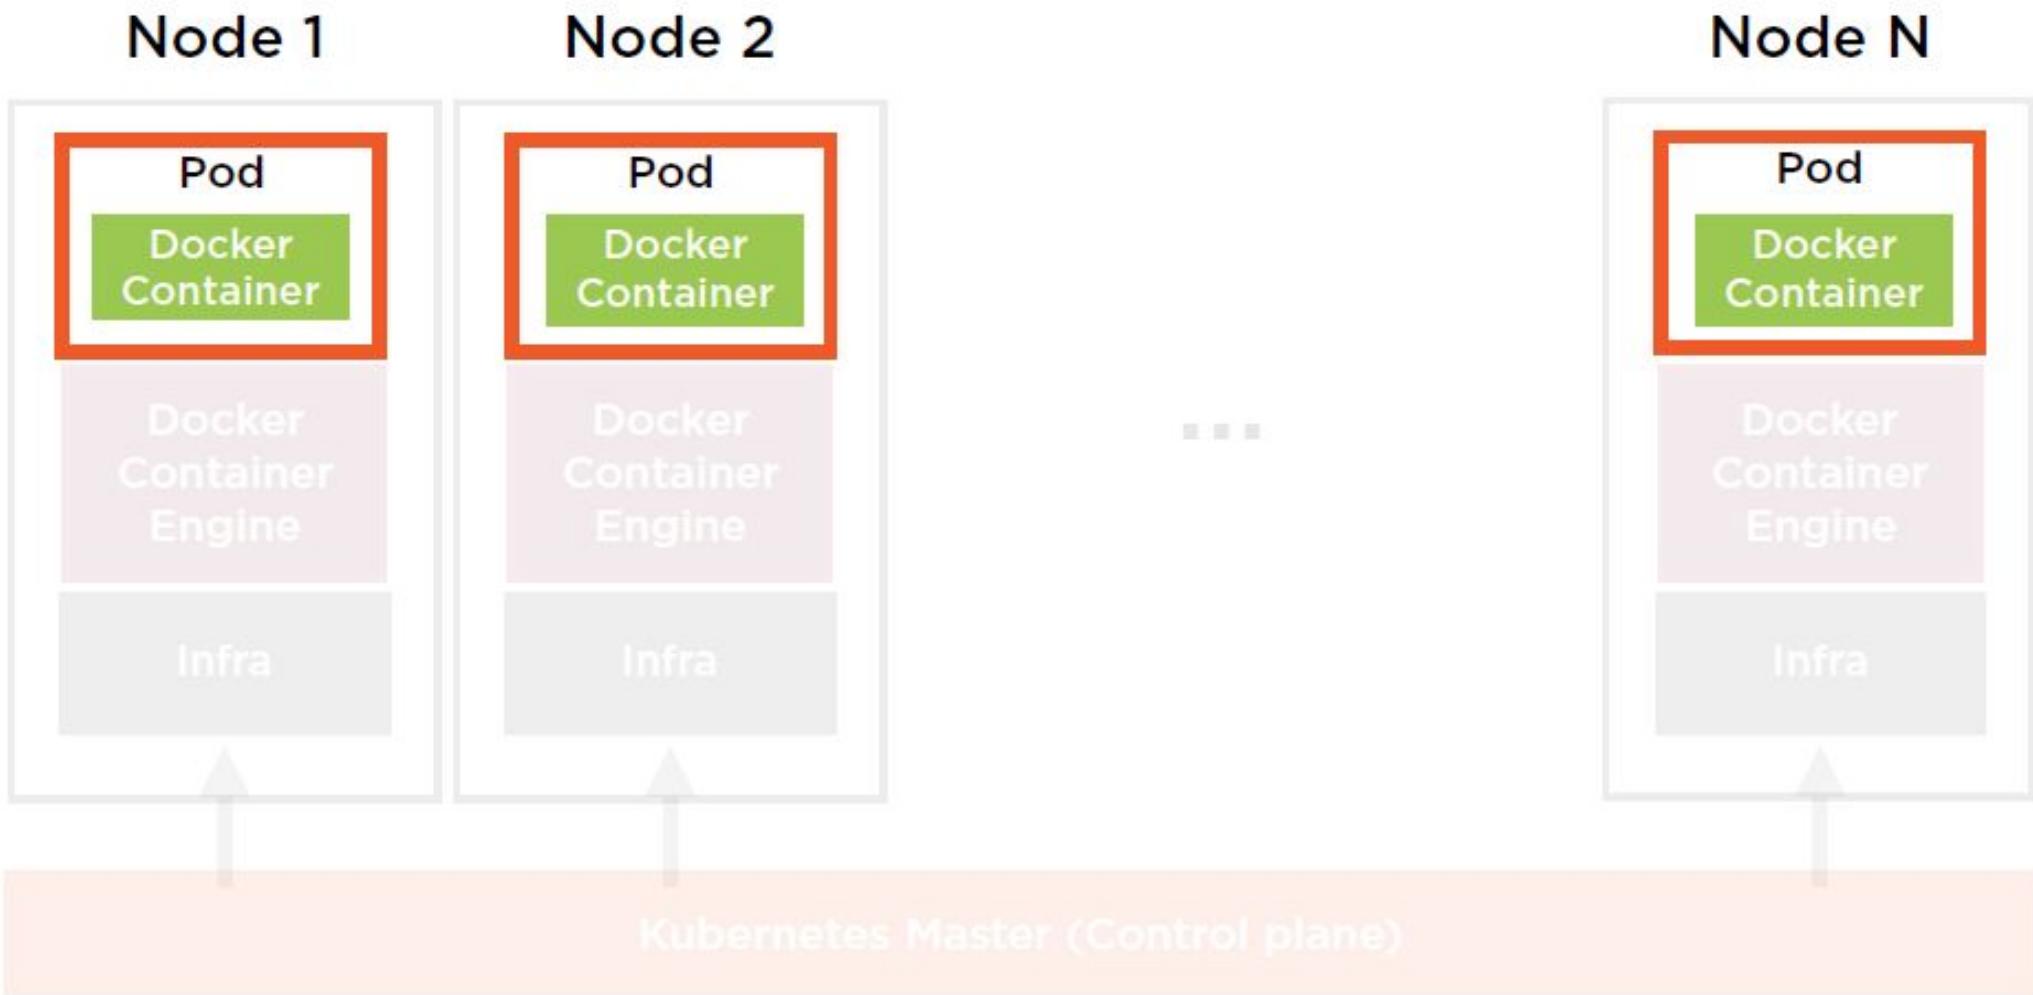
# Kubernetes: Cluster Orchestration



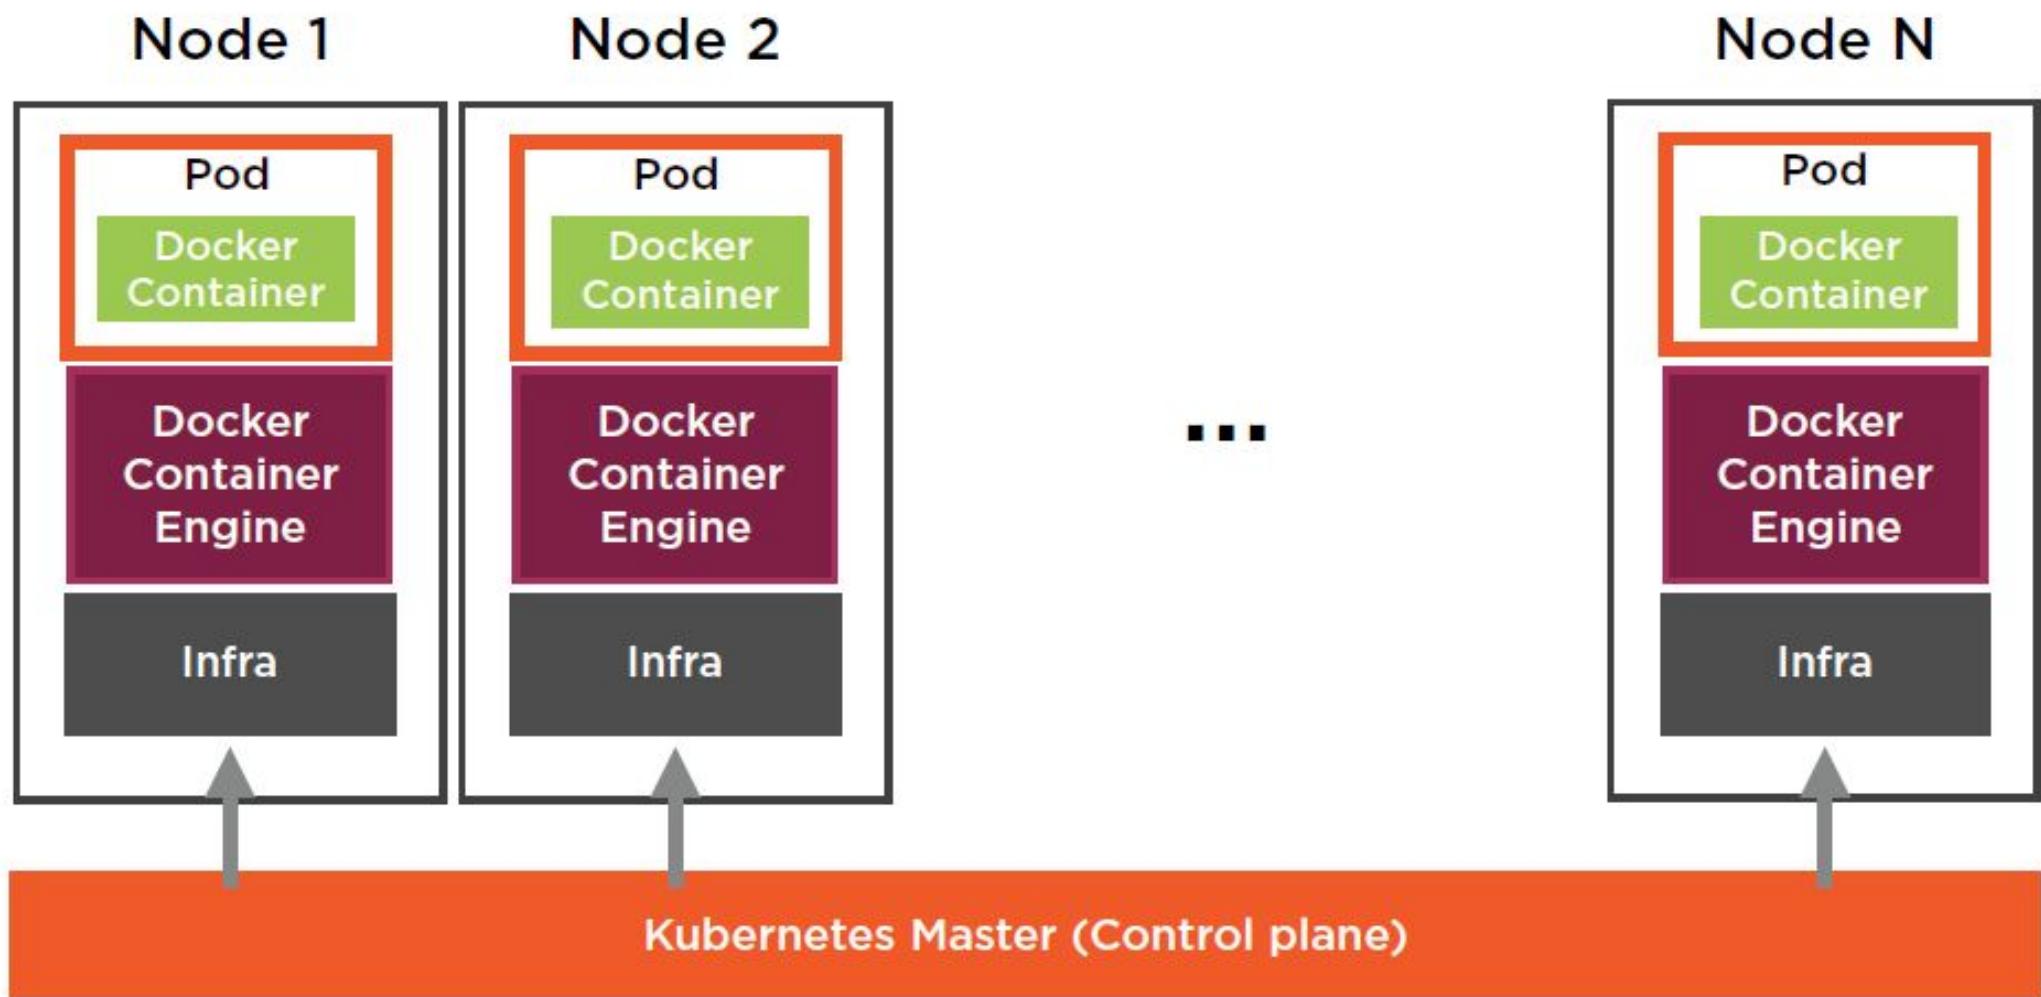
# Kubernetes: Cluster Orchestration



# Kubernetes: Containers Run Within Pods



# Kubernetes: Cluster Orchestration



# Pods as Atomic Units

## Container deployment

All containers in pod are deployed,  
or none are

## Node association

Entire pod is hosted on the same  
node

**Pod is atomic unit of deployment in Kubernetes**

# The ReplicaSet Object

**Multiple identical pods which  
are replicas of each other**

# ReplicaSet

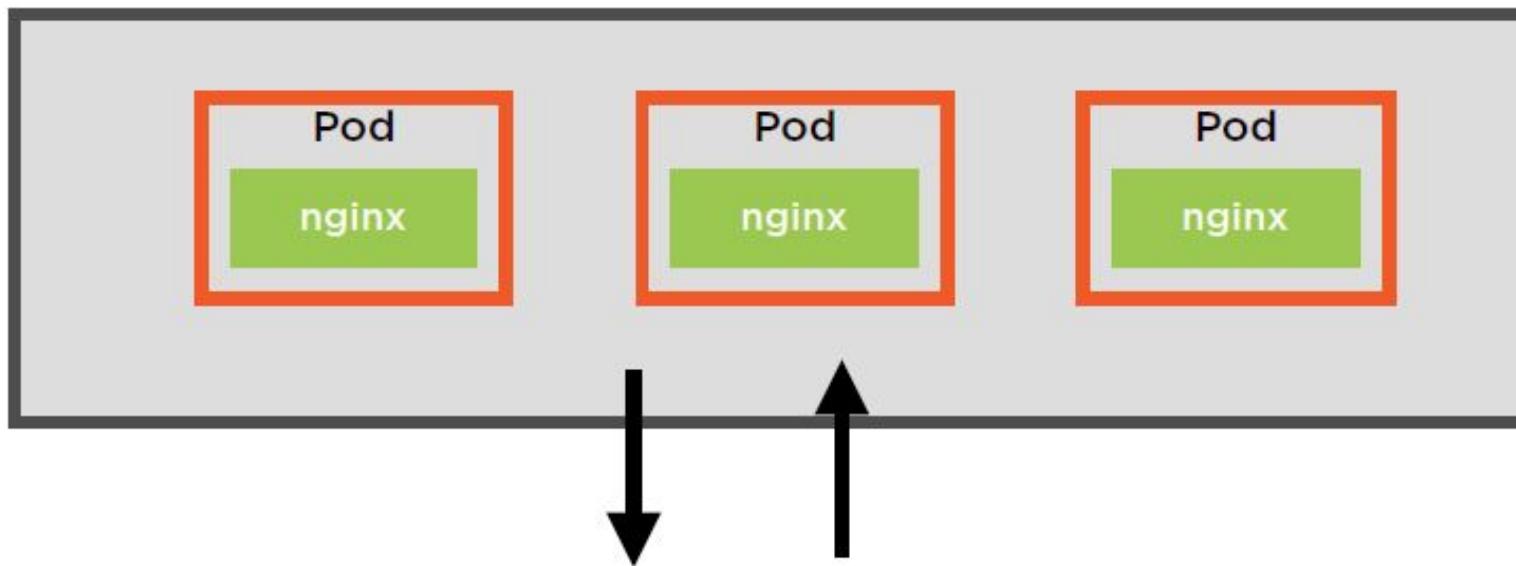


Self healing and autoscaling for our pods

# The Deployment Object

**Adds on deployment and  
rollback functionality**

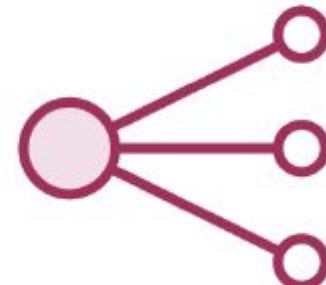
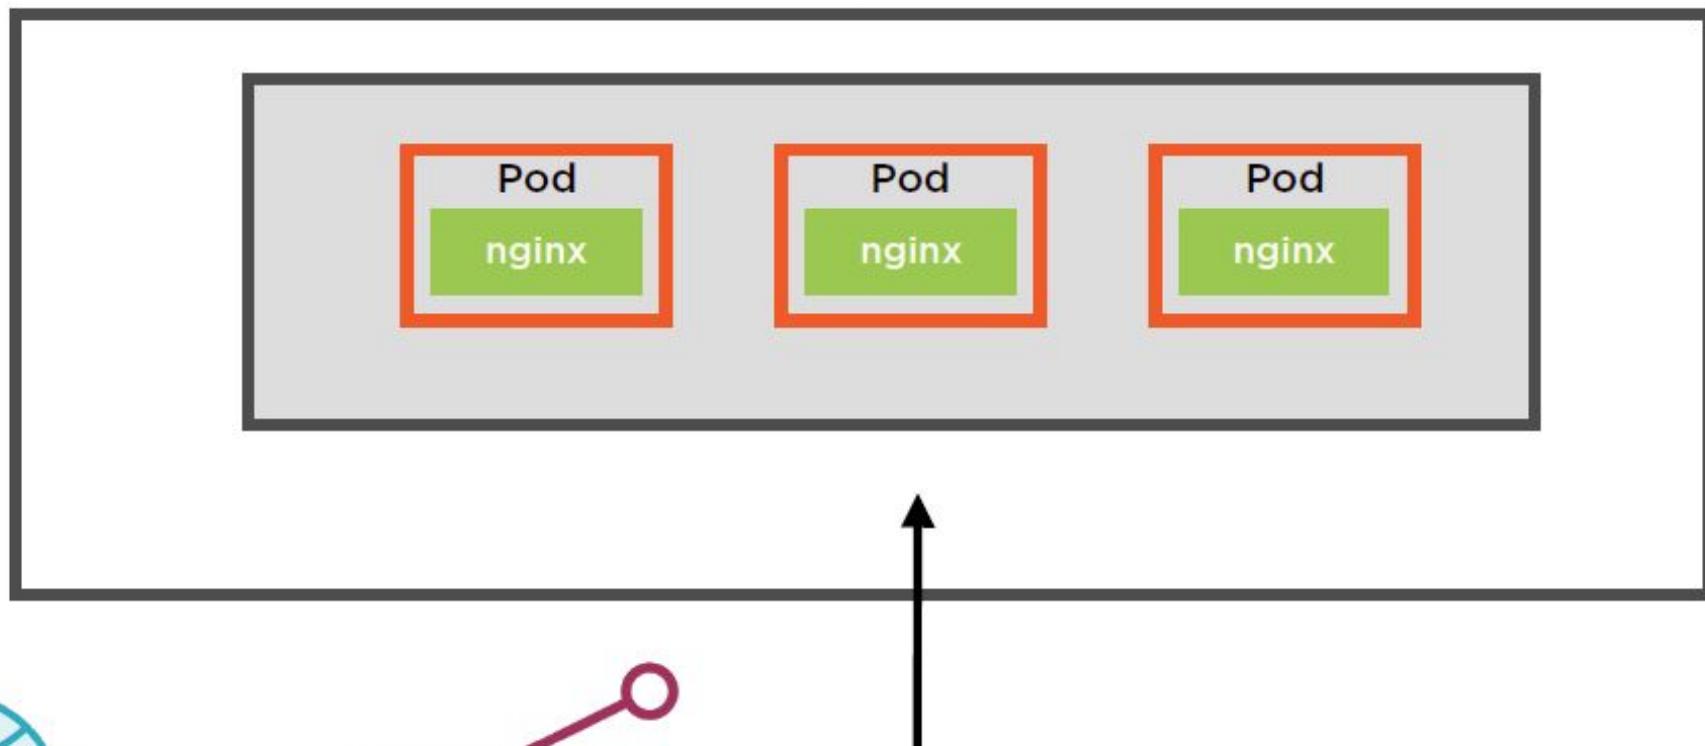
# Deployment



Support for versions, and production-level operations such as rollbacks

Services provide stable IP  
addresses for external  
connections and load balancing

## Service



**Lab:**

**Les fondamentaux de K8S**

---



# CGroups and Namespaces

---

# CGroups and Namespaces

Dans cet exercice, vous apprendrez les bases de CGroups (Control Groups) et Namespces pour appliquer des restrictions de sécurité aux conteneurs. Voici quelques exemples des types de groupes de contrôle et d'espaces de noms qui existent.

## Exemples de Cgroups:

- cpu-shares
- cpuset-cpus
- memory-reservation
- kernel-memory
- blkio-weight (block IO)
- device-read-iops
- device-write-iops

# CGroups and Namespaces

## Exemples de Namespaces:

Cgroup	CLONE_NEWCGROUP	Cgroup root directory
IPC	CLONE_NEWIPC	System V IPC, POSIX message queues
Network	CLONE_NEWWNET	Network devices, stacks, ports, etc.
Mount	CLONE_NEWNS	Mount points
PID	CLONE_NEWPID	Process IDs
User	CLONE_NEWUSER	User and group IDs
UTS	CLONE_NEWUTS	Hostname and NIS domain name

## 1- Définir les limites de mémoire

Les CGroups contrôlent la quantité de ressources qu'un processus peut utiliser. En ajoutant des restrictions, vous pouvez offrir une qualité de service garantie aux applications en vous assurant qu'elles disposent de suffisamment d'espace disponible.

Il est également possible de protéger le système contre les utilisateurs ou applications potentiellement malveillants visant à exécuter des applications de déni de service (DoS) via l'épuisement des ressources.

Cela peut également aider à limiter les applications des fuites de mémoire ou d'autres bogues de programmation en définissant des limites supérieures.

## 1- Définir les limites de mémoire

- Exemple:

> **docker run -d --name mb100 --memory 100m alpine top**

```
[root @ host01 ~] # docker run -d --name mb100 --memory 100m sommet alpin  
7849578514a5da111b1ba6d868d36a8daf2e82d6a0e1cfa6e10ff292820110b4
```

L'utilisation de la mémoire et les limites des conteneurs peuvent être identifiées via la commande *docker stats* .

> **docker stats --no-stream**

```
[root @ host01 ~] # stats docker --no-stream  
CONTAINER ID NAME CPU% MEM USAGE / LIMIT MEM% NET I / O BLOCK I / O PI  
DS  
7849578514a5 mb100 0,00% 436KiB / 100MiB 0,43% 5,33kB / 90B 1,06MB / 0B 1
```

## 2- Définir les partages CPU

- Alors que les limites de mémoire définissent un maximum défini, les limites de processeur sont basées sur les partages.
- Ces parts représentent un poids entre le temps de traitement d'un processus par rapport à un autre.
- Si un processeur est inactif, le processus utilisera toutes les ressources disponibles. Si un deuxième processus nécessite la CPU, le temps CPU disponible sera partagé en fonction de la pondération.

## 2- Définir les partages CPU

- Prenons un exemple de démarrage d'un conteneur avec différents partages. Le paramètre **--cpu-shares** définit un partage entre 0 et 768.
- Si un conteneur définit une part de 768, tandis qu'un autre définit une part de 256, le premier conteneur aura une part de 75%, l'autre ayant 25% du total de part disponible.
- Ces chiffres sont dus à l'approche de pondération pour le partage du processeur au lieu d'une capacité fixe.
- Lançons le premier conteneur qui sera autorisé à avoir 75% de la part et Le deuxième conteneur qui sera limité à 25%.

## 2- Définir les partages CPU

> docker run -d --name c768 --cpuset-cpus 0 --cpu-shares 768 benhall/stress

```
[root@host01 ~]# docker run -d --name c256 --cpuset-cpus 0 --cpu-shares 256 benhall/stress
0f77bb0f7950fcc3e6982bd76202caae414e3919ba61dc90c1ab80ddce932a0
[root@host01 ~]# sleep 5
[root@host01 ~]# docker stats --no-stream
CONTAINER ID        NAME          CPU %     MEM USAGE / LIMIT   MEM %     NET I/O           BLOCK I/O          PI
DS
0f77bb0f7950        c256          25.30%    764KiB / 737.5MiB  0.10%    3.64kB / 90B      0B / 0B            3
d937be5b4b3b        c768          74.51%    712KiB / 737.5MiB  0.09%    4.77kB / 90B      0B / 0B            3
[root@host01 ~]# docker rm -f c768 c256
c768
c256
>docker run -d --name c256 --cpuset-cpus 0 --cpu-shares 256 benhall/stress
```

> sleep 5

> docker stats --no-stream

> docker rm -f c768 c256

```
[root@host01 ~]# docker run -d --name c256 --cpuset-cpus 0 --cpu-shares 256 benhall/stress
0f77bb0f7950fcc3e6982bd76202caae414e3919ba61dc90c1ab80ddce932a0
[root@host01 ~]# sleep 5
[root@host01 ~]# docker stats --no-stream
CONTAINER ID        NAME          CPU %     MEM USAGE / LIMIT   MEM %     NET I/O           BLOCK I/O          PI
DS
0f77bb0f7950        c256          25.30%    764KiB / 737.5MiB  0.10%    3.64kB / 90B      0B / 0B            3
d937be5b4b3b        c768          74.51%    712KiB / 737.5MiB  0.09%    4.77kB / 90B      0B / 0B            3
[root@host01 ~]# docker rm -f c768 c256
c768
c256
```

- Il est important de noter qu'un processus peut avoir 100% du partage, quel que soit le poids défini, si aucun autre processus n'est en cours d'exécution.

## 3- Utiliser les « Network Namespace »

- Alors que les groupes de contrôle contrôlent la quantité de ressources qu'un processus peut utiliser, les namespaces contrôlent ce qu'un processus voit et accéder.

> **docker run -it alpine ip addr show**

```
[root@ghost01 ~]# docker run -it alpine ip addr show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
8: eth0@if9: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1500 qdisc noqueue state UP
    link/ether 02:42:ac:12:00:02 brd ff:ff:ff:ff:ff:ff
    inet 172.18.0.2/24 brd 172.18.0.255 scope global eth0
        valid_lft forever preferred_lft forever
```

- En changeant l'espace de noms en host , au lieu que le réseau du conteneur soit isolé avec son interface, le processus aura accès à l'interface réseau des machines hôtes.
- Si le processus écoute sur les ports, ils seront écoutés sur l'interface hôte et mappés sur le conteneur.

## 3- Utiliser les « Network Namespace »

> docker run -it --net=host alpine ip addr show

```
[root@host01 ~]# docker run -it --net=host alpine ip addr show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN qlen 1000
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
        inet 127.0.0.1/8 scope host lo
            valid_lft forever preferred_lft forever
        inet6 ::1/128 scope host
            valid_lft forever preferred_lft forever
2: ens3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP qlen 1000
    link/ether 02:42:ac:11:00:2b brd ff:ff:ff:ff:ff:ff
        inet 172.17.0.43/16 brd 172.17.255.255 scope global ens3
            valid_lft forever preferred_lft forever
        inet6 fe80::bb30:716c:4fa4:9465/64 scope link
            valid_lft forever preferred_lft forever
3: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue state DOWN
    link/ether 02:42:1a:78:f1:bd brd ff:ff:ff:ff:ff:ff
        inet 172.18.0.1/24 brd 172.18.0.255 scope global docker0
            valid_lft forever preferred_lft forever
```

## 4- Utiliser le Namespace Pid

- Comme pour les réseaux, les processus qu'un conteneur peut voir dépendent également de l'espace de noms auquel il appartient.
- En modifiant l'espace de noms Pid, un conteneur peut interagir avec des processus au-delà de sa portée normale.
- **Exemple:**

Le premier conteneur s'exécutera dans son espace de noms de processus. En tant que tels, les seuls processus auxquels il peut accéder sont ceux lancés dans le conteneur.

**> docker run -it alpine ps aux**

```
[root@host01 ~]# docker run -it alpine ps aux
PID  USER      TIME  COMMAND
 1 root      0:00 ps aux
```

## 4- Utiliser le Namespace Pid

En remplaçant le namespace par l'hôte, le conteneur peut également voir tous les autres processus en cours d'exécution sur le système.

```
> docker run -it --pid=host alpine ps aux
```

```
296 root      0:00 -bash
299 root      0:02 docker-containerd --config /var/run/docker/containerd/cont
3026 root     0:00 [kworker/u2:0]
497 root      0:00 [kworker/0:0]
1884 root     0:00 [kworker/0:1]
1942 root     0:00 [kworker/0:2]
2183 root     0:00 docker run -it --pid=host alpine ps aux
2198 root     0:00 docker-containerd-shim -namespace moby -workdir /var/lib/d
2227 root     0:00 ps aux
2270 root     0:00 [kworker/0:3]
2289 root     0:00 {dhcpcd-run-hook} /bin/sh /usr/lib/dhcpcd/dhcpcd-run-hooks
2299 root     0:00 docker-runc --root /var/run/docker/runtime-runc/moby --log
2302 root     0:00 [dhcpcd-run-hook]
```

## 5- Partage de Namespaces

---

- Fournir aux conteneurs l'accès à l'espace de noms de l'hôte est parfois nécessaire, comme pour les outils de débogage, mais est considéré comme une mauvaise pratique. En effet, vous sortez du modèle de sécurité du conteneur, ce qui peut introduire des vulnérabilités
- Au lieu de cela, si nécessaire, utilisez un espace de noms partagé pour fournir un accès uniquement aux espaces de noms dont le conteneur a besoin.

## 5- Partage de Namespaces

### Exemple:

- Le premier conteneur démarre un serveur Nginx. Cela définira un nouvel espace de noms de réseau et de processus.
- Le serveur Nginx se liera au port 80 de l'interface réseau nouvellement définie.

> **docker run -d --name http nginx:alpine**

```
[root@ghost01 ~]# docker run -d --name http nginx:alpine
Unable to find image 'nginx:alpine' locally
alpine: Pulling from library/nginx
801bfaa63ef2: Pull complete
b1242e25d284: Pull complete
7453d3e6b909: Pull complete
07ce7418c4f8: Pull complete
e295e0624aa3: Pull complete
Digest: sha256:c2ce58e024275728b00a554ac25628af25c54782865b3487b11c21caf7fabda
Status: Downloaded newer image for nginx:alpine
b588fb6334b3293231c850114b2af4704df914b2fcfa6b0e43d74cf95fc6a2d5
```

## 5- Partage de Namespaces

- Le premier conteneur démarre un serveur Nginx. Cela définira un nouvel espace de noms de réseau et de processus. Le serveur Nginx se liera au port 80 de l'interface réseau nouvellement définie.

**> docker run -d --name http nginx:alpine**

```
[root @ host01 ~] # docker run -d --name http nginx: alpin  
865f7f6c19422e34ed55dfd940d2e44bb877def13761927c509fd21327c4d495
```

- D'autres conteneurs peuvent désormais réutiliser cet espace de noms à l'aide de la syntaxe **container: <nom>**.
- Ci-dessous, la commande curl peut accéder au serveur HTTP s'exécutant sur localhost car ils partagent la même interface réseau.

## 5- Partage de Namespaces

```
> docker run --net=container:http benhall/curl curl -s
localhost
docker run --net=container:http benhall/curl
curl -s localhost
```

```
[root@host01 ~]# docker run --net=container:http benhall/curl curl -s localhost
Unable to find image 'benhall/curl:latest' locally
latest: Pulling from benhall/curl
12b41071e6ce: Pull complete
fb1cef6edba2: Pull complete
a3ed95caeb02: Pull complete
Digest: sha256:637b3e063550593071b172916c896ea3590d4ecd67ae2399e54672b5770e114e
Status: Downloaded newer image for benhall/curl:latest
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
}
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
```

## 5- Partage de Namespaces

```
> docker run --pid=container:http alpine ps aux
```

- Il peut également voir et s'interfacer avec les processus dans le conteneur partagé.

```
[root@host01 ~]# docker run --pid=container:http alpine ps aux
 PID  USER      TIME  COMMAND
   1 root      0:00 nginx: master process nginx -g daemon off;
   29 101      0:00 nginx: worker process
   30 root      0:00 ps aux
```

- Ceci est utile pour les outils de débogage, tels que **strace**.
- Cela vous permet d'accorder plus d'autorisations à des conteneurs spécifiques sans modifier ni redémarrer l'application.

# Conclusion

## Objectifs de la formation

---

- Comprendre l'état de l'art des solutions d'orchestration de containers et de leur écosystème ainsi que la mise en œuvre une plateforme de type CaaS (Container as a Service).
- Comprendre le fonctionnement, la mise en place et l'utilisation de conteneurs dans une organisation.
- Ce qu'on a couvert:
  - 1) Les fondamentaux
  - 2) Docker
  - 3) Kubernetes, orchestrateur de containers
  - 4) Container as a Service (CaaS)
  - 5) Sécurité du CaaS/Kubernetes/Docker

A bientôt ☺

